

# Machine Learning Assignment 2

By: Nikolaos Gournakis

Student ID: it22023

## Loading the data

For the loading of the data I created the class MLProject2Dataset which is a subclass of the Dataset class, this is important so that we can easily use the Dataloader class later on.

```
In [1]: # Imports for MLProject2Dataset
import glob
import os
import PIL.Image
import pandas as pd
import torchvision
from torch.utils.data import Dataset
```

```
In [2]: class MLProject2Dataset(Dataset):
    # I added type-annotations to make the code more readable/usable
    def __init__(self, data_dir: str, metadata_fname='metadata.csv', transform: torchvision.transforms = None):
        self.data_dir = data_dir
        self.transform = transform

        self.df = pd.DataFrame(columns=['image_id', 'path'])

        # Get all files in data_dir
        files = glob.glob(data_dir + '/*/*.jpg')
        # If on windows, replace backslash with forward slash
        # on linux, this does nothing
        files = [file.replace(os.sep, '/') for file in files]
        # The line commented bellow was used for testing on a smaller dataset
        # files = files[:3000]
```

```

# Add all files to dataframe
for file in files:
    # the `file` is in the format `data_dir/every/sub/directory/image_id.jpg`
    # that's why we split it by `/` and get the last element which will always be
    # `image_id.jpg`. Then just remove the `.jpg` extension
    image_id = file.split('/')[ -1 ].split('.')[ 0 ]

    # `file` is the full path to the image, so we just create a local variable for readability
    path = file

    # Add new row to dataframe
    self.df.loc[len(self.df)] = [image_id, path]

# Add metadata
self.metadata = pd.read_csv(os.path.join(data_dir, metadata_fname))
dx_categorical = pd.Categorical(self.metadata['dx'])

# number_to_label_array is used to convert the prediction from a number back to
# "Human Readable" format, this also used to make the confusion matrix have the labels as
# cells instead of just numbers.
# Which are the seven string labels given in the assignment
self.number_to_label_array = [dx for dx in dx_categorical.categories]

# dx_categorical.codes is a list of numbers from 0 to 6, which represent the seven labels
# Effectively, we are converting the seven string labels to a number encoding from 0 to 6
self.metadata['dx'] = dx_categorical.codes

# Merge metadata with dataframe
# This creates a new dataframe with all the columns from both dataframes
# and merges them where the `image_id` is the same
self.df = self.df.merge(self.metadata, on='image_id')

# In the assignment, we are asked for the dataframe to have columns = ['image_id', 'path', 'dx']
# so we drop the excess columns
self.df = self.df.drop(columns=["lesion_id", "dx_type", "age", "sex", "localization"])

def __len__(self):
    return len(self.df)

def __getitem__(self, idx):
    # Get image path

```

```

image_path = self.df.iloc[idx]['path']

# Load image
image = PIL.Image.open(image_path)

# Change pixel values to float
# This is not really needed, but I do it just in case
image = image.convert('RGB')

# We can convert the image to tensor here or later during the data loading
# image = torchvision.transforms.ToTensor()(image)

# Get Label
label = self.df.iloc[idx]['dx']

if self.transform:
    image = self.transform(image)

return image, label

```

```

In [3]: # We use this function for the Simple CNN and the Complex CNN
# with different values for M and N, but the rest of the code is the same
def create_transform(M: int, N: int) -> torchvision.transforms:
    return torchvision.transforms.Compose([
        torchvision.transforms.Resize((M, N)),
        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
    ])

```

## train\_net

```

In [4]: import torch
from torch.utils.data import DataLoader
from torch import nn, optim

# These lists are used to store the loss and accuracy for each print_period
# so that we can plot them later
val_loss = []
train_loss = []

```

```

val_accuracy = []
train_accuracy = []

def train_net(model: nn.Module, trainloader: DataLoader, valloader: DataLoader = None, epochs: int = 10,
              optimizer: optim.Optimizer = None, loss: nn.modules.loss = None, device: str = 'cpu',
              print_period: int = 10) -> None:
    # Reset the lists so that we don't append to them if we train the model again / train another model
    global val_loss, train_loss, val_accuracy, train_accuracy
    val_loss = []
    train_loss = []
    val_accuracy = []
    train_accuracy = []

    for epoch in range(epochs):
        total = 0
        correct = 0
        running_loss = 0.0

        # batch is the index of the batch,
        # X is `BatchSize` number of images, y is `BatchSize` number of labels
        # In my experiments, I used `BatchSize` = 32
        for batch, (X, y) in enumerate(trainloader, 0):
            # Set model to train mode
            model.train()

            # Move the images and labels to the device specified
            X = X.to(device)
            y = y.to(device)

            # Zero the gradients
            optimizer.zero_grad()

            # Forward pass
            y_pred = model(X)

            # If I don't cast y to long, I get an error:
            # RuntimeError: "nll_loss_forward_reduce_cuda_kernel_2d_index" not implemented for 'Char'
            # solution: https://stackoverflow.com/a/71126544/13250408
            current_loss = loss(y_pred, y.long())

            # Backward pass
            current_loss.backward()

```

```
# Update weights
optimizer.step()

# Get the total number of images processed so far
total += y.size(0)
# Get the number of correct predictions so far
correct += (y_pred.argmax(1) == y).type(torch.float).sum().item()
# Get the running loss
running_loss += current_loss.item()

if batch % print_period == print_period - 1:

    avg_loss = running_loss / print_period
    avg_accuracy = correct / total * 100
    # Append the loss and accuracy for the current print_period
    train_loss.append(avg_loss)
    train_accuracy.append(avg_accuracy)
    # This is equal to `total` but for the validation set
    total_val = 0
    # This is equal to `correct` but for the validation set
    correct_val = 0
    # This is equal to `running_loss` but for the validation set
    running_loss_val = 0.0

    # If we have a validation set, we also want to print the validation loss and accuracy
    # noinspection GrazieInspection
    if valloader is not None:
        # Set model to evaluation mode
        model.eval()
        # We don't want to update the weights, so we disable the gradient calculation
        with torch.no_grad():
            # Iterate over the validation set
            for (X_val, y_val) in valloader:
                # Move the images and labels to the device specified
                X_val = X_val.to(device)
                y_val = y_val.to(device)
                # Get the predictions
                y_pred_val = model(X_val)
                # Get the total number of validation images processed so far
                total_val += y_val.size(0)
                # Get the running loss for the validation set
```

```

        running_loss_val += loss(y_pred_val, y_val.long()).item()
        # Get the number of correct predictions for the validation set
        correct_val += (y_pred_val.argmax(1) == y_val).type(torch.float).sum().item()

        # To get the average loss, we divide the running loss by the `print_period`
        # to do the equivalent for the validation set, we need to do it differently
        # because we test all of the validation set at once, not in batches
        # so to get the average loss for the validation set, we divide the running loss
        # with the total number of batches in the validation set, which is equal to
        # the length of the validation loader
        avg_loss_val = running_loss_val / len(valloader)
        avg_accuracy_val = correct_val / total_val * 100
        # Append the loss and accuracy for the current print_period
        val_loss.append(avg_loss_val)
        val_accuracy.append(avg_accuracy_val)
        # Print the training loss, training accuracy, validation loss and validation accuracy
        print(
            f'[{epoch + 1}, {batch + 1}] loss: {avg_loss} | accuracy: {avg_accuracy:.2f}% | '
            f' val_loss: {avg_loss_val} | val_accuracy: {avg_accuracy_val:.2f}%'
        )

        # If we don't have a validation set, we only print the training loss and accuracy
    else:
        print(f'[{epoch + 1}, {batch + 1}] loss: {avg_loss} | accuracy: {avg_accuracy :.4f}%')
    # Reset the running loss, total and correct for the next print_period
    running_loss = 0.0
    total = 0
    correct = 0

```

## test\_net

```

In [5]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
def test_net(model: nn.Module, testloader: DataLoader, number_to_label_array: list[str],
            loss: nn.modules.loss = None, device: str = 'cpu',) -> None:
    # Set model to evaluation mode
    model.eval()
    correct = 0
    loss_test = 0.0

```

```

y_trues = []
y_preds = []
# We don't want to update the weights, so we disable the gradient calculation
with torch.no_grad():
    # Iterate over the test set
    for (X, y) in testloader:
        # Move the images and labels to the device specified
        X = X.to(device)
        y = y.to(device)
        # Get the predictions
        y_pred_test = model(X)

        y_trues.extend(y.cpu().numpy())
        y_preds.extend(y_pred_test.argmax(1).cpu().numpy())
        # Get the total number of test images processed so far
        loss_test += loss(y_pred_test, y.long())
        # Get the number of correct predictions for the test set
        correct += (y_pred_test.argmax(1) == y).type(torch.float).sum().item()
# Print the test loss and test accuracy
print(f"Test loss: {loss_test / len(testloader)} | "
      f" Test accuracy: {correct / len(testloader.dataset) * 100:.2f}% ")

# Create the confusion matrix
conf_matrix = confusion_matrix(y_trues, y_preds)

# Turn the confusion matrix into a dataframe for easier plotting
df_cm = pd.DataFrame(conf_matrix, index=number_to_label_array,
                     columns=number_to_label_array)

# Plot the confusion matrix
sns.heatmap(df_cm, annot=True, fmt='g', cmap='PiYG')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

```

In [6]: # Setting the device
device = 'cuda' if torch.cuda.is_available() else 'cpu'

```

## Simple CNN

```
In [7]: class SimpleConvNN(torch.nn.Module):
    def __init__(self):
        super(SimpleConvNN, self).__init__()
        # In the specification of the assignment, it says that
        # the images will be resized to [50 x 62] x 3 channels
        # the comments bellow are the output image size after each layer
        self.conv1 = torch.nn.Conv2d(3, 32, 3) # 48 x 60 x 32
        # MaxPool 2x2 24 x 30 x 32
        self.conv2 = torch.nn.Conv2d(32, 64, 3) # 22 x 28 x 64
        # MaxPool 2x2 11 x 14 x 64
        self.conv3 = torch.nn.Conv2d(64, 64, 3) # 9 x 12 x 64
        # MaxPool 2x2 4 x 6 x 64
        self.fc = torch.nn.Linear(4 * 6 * 64, 7)
        self.pool = torch.nn.MaxPool2d(2)

    def forward(self, x):
        # Ordering the Layers as given in the assignment
        x = self.pool(torch.nn.functional.relu(self.conv1(x)))
        x = self.pool(torch.nn.functional.relu(self.conv2(x)))
        x = self.pool(torch.nn.functional.relu(self.conv3(x)))
        x = torch.flatten(x, 1)
        x = self.fc(x)
        return x
```

## Now let's train the model and test it

```
In [8]: from torchsummary import summary

# Constants from the assignment
M = 50
N = 62
# Create the transforms
dataset_transforms = create_transform(M, N)
# Create the dataset
dataset = MLProject2Dataset(data_dir='data', metadata_fname='metadata.csv', transform=dataset_transforms)
# Split the dataset into train, validation and test
train_dataset, val_dataset, test_dataset = torch.utils.data.random_split(dataset, [.6, .1, .3],
                                                                           generator=torch.Generator().manual_seed(42))
```



```
)  
  
# Create the dataloaders  
BATCH_SIZE = 32  
trainloader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)  
valloader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=True)  
testloader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=True)  
  
# Create the model  
model = SimpleConvNN()  
model = model.to(device)  
# Print the model summary  
summary(model, (3, M, N))  
# Create the optimizer  
optimizer = optim.SGD(model.parameters(), lr=0.1)  
# Create the loss function  
loss = nn.CrossEntropyLoss()  
# Train the model  
train_net(model, trainloader, valloader, epochs=20, optimizer=optimizer, loss=loss, device=device)  
# This cell took 38 minutes to run on my machine
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 48, 60]	896
MaxPool2d-2	[-1, 32, 24, 30]	0
Conv2d-3	[-1, 64, 22, 28]	18,496
MaxPool2d-4	[-1, 64, 11, 14]	0
Conv2d-5	[-1, 64, 9, 12]	36,928
MaxPool2d-6	[-1, 64, 4, 6]	0
Linear-7	[-1, 7]	10,759

Total params: 67,079

Trainable params: 67,079

Non-trainable params: 0

Input size (MB): 0.04

Forward/backward pass size (MB): 1.32

Params size (MB): 0.26

Estimated Total Size (MB): 1.61

[1, 10]	loss: 1.518793785572052		accuracy: 58.13%		val_loss: 1.2464358545839787		val_accuracy: 65.13%
[1, 20]	loss: 1.3173851668834686		accuracy: 68.75%		val_loss: 1.1391955614089966		val_accuracy: 65.13%
[1, 30]	loss: 1.1066274404525758		accuracy: 66.25%		val_loss: 1.0768212135881186		val_accuracy: 65.13%
[1, 40]	loss: 1.0149292349815369		accuracy: 69.38%		val_loss: 1.0525662787258625		val_accuracy: 65.13%
[1, 50]	loss: 1.139615136384964		accuracy: 61.25%		val_loss: 1.1996386479586363		val_accuracy: 65.33%
[1, 60]	loss: 1.097195464372635		accuracy: 65.62%		val_loss: 1.0418086061254144		val_accuracy: 65.13%
[1, 70]	loss: 1.0688795149326324		accuracy: 65.94%		val_loss: 0.9865861423313618		val_accuracy: 65.13%
[1, 80]	loss: 0.9770754396915435		accuracy: 67.81%		val_loss: 1.142379056662321		val_accuracy: 65.13%
[1, 90]	loss: 1.0007180273532867		accuracy: 63.75%		val_loss: 1.0331340655684471		val_accuracy: 65.13%
[1, 100]	loss: 0.9722979307174683		accuracy: 65.94%		val_loss: 1.2319824155420065		val_accuracy: 65.33%
[1, 110]	loss: 1.0412011325359345		accuracy: 66.88%		val_loss: 0.9679132308810949		val_accuracy: 65.43%
[1, 120]	loss: 0.9695410609245301		accuracy: 64.69%		val_loss: 0.958821153268218		val_accuracy: 65.53%
[1, 130]	loss: 0.9361404955387116		accuracy: 66.25%		val_loss: 0.9314349181950092		val_accuracy: 65.93%
[1, 140]	loss: 1.058273833990097		accuracy: 65.31%		val_loss: 0.950197359547019		val_accuracy: 65.43%
[1, 150]	loss: 0.8888957321643829		accuracy: 68.12%		val_loss: 0.9497997667640448		val_accuracy: 66.03%
[1, 160]	loss: 0.9848702192306519		accuracy: 66.88%		val_loss: 0.9146623685956001		val_accuracy: 66.23%
[1, 170]	loss: 0.9308148145675659		accuracy: 68.12%		val_loss: 0.9495742432773113		val_accuracy: 66.23%
[1, 180]	loss: 0.8736093640327454		accuracy: 70.94%		val_loss: 1.0582870300859213		val_accuracy: 65.13%
[2, 10]	loss: 0.8881802827119827		accuracy: 70.62%		val_loss: 0.9386426508426666		val_accuracy: 67.03%
[2, 20]	loss: 0.9549904286861419		accuracy: 66.25%		val_loss: 0.918483629822731		val_accuracy: 65.73%
[2, 30]	loss: 0.8058034777641296		accuracy: 71.25%		val_loss: 0.9602652927860618		val_accuracy: 65.93%
[2, 40]	loss: 0.972988098859787		accuracy: 65.94%		val_loss: 0.9043727740645409		val_accuracy: 66.83%

[2, 50]	loss: 0.8965781688690185	accuracy: 68.12%	val_loss: 0.8962981477379799	val_accuracy: 65.33%
[2, 60]	loss: 0.9681354999542237	accuracy: 66.25%	val_loss: 0.9129490694031119	val_accuracy: 65.93%
[2, 70]	loss: 0.8769855678081513	accuracy: 70.00%	val_loss: 0.9951959066092968	val_accuracy: 65.73%
[2, 80]	loss: 0.8566119372844696	accuracy: 70.31%	val_loss: 0.9078674167394638	val_accuracy: 64.64%
[2, 90]	loss: 0.8813894808292388	accuracy: 70.31%	val_loss: 1.012676689773798	val_accuracy: 66.03%
[2, 100]	loss: 0.9491555988788605	accuracy: 67.50%	val_loss: 0.8930072579532862	val_accuracy: 66.33%
[2, 110]	loss: 0.9263675212860107	accuracy: 65.31%	val_loss: 0.891107683070004	val_accuracy: 67.23%
[2, 120]	loss: 0.9533082544803619	accuracy: 68.12%	val_loss: 0.8923446051776409	val_accuracy: 66.43%
[2, 130]	loss: 0.9229607105255127	accuracy: 64.38%	val_loss: 0.8439620733261108	val_accuracy: 68.93%
[2, 140]	loss: 0.8775909841060638	accuracy: 68.12%	val_loss: 0.9100127127021551	val_accuracy: 67.23%
[2, 150]	loss: 1.0363579750061036	accuracy: 63.44%	val_loss: 0.8825412821024656	val_accuracy: 67.13%
[2, 160]	loss: 0.7596043646335602	accuracy: 73.12%	val_loss: 0.8908407762646675	val_accuracy: 65.73%
[2, 170]	loss: 0.8432924270629882	accuracy: 70.00%	val_loss: 0.8675894299522042	val_accuracy: 66.33%
[2, 180]	loss: 0.973846846818924	accuracy: 62.50%	val_loss: 0.8677255157381296	val_accuracy: 66.93%
[3, 10]	loss: 0.8693842113018035	accuracy: 70.00%	val_loss: 0.9068197906017303	val_accuracy: 68.33%
[3, 20]	loss: 0.9245539367198944	accuracy: 69.69%	val_loss: 0.8482250161468983	val_accuracy: 66.43%
[3, 30]	loss: 0.8058761715888977	accuracy: 71.88%	val_loss: 0.8879662547260523	val_accuracy: 66.33%
[3, 40]	loss: 0.8018431305885315	accuracy: 70.62%	val_loss: 0.8751243390142918	val_accuracy: 68.73%
[3, 50]	loss: 0.8726075828075409	accuracy: 70.94%	val_loss: 1.0664616413414478	val_accuracy: 64.44%
[3, 60]	loss: 0.9015521109104156	accuracy: 64.38%	val_loss: 0.8642947552725673	val_accuracy: 68.33%
[3, 70]	loss: 0.9306747198104859	accuracy: 63.75%	val_loss: 0.8376922477036715	val_accuracy: 68.03%
[3, 80]	loss: 0.7847657382488251	accuracy: 71.88%	val_loss: 0.8478030357509851	val_accuracy: 67.03%
[3, 90]	loss: 0.9061262130737304	accuracy: 66.88%	val_loss: 0.8297563698142767	val_accuracy: 68.93%
[3, 100]	loss: 0.870602947473526	accuracy: 66.56%	val_loss: 0.8714733850210905	val_accuracy: 67.23%
[3, 110]	loss: 0.8912302553653717	accuracy: 66.25%	val_loss: 0.8362489119172096	val_accuracy: 69.13%
[3, 120]	loss: 0.8051473140716553	accuracy: 71.88%	val_loss: 0.843952676281333	val_accuracy: 68.83%
[3, 130]	loss: 0.9509417712688446	accuracy: 66.25%	val_loss: 0.8772696387022734	val_accuracy: 69.43%
[3, 140]	loss: 0.8863458216190339	accuracy: 69.06%	val_loss: 0.8979610204696655	val_accuracy: 65.53%
[3, 150]	loss: 0.8993420749902725	accuracy: 67.19%	val_loss: 0.9754824060946703	val_accuracy: 68.53%
[3, 160]	loss: 0.8832664847373962	accuracy: 66.25%	val_loss: 0.8510226430371404	val_accuracy: 68.03%
[3, 170]	loss: 0.7881470412015915	accuracy: 73.12%	val_loss: 0.9633007626980543	val_accuracy: 61.24%
[3, 180]	loss: 0.8733445048332215	accuracy: 67.19%	val_loss: 0.9262318713590503	val_accuracy: 65.63%
[4, 10]	loss: 0.8964118480682373	accuracy: 68.44%	val_loss: 0.8459508959203959	val_accuracy: 66.93%
[4, 20]	loss: 0.7726080775260925	accuracy: 74.38%	val_loss: 0.8318781824782491	val_accuracy: 67.53%
[4, 30]	loss: 0.7636393129825592	accuracy: 72.50%	val_loss: 0.8219802044332027	val_accuracy: 68.33%
[4, 40]	loss: 0.7814334988594055	accuracy: 71.56%	val_loss: 0.8352053444832563	val_accuracy: 67.13%
[4, 50]	loss: 0.9669749915599823	accuracy: 64.69%	val_loss: 0.9148869477212429	val_accuracy: 67.03%
[4, 60]	loss: 0.8185674667358398	accuracy: 70.94%	val_loss: 0.8547702115029097	val_accuracy: 68.83%
[4, 70]	loss: 0.7825084030628204	accuracy: 71.88%	val_loss: 0.8541211653500795	val_accuracy: 68.73%
[4, 80]	loss: 0.7922193884849549	accuracy: 72.19%	val_loss: 0.8031186126172543	val_accuracy: 68.43%
[4, 90]	loss: 0.8617696523666382	accuracy: 67.19%	val_loss: 0.8559656739234924	val_accuracy: 67.73%
[4, 100]	loss: 0.7899049997329712	accuracy: 70.62%	val_loss: 0.8143490580841899	val_accuracy: 69.23%

[4, 110]	loss: 0.7909148812294007	accuracy: 73.12%	val_loss: 0.839736714027822	val_accuracy: 68.63%
[4, 120]	loss: 0.9073679387569428	accuracy: 69.38%	val_loss: 0.8112899251282215	val_accuracy: 70.23%
[4, 130]	loss: 0.8610491096973419	accuracy: 67.81%	val_loss: 0.8687341436743736	val_accuracy: 68.13%
[4, 140]	loss: 0.8330842733383179	accuracy: 69.38%	val_loss: 0.812447052448988	val_accuracy: 70.33%
[4, 150]	loss: 0.9180149853229522	accuracy: 66.25%	val_loss: 0.8030542824417353	val_accuracy: 70.13%
[4, 160]	loss: 0.8498059511184692	accuracy: 70.00%	val_loss: 0.8029752094298601	val_accuracy: 70.63%
[4, 170]	loss: 0.8327560901641846	accuracy: 70.94%	val_loss: 0.7931884033605456	val_accuracy: 70.53%
[4, 180]	loss: 0.8189090669155121	accuracy: 68.44%	val_loss: 0.7922591688111424	val_accuracy: 69.93%
[5, 10]	loss: 0.6943336248397827	accuracy: 76.56%	val_loss: 0.8273032214492559	val_accuracy: 70.33%
[5, 20]	loss: 0.8296578228473663	accuracy: 69.06%	val_loss: 0.8268737401813269	val_accuracy: 67.03%
[5, 30]	loss: 0.8139205902814866	accuracy: 69.06%	val_loss: 0.8196424394845963	val_accuracy: 69.53%
[5, 40]	loss: 0.8774439513683319	accuracy: 69.06%	val_loss: 0.8302071988582611	val_accuracy: 68.33%
[5, 50]	loss: 0.9478014469146728	accuracy: 64.06%	val_loss: 0.8248966718092561	val_accuracy: 70.23%
[5, 60]	loss: 0.8615355134010315	accuracy: 69.06%	val_loss: 0.8212379980832338	val_accuracy: 69.23%
[5, 70]	loss: 0.7926372230052948	accuracy: 69.69%	val_loss: 1.0204764977097511	val_accuracy: 62.14%
[5, 80]	loss: 0.8106517910957336	accuracy: 69.69%	val_loss: 0.7966296840459108	val_accuracy: 69.83%
[5, 90]	loss: 0.7340528309345246	accuracy: 72.19%	val_loss: 0.7859837468713522	val_accuracy: 71.03%
[5, 100]	loss: 0.7939515560865402	accuracy: 69.38%	val_loss: 0.8159124702215195	val_accuracy: 69.13%
[5, 110]	loss: 0.8378702044487	accuracy: 69.06%	val_loss: 0.7837748015299439	val_accuracy: 71.63%
[5, 120]	loss: 0.8507627487182617	accuracy: 71.56%	val_loss: 0.8011882528662682	val_accuracy: 69.93%
[5, 130]	loss: 0.7834202826023102	accuracy: 71.25%	val_loss: 0.8648223374038935	val_accuracy: 68.83%
[5, 140]	loss: 0.7988188207149506	accuracy: 72.50%	val_loss: 0.7861625403165817	val_accuracy: 69.83%
[5, 150]	loss: 0.7317732393741607	accuracy: 73.75%	val_loss: 0.8254772759974003	val_accuracy: 70.53%
[5, 160]	loss: 0.7997078239917755	accuracy: 70.00%	val_loss: 0.7750689536333084	val_accuracy: 70.43%
[5, 170]	loss: 0.9105548083782196	accuracy: 64.38%	val_loss: 0.7684560809284449	val_accuracy: 70.73%
[5, 180]	loss: 0.8114823758602142	accuracy: 74.06%	val_loss: 0.7881467174738646	val_accuracy: 70.93%
[6, 10]	loss: 0.7293206214904785	accuracy: 73.75%	val_loss: 0.7530530923977494	val_accuracy: 71.33%
[6, 20]	loss: 0.7692529916763305	accuracy: 69.06%	val_loss: 0.8168331105262041	val_accuracy: 69.63%
[6, 30]	loss: 0.6856015205383301	accuracy: 74.69%	val_loss: 0.7832402596250176	val_accuracy: 69.53%
[6, 40]	loss: 0.7093329280614853	accuracy: 74.38%	val_loss: 0.7900243597105145	val_accuracy: 70.23%
[6, 50]	loss: 0.7612624347209931	accuracy: 72.19%	val_loss: 0.8026740187779069	val_accuracy: 69.83%
[6, 60]	loss: 0.7682025134563446	accuracy: 70.31%	val_loss: 0.7949044657871127	val_accuracy: 70.03%
[6, 70]	loss: 0.8108818888664245	accuracy: 70.31%	val_loss: 0.7516264021396637	val_accuracy: 71.83%
[6, 80]	loss: 0.8695433497428894	accuracy: 68.44%	val_loss: 0.7606173325330019	val_accuracy: 71.33%
[6, 90]	loss: 0.8544639945030212	accuracy: 68.12%	val_loss: 0.971032053232193	val_accuracy: 60.54%
[6, 100]	loss: 0.7887440085411072	accuracy: 72.50%	val_loss: 0.7860160954296589	val_accuracy: 70.53%
[6, 110]	loss: 0.7232416808605194	accuracy: 72.50%	val_loss: 0.7712192134931684	val_accuracy: 72.13%
[6, 120]	loss: 0.7800049692392349	accuracy: 68.75%	val_loss: 0.7719339597970247	val_accuracy: 70.73%
[6, 130]	loss: 0.7867901444435119	accuracy: 72.19%	val_loss: 0.7700664168223739	val_accuracy: 70.03%
[6, 140]	loss: 0.7703287363052368	accuracy: 72.50%	val_loss: 0.8238052688539028	val_accuracy: 69.83%
[6, 150]	loss: 0.8538231074810028	accuracy: 69.38%	val_loss: 0.7674029665067792	val_accuracy: 71.03%
[6, 160]	loss: 0.6847216010093689	accuracy: 73.44%	val_loss: 0.8001913335174322	val_accuracy: 71.33%

[6, 170]	loss: 0.7412172734737397	accuracy: 72.81%	val_loss: 0.7916136980056763	val_accuracy: 69.93%
[6, 180]	loss: 0.8604855358600616	accuracy: 69.38%	val_loss: 0.7771682403981686	val_accuracy: 72.53%
[7, 10]	loss: 0.7201659739017486	accuracy: 72.19%	val_loss: 0.7655354756861925	val_accuracy: 71.53%
[7, 20]	loss: 0.8270796656608581	accuracy: 70.62%	val_loss: 0.7727065803483129	val_accuracy: 71.13%
[7, 30]	loss: 0.7508438289165497	accuracy: 70.94%	val_loss: 0.7667716834694147	val_accuracy: 71.33%
[7, 40]	loss: 0.7167212486267089	accuracy: 73.75%	val_loss: 0.7659776322543621	val_accuracy: 71.53%
[7, 50]	loss: 0.7115629225969314	accuracy: 76.25%	val_loss: 0.8114396939054132	val_accuracy: 71.73%
[7, 60]	loss: 0.8154988765716553	accuracy: 69.38%	val_loss: 0.7904073316603899	val_accuracy: 71.43%
[7, 70]	loss: 0.6974685281515122	accuracy: 75.62%	val_loss: 0.7728052232414484	val_accuracy: 70.93%
[7, 80]	loss: 0.8187405109405518	accuracy: 69.38%	val_loss: 0.759864516556263	val_accuracy: 70.73%
[7, 90]	loss: 0.7157119929790496	accuracy: 75.31%	val_loss: 0.7406004630029202	val_accuracy: 72.53%
[7, 100]	loss: 0.5760733723640442	accuracy: 77.19%	val_loss: 0.9259790424257517	val_accuracy: 66.13%
[7, 110]	loss: 0.7420385748147964	accuracy: 73.44%	val_loss: 0.7510803192853928	val_accuracy: 70.23%
[7, 120]	loss: 0.7888030230998992	accuracy: 70.62%	val_loss: 0.7914383923634887	val_accuracy: 69.63%
[7, 130]	loss: 0.8440999865531922	accuracy: 66.88%	val_loss: 0.7471137382090092	val_accuracy: 70.83%
[7, 140]	loss: 0.7159431368112564	accuracy: 72.50%	val_loss: 0.7437009466812015	val_accuracy: 71.43%
[7, 150]	loss: 0.8987858295440674	accuracy: 68.75%	val_loss: 0.7517248010262847	val_accuracy: 71.33%
[7, 160]	loss: 0.7482125192880631	accuracy: 70.31%	val_loss: 0.7395641524344683	val_accuracy: 72.03%
[7, 170]	loss: 0.7596843123435975	accuracy: 70.00%	val_loss: 0.7621422903612256	val_accuracy: 71.33%
[7, 180]	loss: 0.7990598022937775	accuracy: 72.19%	val_loss: 0.7248687986284494	val_accuracy: 71.83%
[8, 10]	loss: 0.7472058653831481	accuracy: 74.38%	val_loss: 0.7357574244961143	val_accuracy: 71.63%
[8, 20]	loss: 0.7074092537164688	accuracy: 71.88%	val_loss: 0.7107342332601547	val_accuracy: 73.03%
[8, 30]	loss: 0.7300812602043152	accuracy: 72.50%	val_loss: 0.7435651160776615	val_accuracy: 72.53%
[8, 40]	loss: 0.6924116909503937	accuracy: 74.69%	val_loss: 0.7405048850923777	val_accuracy: 72.53%
[8, 50]	loss: 0.7438171476125717	accuracy: 72.19%	val_loss: 0.734918893314898	val_accuracy: 72.63%
[8, 60]	loss: 0.7346929490566254	accuracy: 70.31%	val_loss: 0.8128311792388558	val_accuracy: 67.73%
[8, 70]	loss: 0.790061891078949	accuracy: 67.50%	val_loss: 0.7798128118738532	val_accuracy: 70.33%
[8, 80]	loss: 0.8174928545951843	accuracy: 65.31%	val_loss: 0.7315768972039223	val_accuracy: 72.33%
[8, 90]	loss: 0.7609184622764588	accuracy: 72.50%	val_loss: 0.7726689288392663	val_accuracy: 70.93%
[8, 100]	loss: 0.7715620577335358	accuracy: 71.56%	val_loss: 0.7555542076006532	val_accuracy: 72.13%
[8, 110]	loss: 0.7240578949451446	accuracy: 73.75%	val_loss: 0.7182535585016012	val_accuracy: 72.13%
[8, 120]	loss: 0.7684918165206909	accuracy: 71.88%	val_loss: 0.7245315983891487	val_accuracy: 71.83%
[8, 130]	loss: 0.6249454170465469	accuracy: 80.00%	val_loss: 0.7145494939759374	val_accuracy: 72.43%
[8, 140]	loss: 0.6757012993097306	accuracy: 74.06%	val_loss: 0.7314581628888845	val_accuracy: 72.93%
[8, 150]	loss: 0.7642005205154419	accuracy: 73.44%	val_loss: 0.7061256598681211	val_accuracy: 73.13%
[8, 160]	loss: 0.7017618745565415	accuracy: 75.31%	val_loss: 0.7944187261164188	val_accuracy: 71.83%
[8, 170]	loss: 0.7317767053842544	accuracy: 77.19%	val_loss: 0.7853284943848848	val_accuracy: 71.43%
[8, 180]	loss: 0.7986599743366242	accuracy: 68.44%	val_loss: 0.7153360703960061	val_accuracy: 72.53%
[9, 10]	loss: 0.7119902431964874	accuracy: 75.31%	val_loss: 0.7258126339875162	val_accuracy: 72.23%
[9, 20]	loss: 0.7213380724191666	accuracy: 72.81%	val_loss: 0.7634491845965385	val_accuracy: 73.43%
[9, 30]	loss: 0.7433912038803101	accuracy: 73.75%	val_loss: 0.7658159993588924	val_accuracy: 71.33%
[9, 40]	loss: 0.7678873538970947	accuracy: 71.88%	val_loss: 0.7220282908529043	val_accuracy: 71.63%

[9, 50]	loss: 0.6558604896068573	accuracy: 75.31%	val_loss: 0.7056128429248929	val_accuracy: 72.93%
[9, 60]	loss: 0.7123411208391189	accuracy: 74.38%	val_loss: 0.7661803336814046	val_accuracy: 71.73%
[9, 70]	loss: 0.7399612456560135	accuracy: 72.19%	val_loss: 0.7125594140961766	val_accuracy: 72.33%
[9, 80]	loss: 0.8456421494483948	accuracy: 67.50%	val_loss: 0.7448426224291325	val_accuracy: 73.13%
[9, 90]	loss: 0.6530835926532745	accuracy: 77.19%	val_loss: 0.7160550113767385	val_accuracy: 71.53%
[9, 100]	loss: 0.710214251279831	accuracy: 75.94%	val_loss: 0.7123290412127972	val_accuracy: 72.93%
[9, 110]	loss: 0.6917033508419991	accuracy: 73.12%	val_loss: 0.8749028444290161	val_accuracy: 65.63%
[9, 120]	loss: 0.7414666444063187	accuracy: 71.25%	val_loss: 0.73248325381428	val_accuracy: 72.13%
[9, 130]	loss: 0.7209202706813812	accuracy: 72.19%	val_loss: 0.729455335997045	val_accuracy: 72.33%
[9, 140]	loss: 0.755241310596466	accuracy: 71.56%	val_loss: 0.6835456592962146	val_accuracy: 73.73%
[9, 150]	loss: 0.7318120658397674	accuracy: 73.75%	val_loss: 0.7952764229848981	val_accuracy: 70.63%
[9, 160]	loss: 0.657338747382164	accuracy: 77.50%	val_loss: 0.7002133000642061	val_accuracy: 73.43%
[9, 170]	loss: 0.6631748795509338	accuracy: 75.94%	val_loss: 0.7222733767703176	val_accuracy: 73.63%
[9, 180]	loss: 0.6594358712434769	accuracy: 73.12%	val_loss: 0.7528277151286602	val_accuracy: 72.13%
[10, 10]	loss: 0.7912569761276245	accuracy: 68.75%	val_loss: 0.765833580866456	val_accuracy: 72.73%
[10, 20]	loss: 0.67460917532444	accuracy: 75.62%	val_loss: 0.7383092353120446	val_accuracy: 72.03%
[10, 30]	loss: 0.7668874740600586	accuracy: 71.88%	val_loss: 0.7020319998264313	val_accuracy: 73.53%
[10, 40]	loss: 0.5933469235897064	accuracy: 76.56%	val_loss: 0.6806788612157106	val_accuracy: 73.33%
[10, 50]	loss: 0.6881805062294006	accuracy: 78.75%	val_loss: 0.7576083978638053	val_accuracy: 73.03%
[10, 60]	loss: 0.6805338025093078	accuracy: 74.38%	val_loss: 0.706424112431705	val_accuracy: 72.03%
[10, 70]	loss: 0.7163186699151993	accuracy: 76.25%	val_loss: 0.7252379199489951	val_accuracy: 72.53%
[10, 80]	loss: 0.6529519617557525	accuracy: 74.06%	val_loss: 0.7161224232986569	val_accuracy: 72.73%
[10, 90]	loss: 0.7548629224300385	accuracy: 71.88%	val_loss: 0.6887117270380259	val_accuracy: 73.03%
[10, 100]	loss: 0.6154628098011017	accuracy: 77.50%	val_loss: 0.7806516317650676	val_accuracy: 70.63%
[10, 110]	loss: 0.7019013822078705	accuracy: 77.19%	val_loss: 0.6912295473739505	val_accuracy: 74.73%
[10, 120]	loss: 0.7110175430774689	accuracy: 75.31%	val_loss: 0.7406770568341017	val_accuracy: 72.93%
[10, 130]	loss: 0.7096666127443314	accuracy: 74.38%	val_loss: 0.6961056375876069	val_accuracy: 72.73%
[10, 140]	loss: 0.7037199079990387	accuracy: 75.31%	val_loss: 0.7036279514431953	val_accuracy: 71.93%
[10, 150]	loss: 0.7470318853855134	accuracy: 71.56%	val_loss: 0.7296681785956025	val_accuracy: 72.93%
[10, 160]	loss: 0.6724254995584488	accuracy: 75.00%	val_loss: 0.7148599708452821	val_accuracy: 73.33%
[10, 170]	loss: 0.6460056304931641	accuracy: 78.12%	val_loss: 0.7627247888594866	val_accuracy: 73.13%
[10, 180]	loss: 0.6986802279949188	accuracy: 71.88%	val_loss: 0.6879905685782433	val_accuracy: 74.73%
[11, 10]	loss: 0.6793292552232743	accuracy: 72.19%	val_loss: 0.6922942642122507	val_accuracy: 74.43%
[11, 20]	loss: 0.7512198835611343	accuracy: 71.25%	val_loss: 0.715288084000349	val_accuracy: 73.63%
[11, 30]	loss: 0.6893922328948975	accuracy: 75.94%	val_loss: 0.7039272990077734	val_accuracy: 74.33%
[11, 40]	loss: 0.6532567828893662	accuracy: 76.25%	val_loss: 0.7341353902593255	val_accuracy: 71.53%
[11, 50]	loss: 0.5727572709321975	accuracy: 80.00%	val_loss: 0.7120752474293113	val_accuracy: 73.73%
[11, 60]	loss: 0.7334563255310058	accuracy: 73.12%	val_loss: 0.7163863126188517	val_accuracy: 73.33%
[11, 70]	loss: 0.7279835641384125	accuracy: 76.25%	val_loss: 0.6962894788011909	val_accuracy: 73.43%
[11, 80]	loss: 0.7550597846508026	accuracy: 70.00%	val_loss: 0.6997461346909404	val_accuracy: 73.13%
[11, 90]	loss: 0.708294153213501	accuracy: 74.69%	val_loss: 0.7009940827265382	val_accuracy: 74.03%
[11, 100]	loss: 0.6725185066461563	accuracy: 76.56%	val_loss: 0.7723301602527499	val_accuracy: 70.63%

[11, 110]	loss: 0.6225936442613602	accuracy: 76.56%	val_loss: 0.7482059467583895	val_accuracy: 73.93%
[11, 120]	loss: 0.6968424797058106	accuracy: 71.56%	val_loss: 0.6819869503378868	val_accuracy: 71.93%
[11, 130]	loss: 0.7123645216226577	accuracy: 74.69%	val_loss: 0.7200260497629642	val_accuracy: 74.03%
[11, 140]	loss: 0.6579011231660843	accuracy: 76.56%	val_loss: 0.7896305872127414	val_accuracy: 74.33%
[11, 150]	loss: 0.5555790215730667	accuracy: 77.81%	val_loss: 0.7146493038162589	val_accuracy: 73.83%
[11, 160]	loss: 0.7250492215156555	accuracy: 76.56%	val_loss: 0.7603372680023313	val_accuracy: 72.23%
[11, 170]	loss: 0.736522427201271	accuracy: 74.38%	val_loss: 0.8144238069653511	val_accuracy: 70.33%
[11, 180]	loss: 0.6190130889415741	accuracy: 78.75%	val_loss: 0.6995085570961237	val_accuracy: 74.63%
[12, 10]	loss: 0.7211729735136032	accuracy: 74.06%	val_loss: 0.7117769010365009	val_accuracy: 74.33%
[12, 20]	loss: 0.6274710267782211	accuracy: 75.62%	val_loss: 0.7002258412539959	val_accuracy: 73.93%
[12, 30]	loss: 0.5933802843093872	accuracy: 79.06%	val_loss: 0.7749050380662084	val_accuracy: 72.63%
[12, 40]	loss: 0.6219657510519028	accuracy: 76.88%	val_loss: 0.6764179794117808	val_accuracy: 74.53%
[12, 50]	loss: 0.6278844922780991	accuracy: 77.50%	val_loss: 0.6724238786846399	val_accuracy: 74.23%
[12, 60]	loss: 0.6802094221115113	accuracy: 74.06%	val_loss: 0.6892040204256773	val_accuracy: 73.93%
[12, 70]	loss: 0.7436098158359528	accuracy: 72.50%	val_loss: 0.8002385254949331	val_accuracy: 70.43%
[12, 80]	loss: 0.6353487700223923	accuracy: 78.12%	val_loss: 0.709632471203804	val_accuracy: 73.23%
[12, 90]	loss: 0.5829992562532424	accuracy: 78.44%	val_loss: 0.685943434946239	val_accuracy: 74.33%
[12, 100]	loss: 0.5867188215255738	accuracy: 76.88%	val_loss: 0.710184920579195	val_accuracy: 72.93%
[12, 110]	loss: 0.6275329828262329	accuracy: 77.50%	val_loss: 0.8252037083730102	val_accuracy: 72.33%
[12, 120]	loss: 0.6847056448459625	accuracy: 75.94%	val_loss: 0.667448160238564	val_accuracy: 74.43%
[12, 130]	loss: 0.6179657518863678	accuracy: 76.56%	val_loss: 0.7003251137211919	val_accuracy: 73.83%
[12, 140]	loss: 0.6324828088283538	accuracy: 74.69%	val_loss: 0.7060683388262987	val_accuracy: 74.33%
[12, 150]	loss: 0.681729644536972	accuracy: 77.50%	val_loss: 0.6884771967306733	val_accuracy: 74.03%
[12, 160]	loss: 0.7363033652305603	accuracy: 73.75%	val_loss: 0.6894932454451919	val_accuracy: 73.93%
[12, 170]	loss: 0.6367806375026703	accuracy: 73.75%	val_loss: 0.7542187459766865	val_accuracy: 72.33%
[12, 180]	loss: 0.6397121131420136	accuracy: 76.25%	val_loss: 0.7097441339865327	val_accuracy: 73.83%
[13, 10]	loss: 0.6552366822957992	accuracy: 77.19%	val_loss: 0.7108350098133087	val_accuracy: 73.13%
[13, 20]	loss: 0.6448940813541413	accuracy: 72.81%	val_loss: 0.659646418876946	val_accuracy: 74.73%
[13, 30]	loss: 0.6685197502374649	accuracy: 72.81%	val_loss: 0.6622613053768873	val_accuracy: 72.93%
[13, 40]	loss: 0.5990982353687286	accuracy: 76.56%	val_loss: 0.6766712348908186	val_accuracy: 74.43%
[13, 50]	loss: 0.619576495885849	accuracy: 73.75%	val_loss: 0.6668750289827585	val_accuracy: 74.63%
[13, 60]	loss: 0.6077278643846512	accuracy: 76.25%	val_loss: 0.7763166483491659	val_accuracy: 72.33%
[13, 70]	loss: 0.662548103928566	accuracy: 75.31%	val_loss: 0.6663905223831534	val_accuracy: 74.73%
[13, 80]	loss: 0.5698891848325729	accuracy: 78.44%	val_loss: 0.7176613211631775	val_accuracy: 73.03%
[13, 90]	loss: 0.6781819313764572	accuracy: 74.06%	val_loss: 0.6956408619880676	val_accuracy: 73.43%
[13, 100]	loss: 0.7258035182952881	accuracy: 75.31%	val_loss: 0.7739758202806115	val_accuracy: 72.73%
[13, 110]	loss: 0.7156279653310775	accuracy: 73.44%	val_loss: 0.7422378566116095	val_accuracy: 71.63%
[13, 120]	loss: 0.6231303006410599	accuracy: 75.31%	val_loss: 0.6996891964226961	val_accuracy: 74.43%
[13, 130]	loss: 0.6647371232509613	accuracy: 77.50%	val_loss: 0.7035988587886095	val_accuracy: 74.93%
[13, 140]	loss: 0.6621907830238343	accuracy: 75.94%	val_loss: 0.7811008766293526	val_accuracy: 71.83%
[13, 150]	loss: 0.6543887615203857	accuracy: 75.94%	val_loss: 0.7298274524509907	val_accuracy: 73.43%
[13, 160]	loss: 0.6028830647468567	accuracy: 80.00%	val_loss: 0.7104420308023691	val_accuracy: 74.03%

[13, 170]	loss: 0.5485974431037903	accuracy: 78.44%	val_loss: 0.7170111071318388	val_accuracy: 73.93%
[13, 180]	loss: 0.656700250506401	accuracy: 78.75%	val_loss: 0.7163003506138921	val_accuracy: 72.93%
[14, 10]	loss: 0.6210664510726929	accuracy: 76.88%	val_loss: 0.6913860877975821	val_accuracy: 74.63%
[14, 20]	loss: 0.5968304127454758	accuracy: 76.88%	val_loss: 0.8990829922258854	val_accuracy: 64.44%
[14, 30]	loss: 0.5854073107242584	accuracy: 78.75%	val_loss: 0.7114640651270747	val_accuracy: 74.93%
[14, 40]	loss: 0.6887169867753983	accuracy: 73.44%	val_loss: 0.6752008963376284	val_accuracy: 75.42%
[14, 50]	loss: 0.6075867474079132	accuracy: 77.50%	val_loss: 0.7140157101675868	val_accuracy: 73.63%
[14, 60]	loss: 0.5836057841777802	accuracy: 80.00%	val_loss: 0.7193344058468938	val_accuracy: 74.73%
[14, 70]	loss: 0.639396271109581	accuracy: 75.31%	val_loss: 0.7545861825346947	val_accuracy: 70.33%
[14, 80]	loss: 0.6121407896280289	accuracy: 76.88%	val_loss: 0.65510512329638	val_accuracy: 75.92%
[14, 90]	loss: 0.640603193640709	accuracy: 75.31%	val_loss: 0.6859643682837486	val_accuracy: 74.83%
[14, 100]	loss: 0.5649170696735382	accuracy: 77.50%	val_loss: 0.6660568313673139	val_accuracy: 75.02%
[14, 110]	loss: 0.6906864762306213	accuracy: 72.81%	val_loss: 0.6848413050174713	val_accuracy: 74.13%
[14, 120]	loss: 0.6257061004638672	accuracy: 79.38%	val_loss: 0.7237070016562939	val_accuracy: 72.83%
[14, 130]	loss: 0.6568994432687759	accuracy: 76.25%	val_loss: 0.8098530760034919	val_accuracy: 66.73%
[14, 140]	loss: 0.6387811183929444	accuracy: 75.94%	val_loss: 0.7568588331341743	val_accuracy: 72.83%
[14, 150]	loss: 0.613511860370636	accuracy: 79.69%	val_loss: 0.7161957258358598	val_accuracy: 73.93%
[14, 160]	loss: 0.580518102645874	accuracy: 77.50%	val_loss: 0.672249655239284	val_accuracy: 75.22%
[14, 170]	loss: 0.5919616669416428	accuracy: 80.94%	val_loss: 0.844207962974906	val_accuracy: 68.83%
[14, 180]	loss: 0.620694363117218	accuracy: 75.31%	val_loss: 0.6816788255237043	val_accuracy: 73.63%
[15, 10]	loss: 0.7027302950620651	accuracy: 71.88%	val_loss: 0.6635789200663567	val_accuracy: 74.43%
[15, 20]	loss: 0.6318568706512451	accuracy: 76.25%	val_loss: 0.6924440702423453	val_accuracy: 73.93%
[15, 30]	loss: 0.5564908057451248	accuracy: 78.12%	val_loss: 0.6932139787822962	val_accuracy: 74.33%
[15, 40]	loss: 0.5283477246761322	accuracy: 80.00%	val_loss: 0.7637495808303356	val_accuracy: 69.43%
[15, 50]	loss: 0.49295041859149935	accuracy: 80.62%	val_loss: 0.7047851569950581	val_accuracy: 76.12%
[15, 60]	loss: 0.48612076342105864	accuracy: 83.44%	val_loss: 0.7555817607790232	val_accuracy: 74.33%
[15, 70]	loss: 0.6971864491701126	accuracy: 75.94%	val_loss: 0.6864071409218013	val_accuracy: 74.43%
[15, 80]	loss: 0.649814561009407	accuracy: 75.62%	val_loss: 0.7215990936383605	val_accuracy: 73.63%
[15, 90]	loss: 0.6081157863140106	accuracy: 77.81%	val_loss: 0.6754725826904178	val_accuracy: 74.93%
[15, 100]	loss: 0.5660370171070099	accuracy: 82.19%	val_loss: 0.7236782675608993	val_accuracy: 72.93%
[15, 110]	loss: 0.5700437128543854	accuracy: 79.06%	val_loss: 0.6828859313391149	val_accuracy: 75.02%
[15, 120]	loss: 0.7321209251880646	accuracy: 70.94%	val_loss: 0.7124384213238955	val_accuracy: 74.73%
[15, 130]	loss: 0.604701042175293	accuracy: 76.88%	val_loss: 0.667386676184833	val_accuracy: 75.62%
[15, 140]	loss: 0.5418462812900543	accuracy: 78.44%	val_loss: 0.7073777318000793	val_accuracy: 74.73%
[15, 150]	loss: 0.5294135987758637	accuracy: 80.00%	val_loss: 0.7401409512385726	val_accuracy: 74.33%
[15, 160]	loss: 0.626316973567009	accuracy: 78.44%	val_loss: 0.6848354106768966	val_accuracy: 75.52%
[15, 170]	loss: 0.6773416697978973	accuracy: 73.44%	val_loss: 0.6725482847541571	val_accuracy: 74.03%
[15, 180]	loss: 0.56961729824543	accuracy: 76.25%	val_loss: 0.6913834568113089	val_accuracy: 74.53%
[16, 10]	loss: 0.5964781373739243	accuracy: 79.69%	val_loss: 0.6999955726787448	val_accuracy: 73.23%
[16, 20]	loss: 0.6769651979207992	accuracy: 75.31%	val_loss: 0.7048727152869105	val_accuracy: 73.43%
[16, 30]	loss: 0.6941670686006546	accuracy: 74.38%	val_loss: 0.6801823717541993	val_accuracy: 74.13%
[16, 40]	loss: 0.5968503683805466	accuracy: 78.44%	val_loss: 0.7077294941991568	val_accuracy: 73.83%




[16, 50]	loss: 0.5328584462404251	accuracy: 78.12%	val_loss: 0.701907797716558	val_accuracy: 74.53%
[16, 60]	loss: 0.6565305113792419	accuracy: 73.44%	val_loss: 0.6783710746094584	val_accuracy: 75.02%
[16, 70]	loss: 0.6286000639200211	accuracy: 79.06%	val_loss: 0.7694631535559893	val_accuracy: 69.43%
[16, 80]	loss: 0.5319704949855805	accuracy: 81.25%	val_loss: 0.6691893581300974	val_accuracy: 75.62%
[16, 90]	loss: 0.5211558490991592	accuracy: 83.12%	val_loss: 0.9185375943779945	val_accuracy: 72.73%
[16, 100]	loss: 0.6595533668994904	accuracy: 75.00%	val_loss: 0.697311645373702	val_accuracy: 73.63%
[16, 110]	loss: 0.6359914600849151	accuracy: 77.81%	val_loss: 0.682224134914577	val_accuracy: 73.93%
[16, 120]	loss: 0.460530486702919	accuracy: 84.06%	val_loss: 0.7135714227333665	val_accuracy: 73.93%
[16, 130]	loss: 0.5834781587123871	accuracy: 77.19%	val_loss: 0.6908708298578858	val_accuracy: 73.93%
[16, 140]	loss: 0.49966336488723756	accuracy: 82.19%	val_loss: 0.7195969270542264	val_accuracy: 74.53%
[16, 150]	loss: 0.5075605839490891	accuracy: 79.69%	val_loss: 0.6616184609010816	val_accuracy: 75.52%
[16, 160]	loss: 0.5394823312759399	accuracy: 77.81%	val_loss: 0.6931201461702585	val_accuracy: 75.42%
[16, 170]	loss: 0.5480892419815063	accuracy: 81.25%	val_loss: 0.678401050157845	val_accuracy: 75.72%
[16, 180]	loss: 0.539176344871521	accuracy: 80.62%	val_loss: 0.6582878576591611	val_accuracy: 75.82%
[17, 10]	loss: 0.5663069456815719	accuracy: 79.06%	val_loss: 0.6348041575402021	val_accuracy: 76.22%
[17, 20]	loss: 0.5843781590461731	accuracy: 78.75%	val_loss: 0.9439464490860701	val_accuracy: 62.44%
[17, 30]	loss: 0.5426694095134735	accuracy: 80.00%	val_loss: 0.654281809926033	val_accuracy: 75.92%
[17, 40]	loss: 0.4895738273859024	accuracy: 84.06%	val_loss: 0.6694999532774091	val_accuracy: 75.52%
[17, 50]	loss: 0.6127139866352082	accuracy: 76.88%	val_loss: 0.6823668917641044	val_accuracy: 75.52%
[17, 60]	loss: 0.5314958095550537	accuracy: 79.38%	val_loss: 0.6546584209427238	val_accuracy: 75.92%
[17, 70]	loss: 0.4693871021270752	accuracy: 83.12%	val_loss: 0.778367917984724	val_accuracy: 74.63%
[17, 80]	loss: 0.4597057715058327	accuracy: 83.44%	val_loss: 0.6755098002031446	val_accuracy: 76.02%
[17, 90]	loss: 0.534672400355339	accuracy: 81.56%	val_loss: 0.6665982119739056	val_accuracy: 75.02%
[17, 100]	loss: 0.5699723839759827	accuracy: 77.81%	val_loss: 0.6565159317106009	val_accuracy: 76.12%
[17, 110]	loss: 0.6394007474184036	accuracy: 77.81%	val_loss: 0.6784810954704881	val_accuracy: 74.73%
[17, 120]	loss: 0.6771276950836181	accuracy: 75.31%	val_loss: 0.6635330151766539	val_accuracy: 75.82%
[17, 130]	loss: 0.5809164106845855	accuracy: 80.62%	val_loss: 0.6598292263224721	val_accuracy: 75.42%
[17, 140]	loss: 0.5501100093126297	accuracy: 77.50%	val_loss: 0.7266619727015495	val_accuracy: 74.53%
[17, 150]	loss: 0.5182750701904297	accuracy: 79.06%	val_loss: 0.6688973223790526	val_accuracy: 75.52%
[17, 160]	loss: 0.5048647135496139	accuracy: 81.56%	val_loss: 0.7044289372861385	val_accuracy: 72.93%
[17, 170]	loss: 0.6183887988328933	accuracy: 75.94%	val_loss: 0.7290096711367369	val_accuracy: 74.83%
[17, 180]	loss: 0.5638150960206986	accuracy: 80.00%	val_loss: 0.7700883261859417	val_accuracy: 70.13%
[18, 10]	loss: 0.4851915866136551	accuracy: 83.75%	val_loss: 0.6997059807181358	val_accuracy: 75.72%
[18, 20]	loss: 0.559866976737976	accuracy: 78.44%	val_loss: 0.6820231368765235	val_accuracy: 74.23%
[18, 30]	loss: 0.530940717458725	accuracy: 82.19%	val_loss: 0.6902415594086051	val_accuracy: 75.22%
[18, 40]	loss: 0.5704767525196075	accuracy: 82.19%	val_loss: 0.6783897113054991	val_accuracy: 75.72%
[18, 50]	loss: 0.5370629787445068	accuracy: 81.88%	val_loss: 0.6840983824804425	val_accuracy: 75.12%
[18, 60]	loss: 0.4762289136648178	accuracy: 81.88%	val_loss: 0.6804696870967746	val_accuracy: 74.33%
[18, 70]	loss: 0.506780418753624	accuracy: 80.00%	val_loss: 0.6630978919565678	val_accuracy: 75.62%
[18, 80]	loss: 0.5853408306837082	accuracy: 78.44%	val_loss: 0.685639831237495	val_accuracy: 74.43%
[18, 90]	loss: 0.569967982172966	accuracy: 79.38%	val_loss: 0.7412049937993288	val_accuracy: 73.03%
[18, 100]	loss: 0.423202720284462	accuracy: 84.06%	val_loss: 0.7597796665504575	val_accuracy: 74.83%

[18, 110]	loss: 0.557905125617981	accuracy: 79.38%	val_loss: 0.7167279990389943	val_accuracy: 74.73%
[18, 120]	loss: 0.4835855334997177	accuracy: 84.06%	val_loss: 0.6999374832957983	val_accuracy: 73.33%
[18, 130]	loss: 0.5112580478191375	accuracy: 81.88%	val_loss: 0.7543759401887655	val_accuracy: 72.33%
[18, 140]	loss: 0.5205389887094498	accuracy: 80.31%	val_loss: 0.7049632994458079	val_accuracy: 75.52%
[18, 150]	loss: 0.5457695126533508	accuracy: 77.81%	val_loss: 0.6737318746745586	val_accuracy: 74.93%
[18, 160]	loss: 0.5713148415088654	accuracy: 79.06%	val_loss: 0.6971875838935375	val_accuracy: 75.32%
[18, 170]	loss: 0.5451025307178498	accuracy: 80.00%	val_loss: 0.765893530100584	val_accuracy: 70.43%
[18, 180]	loss: 0.5250411629676819	accuracy: 80.00%	val_loss: 0.6648302497342229	val_accuracy: 74.63%
[19, 10]	loss: 0.5170228570699692	accuracy: 81.25%	val_loss: 0.703082007355988	val_accuracy: 75.32%
[19, 20]	loss: 0.4787963777780533	accuracy: 82.81%	val_loss: 0.7045997567474842	val_accuracy: 74.03%
[19, 30]	loss: 0.5419222474098205	accuracy: 81.25%	val_loss: 0.6788358530029655	val_accuracy: 76.22%
[19, 40]	loss: 0.4324310928583145	accuracy: 84.38%	val_loss: 0.7502948939800262	val_accuracy: 72.43%
[19, 50]	loss: 0.4886013180017471	accuracy: 81.56%	val_loss: 0.6995030250400305	val_accuracy: 75.22%
[19, 60]	loss: 0.531373479962349	accuracy: 78.75%	val_loss: 0.6821174351498485	val_accuracy: 75.62%
[19, 70]	loss: 0.5440706044435502	accuracy: 80.31%	val_loss: 0.6965955151244998	val_accuracy: 75.32%
[19, 80]	loss: 0.48395558893680574	accuracy: 83.44%	val_loss: 0.6783685479313135	val_accuracy: 74.53%
[19, 90]	loss: 0.4902141153812408	accuracy: 80.62%	val_loss: 0.6702923877164721	val_accuracy: 74.73%
[19, 100]	loss: 0.47545807957649233	accuracy: 83.75%	val_loss: 0.7165146072511561	val_accuracy: 74.43%
[19, 110]	loss: 0.46781986951828003	accuracy: 83.75%	val_loss: 0.6631583878770471	val_accuracy: 74.63%
[19, 120]	loss: 0.4769197329878807	accuracy: 83.12%	val_loss: 0.7348877303302288	val_accuracy: 72.23%
[19, 130]	loss: 0.5277956962585449	accuracy: 80.31%	val_loss: 0.7261274200864136	val_accuracy: 74.73%
[19, 140]	loss: 0.4980957359075546	accuracy: 82.81%	val_loss: 0.7745617171749473	val_accuracy: 75.22%
[19, 150]	loss: 0.5838368773460388	accuracy: 76.25%	val_loss: 0.7034825701266527	val_accuracy: 73.33%
[19, 160]	loss: 0.5631841003894806	accuracy: 81.56%	val_loss: 0.7905402733013034	val_accuracy: 75.12%
[19, 170]	loss: 0.5280616641044616	accuracy: 80.31%	val_loss: 0.7506275428459048	val_accuracy: 74.83%
[19, 180]	loss: 0.5619235873222351	accuracy: 79.06%	val_loss: 0.7063659923151135	val_accuracy: 74.83%
[20, 10]	loss: 0.46805990040302276	accuracy: 83.44%	val_loss: 0.7157301530241966	val_accuracy: 74.93%
[20, 20]	loss: 0.5928396433591843	accuracy: 77.50%	val_loss: 0.6807396290823817	val_accuracy: 75.02%
[20, 30]	loss: 0.46932328641414645	accuracy: 82.50%	val_loss: 0.7149901231750846	val_accuracy: 74.33%
[20, 40]	loss: 0.3826304703950882	accuracy: 85.62%	val_loss: 0.6916065057739615	val_accuracy: 74.53%
[20, 50]	loss: 0.43651013970375063	accuracy: 82.81%	val_loss: 0.747936071595177	val_accuracy: 75.02%
[20, 60]	loss: 0.465143084526062	accuracy: 84.38%	val_loss: 0.7074233908206224	val_accuracy: 75.72%
[20, 70]	loss: 0.47701441645622256	accuracy: 81.88%	val_loss: 0.8409836292266846	val_accuracy: 74.53%
[20, 80]	loss: 0.47115504145622256	accuracy: 81.25%	val_loss: 0.713282996788621	val_accuracy: 75.42%
[20, 90]	loss: 0.45367157012224196	accuracy: 84.06%	val_loss: 0.707571784965694	val_accuracy: 74.33%
[20, 100]	loss: 0.5789996743202209	accuracy: 75.94%	val_loss: 0.7057148739695549	val_accuracy: 74.93%
[20, 110]	loss: 0.5340734601020813	accuracy: 78.75%	val_loss: 0.7190236747264862	val_accuracy: 73.43%
[20, 120]	loss: 0.5389476895332337	accuracy: 79.38%	val_loss: 0.7094653332605958	val_accuracy: 75.52%
[20, 130]	loss: 0.48544315695762635	accuracy: 82.19%	val_loss: 0.7067440040409565	val_accuracy: 75.42%
[20, 140]	loss: 0.4435017019510269	accuracy: 85.00%	val_loss: 0.724676557816565	val_accuracy: 76.52%
[20, 150]	loss: 0.5068931370973587	accuracy: 81.25%	val_loss: 0.7021793685853481	val_accuracy: 75.32%
[20, 160]	loss: 0.4611221671104431	accuracy: 80.00%	val_loss: 0.7594579085707664	val_accuracy: 74.23%

[20, 170] loss: 0.5360630750656128 | accuracy: 79.38% | val\_loss: 0.8036945099011064 | val\_accuracy: 74.73%  
 [20, 180] loss: 0.4569655045866966 | accuracy: 80.31% | val\_loss: 0.7320058541372418 | val\_accuracy: 72.93%

```
In [33]: # Get the number_to_label_array so that we can use it later to plot the confusion matrix
simple_number_to_label_array = testloader.dataset.dataset.number_to_label_array
# Test the model
test_net(model, testloader, simple_number_to_label_array, loss=loss, device=device)
```

Test loss: 0.7584205269813538 | Test accuracy: 73.74%




## Observations

- We can see that the model has a strong inclination to predict the label `nv` and the label `mel`. This might be because the dataset is imbalanced, as we can see below

- The model only predicts `df` once, which again is the least common label in the dataset
- Even though `vasc` is the second least common label, the model does a very good job at predicting it, both precision and recall are very good
- A lot of the predictions of `mel` are actually `nv`, which might mean that either the model is having a hard time distinguishing between the two, or that there are a lot of similarities between the two illnesses
- Even though `bk1` is the second most common label (close with `mel`), the model isn't very good at distinguishing it from `mel` and `nv`.
- We can also see that for `bk1` that has about the same number of samples as `mel`, that when the model was given a `bk1` image, that the predictions are more spread out, whereas for `mel` the predictions are more concentrated on `nv` and `mel`. Again this might strongly suggest that there are a lot of similarities between `mel` and `nv`
- The model's performance on the rest of the labels `akiec`, `bcc`, are very good (recall)

```
In [10]: label_counts = dataset.df['dx'].value_counts()
label_counts = [dataset.number_to_label_array[i] for i in label_counts.index]

plt.figure(figsize=(10,6))
plt.bar(label_counts, dataset.df['dx'].value_counts())
plt.show()
```



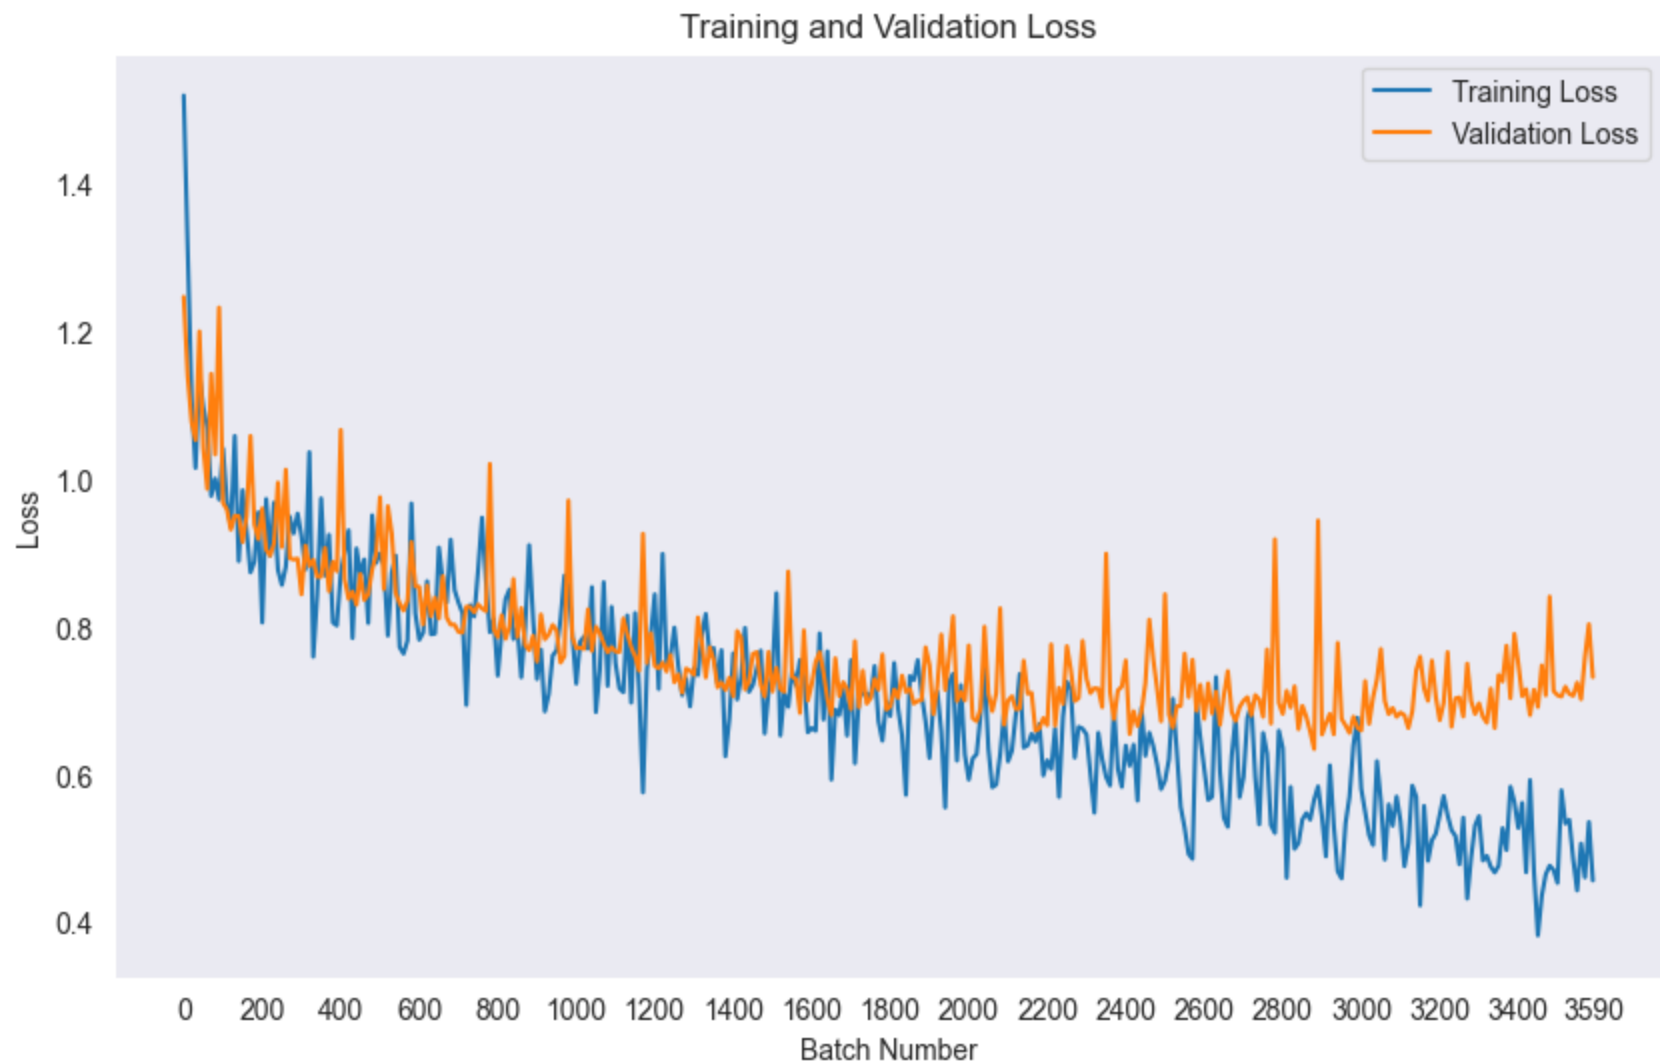
Now let's plot the loss and accuracy

```
In [11]: import copy  
         from torchsummary import summary
```

```
In [12]: # Copying the lists so that I don't overwrite them  
         simple_train_loss = copy.deepcopy(train_loss)  
         simple_val_loss = copy.deepcopy(val_loss)
```

```
simple_train_accuracy = copy.deepcopy(train_accuracy)
simple_val_accuracy = copy.deepcopy(val_accuracy)
```

```
In [13]: x_axis = list(map(lambda x:x*10,range(len(simple_train_loss))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,simple_train_loss, label='Training Loss')
plt.plot(x_axis , simple_val_loss,label='Validation Loss')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 200.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid()
plt.show()
```




## Observations

- From the loss plots I can conclude that the model might start to overfit at around 2500 batches, so we might want to stop the training at around that point
- The validation loss and the training loss are very close together, which is a good thing, because while the model is improving on the training set, it is also improving on the validation set.

- A very negative observation is that both losses are fluctuating a lot, which might mean that the model is "overshooting" the minimum loss, therefore requiring more epochs to train

```
In [14]: x_axis = list(map(lambda x:x*10,range(len(simple_train_accuracy))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,simple_train_accuracy, label='Training Accuracy')
plt.plot(x_axis , simple_val_accuracy,label='Validation Accuracy')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 200.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.grid()
plt.show()
```





## Observations

- From the accuracy plots we can still conclude that the model might start to overfit at around 2500 batches, which is good that both the loss and accuracy plots are showing the same thing
- The accuracy for both the training and validation set are increasing, which is a good thing, because the model is improving on both sets

- We can still see the big fluctuation we noticed on the loss plot before (negative).

## Complex CNN

```
In [15]: from torch.nn.functional import relu

class ComplexConvNN(nn.Module):
    def __init__(self):
        super(ComplexConvNN, self).__init__()
        # In the specification of the assignment, it says that
        # the images should be resized to [100 x 125] x 3 channels
        # the comments bellow are the output image size after each layer
        self.conv1 = nn.Conv2d(3, 32, 3) # 98 x 123 x 32
        self.batchnorm1 = nn.BatchNorm2d(32) # 98 x 123 x 32
        # MaxPool 2x2 49 x 61 x 32
        self.conv2 = nn.Conv2d(32, 64, 3) # 47 x 59 x 64
        self.batchnorm2 = nn.BatchNorm2d(64) # 47 x 59 x 64
        # MaxPool 2x2 23 x 29 x 64
        self.conv3 = nn.Conv2d(64, 128, 3) # 21 x 27 x 128
        self.batchnorm3 = nn.BatchNorm2d(128) # 21 x 27 x 128
        # MaxPool 2x2 10 x 13 x 128
        self.conv4 = nn.Conv2d(128, 256, 3) # 8 x 11 x 256
        self.batchnorm4 = nn.BatchNorm2d(256) # 8 x 11 x 256
        # MaxPool 2x2 4 x 5 x 256
        self.conv5 = nn.Conv2d(256, 512, 3) # 2 x 3 x 512
        self.batchnorm5 = nn.BatchNorm2d(512) # 2 x 3 x 512
        # MaxPool 2x2 1 x 1 x 512
        self.fc = nn.Linear(512, 7)
        self.maxpool = nn.MaxPool2d(2)
        self.flatten = nn.Flatten()

    def forward(self, x):
        # Ordering the layers as given in the assignment
        x = self.maxpool(self.batchnorm1(relu(self.conv1(x))))
        x = self.maxpool(self.batchnorm2(relu(self.conv2(x))))
        x = self.maxpool(self.batchnorm3(relu(self.conv3(x))))
        x = self.maxpool(self.batchnorm4(relu(self.conv4(x))))
        x = self.maxpool(self.batchnorm5(relu(self.conv5(x))))
```

```
x = nn.functional.adaptive_avg_pool2d(x,(1,1))
x = self.flatten(x)
x = self.fc(x)
return x
```

## Now let's train the Complex model and test it and see the difference

Note: naming everything complex\_ so that I don't overwrite the simple model's variables

```
In [16]: # Constants from the assignment
complex_M = 100
complex_N = 125
# Create the transforms
complex_transforms = create_transform(complex_M, complex_N)
# Create the dataset
complex_dataset = MLProject2Dataset(data_dir='data', metadata_fname='metadata.csv', transform=complex_transforms)
# Split the dataset into train, validation and test
complex_train_dataset, complex_val_dataset, complex_test_dataset = torch.utils.data.random_split(
    complex_dataset,
    [.6, .1, .3],
    generator=torch.Generator().manual_seed(42))

# Create the dataLoaders
BATCH_SIZE = 32
complex_trainloader = DataLoader(complex_train_dataset, batch_size=BATCH_SIZE, shuffle=True)
complex_valloader = DataLoader(complex_val_dataset, batch_size=BATCH_SIZE, shuffle=True)
complex_testloader = DataLoader(complex_test_dataset, batch_size=BATCH_SIZE, shuffle=True)

# Create the model
complex_model = ComplexConvNN()
complex_model = complex_model.to(device)
# Print the model summary
summary(complex_model, (3, complex_M, complex_N))
# Create the optimizer
complex_optimizer = optim.Adam(complex_model.parameters(), lr=1e-3)
# Create the loss function
complex_loss = nn.CrossEntropyLoss()
# Train the model
```

```
train_net(complex_model, complex_trainloader, complex_valloader, epochs=20, optimizer=complex_optimizer,  
          loss=complex_loss, device=device)  
# This cell took 44 minutes to run on my machine
```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 32, 98, 123]	896
BatchNorm2d-2	[-1, 32, 98, 123]	64
MaxPool2d-3	[-1, 32, 49, 61]	0
Conv2d-4	[-1, 64, 47, 59]	18,496
BatchNorm2d-5	[-1, 64, 47, 59]	128
MaxPool2d-6	[-1, 64, 23, 29]	0
Conv2d-7	[-1, 128, 21, 27]	73,856
BatchNorm2d-8	[-1, 128, 21, 27]	256
MaxPool2d-9	[-1, 128, 10, 13]	0
Conv2d-10	[-1, 256, 8, 11]	295,168
BatchNorm2d-11	[-1, 256, 8, 11]	512
MaxPool2d-12	[-1, 256, 4, 5]	0
Conv2d-13	[-1, 512, 2, 3]	1,180,160
BatchNorm2d-14	[-1, 512, 2, 3]	1,024
MaxPool2d-15	[-1, 512, 1, 1]	0
Flatten-16	[-1, 512]	0
Linear-17	[-1, 7]	3,591

=====

Total params: 1,574,151

Trainable params: 1,574,151

Non-trainable params: 0

-----

Input size (MB): 0.14

Forward/backward pass size (MB): 11.32

Params size (MB): 6.00

Estimated Total Size (MB): 17.47

-----

[1, 10]	loss: 1.525368857383728		accuracy: 54.69%		val_loss: 1.754875309765339		val_accuracy: 61.14%
[1, 20]	loss: 1.048032408952713		accuracy: 66.56%		val_loss: 1.4846876803785563		val_accuracy: 66.03%
[1, 30]	loss: 0.9679681122303009		accuracy: 64.69%		val_loss: 1.3263956252485514		val_accuracy: 64.34%
[1, 40]	loss: 0.8475734531879425		accuracy: 70.31%		val_loss: 1.196918735280633		val_accuracy: 65.43%
[1, 50]	loss: 0.7968601644039154		accuracy: 71.56%		val_loss: 1.0745130609720945		val_accuracy: 66.03%
[1, 60]	loss: 0.925772750377655		accuracy: 65.00%		val_loss: 0.9131527952849865		val_accuracy: 67.23%
[1, 70]	loss: 0.8625742673873902		accuracy: 68.44%		val_loss: 0.9154977351427078		val_accuracy: 68.03%
[1, 80]	loss: 0.9956126093864441		accuracy: 63.75%		val_loss: 0.8672284316271544		val_accuracy: 69.13%
[1, 90]	loss: 0.8503584682941436		accuracy: 67.50%		val_loss: 0.9038623012602329		val_accuracy: 67.73%
[1, 100]	loss: 0.8787101924419403		accuracy: 68.12%		val_loss: 0.8230524556711316		val_accuracy: 68.83%
[1, 110]	loss: 0.848471587896347		accuracy: 69.06%		val_loss: 0.8648187052458525		val_accuracy: 64.24%
[1, 120]	loss: 0.846266633272171		accuracy: 69.69%		val_loss: 0.9255834827199578		val_accuracy: 71.03%

[1, 130]	loss: 0.9086926519870758	accuracy: 70.00%	val_loss: 0.9583376441150904	val_accuracy: 61.64%
[1, 140]	loss: 0.934523731470108	accuracy: 67.19%	val_loss: 0.9434949308633804	val_accuracy: 69.83%
[1, 150]	loss: 0.9806777119636536	accuracy: 64.69%	val_loss: 0.8877771589905024	val_accuracy: 66.63%
[1, 160]	loss: 0.9341613352298737	accuracy: 66.25%	val_loss: 0.8553004846908152	val_accuracy: 68.13%
[1, 170]	loss: 0.7616441905498504	accuracy: 74.06%	val_loss: 0.810623412951827	val_accuracy: 70.63%
[1, 180]	loss: 0.8601815044879914	accuracy: 69.69%	val_loss: 0.7609581723809242	val_accuracy: 70.13%
[2, 10]	loss: 0.7945726901292801	accuracy: 67.81%	val_loss: 0.7938999701291323	val_accuracy: 71.23%
[2, 20]	loss: 0.7408511936664581	accuracy: 74.69%	val_loss: 0.8276758911088109	val_accuracy: 69.23%
[2, 30]	loss: 0.7311596751213074	accuracy: 71.56%	val_loss: 0.7480264557525516	val_accuracy: 72.33%
[2, 40]	loss: 0.7770373463630676	accuracy: 70.00%	val_loss: 0.8346471888944507	val_accuracy: 67.83%
[2, 50]	loss: 0.7937785059213638	accuracy: 72.81%	val_loss: 0.8024060428142548	val_accuracy: 71.83%
[2, 60]	loss: 0.7525808155536652	accuracy: 71.25%	val_loss: 0.8016196843236685	val_accuracy: 71.23%
[2, 70]	loss: 0.8220077574253082	accuracy: 68.44%	val_loss: 0.7701370064169168	val_accuracy: 73.13%
[2, 80]	loss: 0.8187563627958298	accuracy: 70.31%	val_loss: 0.8546833172440529	val_accuracy: 70.13%
[2, 90]	loss: 0.8573897421360016	accuracy: 67.81%	val_loss: 0.7251671906560659	val_accuracy: 71.03%
[2, 100]	loss: 0.7590774774551392	accuracy: 71.25%	val_loss: 0.736796161159873	val_accuracy: 72.13%
[2, 110]	loss: 0.7323598563671112	accuracy: 72.81%	val_loss: 0.7469757916405797	val_accuracy: 70.83%
[2, 120]	loss: 0.7378456771373749	accuracy: 74.38%	val_loss: 0.7975301695987582	val_accuracy: 72.83%
[2, 130]	loss: 0.7270847409963608	accuracy: 73.12%	val_loss: 0.7631562286987901	val_accuracy: 72.33%
[2, 140]	loss: 0.8094313263893127	accuracy: 71.56%	val_loss: 0.7409283518791199	val_accuracy: 72.33%
[2, 150]	loss: 0.7181434273719788	accuracy: 73.12%	val_loss: 0.7025347780436277	val_accuracy: 72.73%
[2, 160]	loss: 0.7273870825767517	accuracy: 73.12%	val_loss: 0.7537658223882318	val_accuracy: 69.23%
[2, 170]	loss: 0.7126692235469818	accuracy: 74.06%	val_loss: 0.8142985589802265	val_accuracy: 70.23%
[2, 180]	loss: 0.6493846714496613	accuracy: 76.25%	val_loss: 0.7875482980161905	val_accuracy: 72.03%
[3, 10]	loss: 0.7981949210166931	accuracy: 72.50%	val_loss: 0.7243238845840096	val_accuracy: 73.33%
[3, 20]	loss: 0.7865974605083466	accuracy: 70.31%	val_loss: 0.7377853617072105	val_accuracy: 72.33%
[3, 30]	loss: 0.7596865117549896	accuracy: 70.31%	val_loss: 0.7134949490427971	val_accuracy: 73.13%
[3, 40]	loss: 0.6869500339031219	accuracy: 75.94%	val_loss: 0.7359898155555129	val_accuracy: 72.63%
[3, 50]	loss: 0.6931857079267502	accuracy: 74.06%	val_loss: 0.6665789466351271	val_accuracy: 75.82%
[3, 60]	loss: 0.6467795431613922	accuracy: 74.38%	val_loss: 0.6891673523932695	val_accuracy: 73.73%
[3, 70]	loss: 0.7569205075502395	accuracy: 69.69%	val_loss: 0.7441888442263007	val_accuracy: 71.93%
[3, 80]	loss: 0.5871208906173706	accuracy: 79.06%	val_loss: 0.755314220674336	val_accuracy: 72.63%
[3, 90]	loss: 0.6852631628513336	accuracy: 73.75%	val_loss: 0.6993652936071157	val_accuracy: 72.93%
[3, 100]	loss: 0.6780262380838394	accuracy: 73.44%	val_loss: 0.718998109921813	val_accuracy: 73.63%
[3, 110]	loss: 0.6798575162887573	accuracy: 73.12%	val_loss: 0.6961178304627538	val_accuracy: 74.33%
[3, 120]	loss: 0.640699279308319	accuracy: 75.62%	val_loss: 0.8220480047166348	val_accuracy: 71.23%
[3, 130]	loss: 0.660409563779831	accuracy: 75.94%	val_loss: 0.764733461663127	val_accuracy: 71.83%
[3, 140]	loss: 0.7896904170513153	accuracy: 70.31%	val_loss: 0.7851889301091433	val_accuracy: 72.13%
[3, 150]	loss: 0.7337817490100861	accuracy: 72.50%	val_loss: 0.7321496540680528	val_accuracy: 73.33%
[3, 160]	loss: 0.762595921754837	accuracy: 73.44%	val_loss: 0.7484336290508509	val_accuracy: 71.33%
[3, 170]	loss: 0.6975354492664337	accuracy: 76.56%	val_loss: 0.818058161996305	val_accuracy: 70.03%
[3, 180]	loss: 0.7101779103279113	accuracy: 74.06%	val_loss: 0.7126108082011342	val_accuracy: 72.53%

[4, 10]	loss: 0.6851392537355423	accuracy: 73.12%	val_loss: 0.6661160225048661	val_accuracy: 74.53%
[4, 20]	loss: 0.6143472164869308	accuracy: 77.50%	val_loss: 0.7308760536834598	val_accuracy: 71.33%
[4, 30]	loss: 0.6243782639503479	accuracy: 77.19%	val_loss: 0.7943532140925527	val_accuracy: 72.93%
[4, 40]	loss: 0.6495967596769333	accuracy: 75.31%	val_loss: 0.690502280369401	val_accuracy: 72.83%
[4, 50]	loss: 0.6017820715904236	accuracy: 77.19%	val_loss: 0.7567709516733885	val_accuracy: 72.63%
[4, 60]	loss: 0.6245465099811554	accuracy: 78.44%	val_loss: 0.6946321725845337	val_accuracy: 73.73%
[4, 70]	loss: 0.5639015108346939	accuracy: 79.69%	val_loss: 0.7111127162352204	val_accuracy: 73.53%
[4, 80]	loss: 0.6008388489484787	accuracy: 79.69%	val_loss: 0.7157345283776522	val_accuracy: 73.83%
[4, 90]	loss: 0.6503967195749283	accuracy: 74.69%	val_loss: 0.7042385125532746	val_accuracy: 74.03%
[4, 100]	loss: 0.7284781575202942	accuracy: 72.50%	val_loss: 0.7170220529660583	val_accuracy: 73.03%
[4, 110]	loss: 0.6417625278234482	accuracy: 75.62%	val_loss: 0.7443004930391908	val_accuracy: 72.23%
[4, 120]	loss: 0.7383839577436447	accuracy: 73.44%	val_loss: 0.730309440754354	val_accuracy: 70.93%
[4, 130]	loss: 0.6641679525375366	accuracy: 73.44%	val_loss: 0.661432571709156	val_accuracy: 75.02%
[4, 140]	loss: 0.5629128009080887	accuracy: 81.25%	val_loss: 0.6822343049570918	val_accuracy: 72.93%
[4, 150]	loss: 0.5912684619426727	accuracy: 78.44%	val_loss: 0.6753487708047032	val_accuracy: 74.03%
[4, 160]	loss: 0.6912270188331604	accuracy: 76.25%	val_loss: 0.7066282192245126	val_accuracy: 74.03%
[4, 170]	loss: 0.6448119044303894	accuracy: 78.44%	val_loss: 0.7190633974969387	val_accuracy: 74.23%
[4, 180]	loss: 0.6659507751464844	accuracy: 75.62%	val_loss: 0.7199621144682169	val_accuracy: 72.63%
[5, 10]	loss: 0.42747653275728226	accuracy: 84.38%	val_loss: 0.7026381734758615	val_accuracy: 74.53%
[5, 20]	loss: 0.628812262415886	accuracy: 77.19%	val_loss: 0.6493992051109672	val_accuracy: 75.52%
[5, 30]	loss: 0.6150345742702484	accuracy: 74.69%	val_loss: 0.752644918859005	val_accuracy: 73.63%
[5, 40]	loss: 0.5681629598140716	accuracy: 80.94%	val_loss: 0.6697120564058423	val_accuracy: 74.93%
[5, 50]	loss: 0.6551063001155853	accuracy: 72.81%	val_loss: 0.6395934196189046	val_accuracy: 76.22%
[5, 60]	loss: 0.5988477528095245	accuracy: 77.81%	val_loss: 0.760942105203867	val_accuracy: 73.53%
[5, 70]	loss: 0.6526502728462219	accuracy: 75.00%	val_loss: 0.7256879229098558	val_accuracy: 73.03%
[5, 80]	loss: 0.5281339973211289	accuracy: 80.00%	val_loss: 0.7300684247165918	val_accuracy: 73.53%
[5, 90]	loss: 0.5579124093055725	accuracy: 80.31%	val_loss: 0.6426172945648432	val_accuracy: 76.82%
[5, 100]	loss: 0.6096403241157532	accuracy: 79.38%	val_loss: 0.7014140486717224	val_accuracy: 73.83%
[5, 110]	loss: 0.6089135110378265	accuracy: 77.19%	val_loss: 0.6939684972167015	val_accuracy: 72.93%
[5, 120]	loss: 0.5378909409046173	accuracy: 79.06%	val_loss: 0.6781051466241479	val_accuracy: 74.33%
[5, 130]	loss: 0.6228233128786087	accuracy: 77.19%	val_loss: 0.6672952435910702	val_accuracy: 74.93%
[5, 140]	loss: 0.6258638113737106	accuracy: 77.50%	val_loss: 0.701608058065176	val_accuracy: 74.33%
[5, 150]	loss: 0.5597929537296296	accuracy: 80.00%	val_loss: 0.742314176633954	val_accuracy: 73.43%
[5, 160]	loss: 0.602263480424881	accuracy: 78.12%	val_loss: 0.6921169459819794	val_accuracy: 74.43%
[5, 170]	loss: 0.6528618723154068	accuracy: 76.88%	val_loss: 0.7696810476481915	val_accuracy: 71.83%
[5, 180]	loss: 0.6011953115463257	accuracy: 77.19%	val_loss: 0.7547702193260193	val_accuracy: 74.23%
[6, 10]	loss: 0.5242511987686157	accuracy: 82.19%	val_loss: 0.7555083157494664	val_accuracy: 73.13%
[6, 20]	loss: 0.531225460767746	accuracy: 80.62%	val_loss: 0.7659792574122548	val_accuracy: 74.03%
[6, 30]	loss: 0.5498735547065735	accuracy: 80.31%	val_loss: 0.7253838591277599	val_accuracy: 74.03%
[6, 40]	loss: 0.5838291585445404	accuracy: 76.56%	val_loss: 0.7196865091100335	val_accuracy: 74.63%
[6, 50]	loss: 0.5515437334775924	accuracy: 81.56%	val_loss: 0.6670173043385148	val_accuracy: 74.73%
[6, 60]	loss: 0.5277738034725189	accuracy: 80.94%	val_loss: 0.6623294213786721	val_accuracy: 75.32%

[6, 70]	loss: 0.5436754882335663	accuracy: 79.69%	val_loss: 0.6594334347173572	val_accuracy: 75.92%
[6, 80]	loss: 0.5011586725711823	accuracy: 83.44%	val_loss: 0.6716408152133226	val_accuracy: 74.83%
[6, 90]	loss: 0.44178597033023836	accuracy: 83.44%	val_loss: 0.7686292761936784	val_accuracy: 73.43%
[6, 100]	loss: 0.572308573126793	accuracy: 77.81%	val_loss: 0.6937563940882683	val_accuracy: 75.82%
[6, 110]	loss: 0.5898539304733277	accuracy: 78.75%	val_loss: 0.762327391654253	val_accuracy: 71.43%
[6, 120]	loss: 0.6022080183029175	accuracy: 76.25%	val_loss: 0.7577684153802693	val_accuracy: 74.23%
[6, 130]	loss: 0.6230683565139771	accuracy: 77.81%	val_loss: 0.7852026764303446	val_accuracy: 70.93%
[6, 140]	loss: 0.6845522135496139	accuracy: 74.38%	val_loss: 0.7362423706799746	val_accuracy: 73.63%
[6, 150]	loss: 0.5716134905815125	accuracy: 80.94%	val_loss: 0.7635836079716682	val_accuracy: 74.93%
[6, 160]	loss: 0.6003013968467712	accuracy: 79.06%	val_loss: 0.6981897838413715	val_accuracy: 74.93%
[6, 170]	loss: 0.618625384569168	accuracy: 78.12%	val_loss: 0.7180526209995151	val_accuracy: 73.23%
[6, 180]	loss: 0.5887133687734604	accuracy: 77.81%	val_loss: 0.70682492852211	val_accuracy: 73.93%
[7, 10]	loss: 0.5634248942136765	accuracy: 76.56%	val_loss: 0.6966182123869658	val_accuracy: 74.83%
[7, 20]	loss: 0.5357971429824829	accuracy: 78.75%	val_loss: 0.7289251117035747	val_accuracy: 75.22%
[7, 30]	loss: 0.4663150429725647	accuracy: 84.38%	val_loss: 0.6549995937384665	val_accuracy: 74.03%
[7, 40]	loss: 0.5001370221376419	accuracy: 79.38%	val_loss: 0.6826645396649837	val_accuracy: 76.42%
[7, 50]	loss: 0.45196916460990905	accuracy: 83.75%	val_loss: 0.6773256994783878	val_accuracy: 76.52%
[7, 60]	loss: 0.5751553356647492	accuracy: 79.06%	val_loss: 0.6490303492173553	val_accuracy: 75.12%
[7, 70]	loss: 0.4658581018447876	accuracy: 82.50%	val_loss: 0.6900851363316178	val_accuracy: 76.12%
[7, 80]	loss: 0.5201054662466049	accuracy: 80.31%	val_loss: 0.8181768758222461	val_accuracy: 69.73%
[7, 90]	loss: 0.49066494405269623	accuracy: 81.25%	val_loss: 0.8877650555223227	val_accuracy: 71.93%
[7, 100]	loss: 0.4751939564943314	accuracy: 83.12%	val_loss: 0.7934413435868919	val_accuracy: 73.03%
[7, 110]	loss: 0.47964825928211213	accuracy: 80.94%	val_loss: 0.7247113166376948	val_accuracy: 74.53%
[7, 120]	loss: 0.5779032170772552	accuracy: 79.06%	val_loss: 0.7125228801742196	val_accuracy: 74.03%
[7, 130]	loss: 0.5329542577266693	accuracy: 80.94%	val_loss: 0.756137597374618	val_accuracy: 73.53%
[7, 140]	loss: 0.5919827550649643	accuracy: 78.44%	val_loss: 0.669350235722959	val_accuracy: 75.32%
[7, 150]	loss: 0.5527170091867447	accuracy: 78.44%	val_loss: 0.6826230697333813	val_accuracy: 75.72%
[7, 160]	loss: 0.5015882283449173	accuracy: 78.44%	val_loss: 0.7122407080605626	val_accuracy: 72.93%
[7, 170]	loss: 0.4937757521867752	accuracy: 82.19%	val_loss: 0.7171920370310545	val_accuracy: 74.83%
[7, 180]	loss: 0.5859994709491729	accuracy: 77.19%	val_loss: 0.7165569718927145	val_accuracy: 73.03%
[8, 10]	loss: 0.4640814781188965	accuracy: 83.75%	val_loss: 0.6844212589785457	val_accuracy: 75.82%
[8, 20]	loss: 0.4802232250571251	accuracy: 84.06%	val_loss: 0.6653440967202187	val_accuracy: 76.02%
[8, 30]	loss: 0.346993026137352	accuracy: 88.44%	val_loss: 0.7225746782496572	val_accuracy: 74.53%
[8, 40]	loss: 0.42597312331199644	accuracy: 85.00%	val_loss: 0.7238705363124609	val_accuracy: 73.43%
[8, 50]	loss: 0.46203380823135376	accuracy: 84.38%	val_loss: 0.7357659935951233	val_accuracy: 75.72%
[8, 60]	loss: 0.4289455056190491	accuracy: 83.75%	val_loss: 0.7429496459662914	val_accuracy: 76.32%
[8, 70]	loss: 0.4391680970788002	accuracy: 85.00%	val_loss: 0.703657578676939	val_accuracy: 74.93%
[8, 80]	loss: 0.4891479820013046	accuracy: 80.62%	val_loss: 0.7675921767950058	val_accuracy: 72.53%
[8, 90]	loss: 0.42904698848724365	accuracy: 80.94%	val_loss: 0.7237844932824373	val_accuracy: 75.72%
[8, 100]	loss: 0.49800845682621003	accuracy: 81.25%	val_loss: 0.7427102653309703	val_accuracy: 75.62%
[8, 110]	loss: 0.4144706666469574	accuracy: 83.12%	val_loss: 0.6443719221279025	val_accuracy: 76.42%
[8, 120]	loss: 0.44645904898643496	accuracy: 84.06%	val_loss: 0.7195924250409007	val_accuracy: 76.32%



[8, 130] loss: 0.41813180297613145 | accuracy: 83.12% | val\_loss: 0.6679299352690578 | val\_accuracy: 76.62%  
[8, 140] loss: 0.4510944664478302 | accuracy: 83.44% | val\_loss: 0.7720366399735212 | val\_accuracy: 74.63%  
[8, 150] loss: 0.43860664069652555 | accuracy: 83.12% | val\_loss: 0.7692750254645944 | val\_accuracy: 74.03%  
[8, 160] loss: 0.460187965631485 | accuracy: 84.38% | val\_loss: 0.8037289464846253 | val\_accuracy: 72.43%  
[8, 170] loss: 0.5126987636089325 | accuracy: 80.00% | val\_loss: 0.7359307929873466 | val\_accuracy: 74.33%  
[8, 180] loss: 0.4864861384034157 | accuracy: 80.62% | val\_loss: 0.7266455860808492 | val\_accuracy: 75.02%  
[9, 10] loss: 0.3166169703006744 | accuracy: 86.56% | val\_loss: 0.7393769845366478 | val\_accuracy: 75.42%  
[9, 20] loss: 0.384126278758049 | accuracy: 88.44% | val\_loss: 0.6952272485941648 | val\_accuracy: 75.12%  
[9, 30] loss: 0.34936831295490267 | accuracy: 84.38% | val\_loss: 0.7418244276195765 | val\_accuracy: 74.03%  
[9, 40] loss: 0.3787523075938225 | accuracy: 86.56% | val\_loss: 0.7207141928374767 | val\_accuracy: 74.73%  
[9, 50] loss: 0.3377505086362362 | accuracy: 85.31% | val\_loss: 0.7638653377071023 | val\_accuracy: 75.12%  
[9, 60] loss: 0.39897020906209946 | accuracy: 85.94% | val\_loss: 0.7349828770384192 | val\_accuracy: 75.42%  
[9, 70] loss: 0.4033952057361603 | accuracy: 84.38% | val\_loss: 0.7435270177666098 | val\_accuracy: 75.02%  
[9, 80] loss: 0.3554788395762444 | accuracy: 86.56% | val\_loss: 0.7502064639702439 | val\_accuracy: 74.33%  
[9, 90] loss: 0.32988495901227 | accuracy: 88.44% | val\_loss: 0.7130730152130127 | val\_accuracy: 76.92%  
[9, 100] loss: 0.33826935589313506 | accuracy: 86.56% | val\_loss: 0.6990932701155543 | val\_accuracy: 74.73%  
[9, 110] loss: 0.3701410606503487 | accuracy: 86.25% | val\_loss: 0.7560901842080057 | val\_accuracy: 75.82%  
[9, 120] loss: 0.3612973004579544 | accuracy: 86.25% | val\_loss: 0.7474024202674627 | val\_accuracy: 71.53%  
[9, 130] loss: 0.425904244184494 | accuracy: 85.00% | val\_loss: 0.8258399926126003 | val\_accuracy: 72.33%  
[9, 140] loss: 0.41219498217105865 | accuracy: 85.62% | val\_loss: 0.7315480476245284 | val\_accuracy: 75.12%  
[9, 150] loss: 0.46191360503435136 | accuracy: 81.88% | val\_loss: 0.8383160904049873 | val\_accuracy: 75.62%  
[9, 160] loss: 0.4874177575111389 | accuracy: 85.31% | val\_loss: 0.7688255310058594 | val\_accuracy: 75.12%  
[9, 170] loss: 0.48077191412448883 | accuracy: 80.94% | val\_loss: 0.7402738416567445 | val\_accuracy: 73.93%  
[9, 180] loss: 0.4294442892074585 | accuracy: 82.50% | val\_loss: 0.8797552864998579 | val\_accuracy: 72.53%  
[10, 10] loss: 0.32960854172706605 | accuracy: 89.38% | val\_loss: 0.7019726652652025 | val\_accuracy: 76.42%  
[10, 20] loss: 0.230664724111557 | accuracy: 93.12% | val\_loss: 0.6882468909025192 | val\_accuracy: 76.92%  
[10, 30] loss: 0.2841650083661079 | accuracy: 89.38% | val\_loss: 0.676266735419631 | val\_accuracy: 75.52%  
[10, 40] loss: 0.2350860580801964 | accuracy: 93.12% | val\_loss: 0.7379265613853931 | val\_accuracy: 75.32%  
[10, 50] loss: 0.2489873692393303 | accuracy: 90.94% | val\_loss: 0.6906312704086304 | val\_accuracy: 77.02%  
[10, 60] loss: 0.2746861234307289 | accuracy: 88.75% | val\_loss: 0.7260429067537189 | val\_accuracy: 76.12%  
[10, 70] loss: 0.25803551375865935 | accuracy: 90.00% | val\_loss: 0.7900746334344149 | val\_accuracy: 75.72%  
[10, 80] loss: 0.3102178543806076 | accuracy: 87.81% | val\_loss: 0.7321698535233736 | val\_accuracy: 73.93%  
[10, 90] loss: 0.4026780426502228 | accuracy: 85.00% | val\_loss: 0.8837671685032547 | val\_accuracy: 73.03%  
[10, 100] loss: 0.42276095151901244 | accuracy: 86.56% | val\_loss: 0.8278008848428726 | val\_accuracy: 72.23%  
[10, 110] loss: 0.322435200214386 | accuracy: 86.88% | val\_loss: 0.8874377887696028 | val\_accuracy: 72.73%  
[10, 120] loss: 0.28960962146520614 | accuracy: 88.75% | val\_loss: 0.7740703811869025 | val\_accuracy: 75.92%  
[10, 130] loss: 0.2674565531313419 | accuracy: 90.62% | val\_loss: 0.7284584818407893 | val\_accuracy: 76.32%  
[10, 140] loss: 0.27849512100219725 | accuracy: 90.94% | val\_loss: 0.7557185404002666 | val\_accuracy: 77.32%  
[10, 150] loss: 0.2491996854543686 | accuracy: 90.31% | val\_loss: 0.7652645688503981 | val\_accuracy: 75.82%  
[10, 160] loss: 0.28879406601190566 | accuracy: 90.62% | val\_loss: 0.7480279952287674 | val\_accuracy: 77.32%  
[10, 170] loss: 0.304297436773777 | accuracy: 88.75% | val\_loss: 0.729223994538188 | val\_accuracy: 76.72%  
[10, 180] loss: 0.4278077632188797 | accuracy: 83.75% | val\_loss: 0.8021510299295187 | val\_accuracy: 73.33%

[11, 10]	loss: 0.20485142916440963	accuracy: 92.19%	val_loss: 0.8090666546486318	val_accuracy: 73.43%
[11, 20]	loss: 0.22460691630840302	accuracy: 90.94%	val_loss: 0.8081153770908713	val_accuracy: 74.63%
[11, 30]	loss: 0.18803773447871208	accuracy: 93.75%	val_loss: 0.7663395954295993	val_accuracy: 77.12%
[11, 40]	loss: 0.18693647906184196	accuracy: 93.44%	val_loss: 0.8091229954734445	val_accuracy: 75.32%
[11, 50]	loss: 0.3013253159821033	accuracy: 87.19%	val_loss: 1.0432086903601885	val_accuracy: 71.63%
[11, 60]	loss: 0.2305863991379738	accuracy: 92.19%	val_loss: 0.7753614038228989	val_accuracy: 76.12%
[11, 70]	loss: 0.4126579537987709	accuracy: 86.25%	val_loss: 0.8146556178107858	val_accuracy: 75.92%
[11, 80]	loss: 0.303260849416256	accuracy: 88.75%	val_loss: 0.8859371379949152	val_accuracy: 73.23%
[11, 90]	loss: 0.285899792611599	accuracy: 88.12%	val_loss: 0.8209557188674808	val_accuracy: 74.13%
[11, 100]	loss: 0.32080396115779874	accuracy: 86.88%	val_loss: 0.80959093850106	val_accuracy: 74.43%
[11, 110]	loss: 0.3036200061440468	accuracy: 88.75%	val_loss: 0.7645796267315745	val_accuracy: 76.42%
[11, 120]	loss: 0.3117756336927414	accuracy: 88.75%	val_loss: 0.7444771328009665	val_accuracy: 75.92%
[11, 130]	loss: 0.3985989898443222	accuracy: 85.94%	val_loss: 0.7414070246741176	val_accuracy: 74.93%
[11, 140]	loss: 0.35424110740423204	accuracy: 84.69%	val_loss: 0.7464679926633835	val_accuracy: 75.92%
[11, 150]	loss: 0.2665899500250816	accuracy: 90.62%	val_loss: 0.8216437450610101	val_accuracy: 75.22%
[11, 160]	loss: 0.2902370892465115	accuracy: 91.25%	val_loss: 0.789897358044982	val_accuracy: 74.53%
[11, 170]	loss: 0.23488557934761048	accuracy: 91.25%	val_loss: 0.8302321615628898	val_accuracy: 74.73%
[11, 180]	loss: 0.25668537318706514	accuracy: 92.19%	val_loss: 0.898239653557539	val_accuracy: 73.43%
[12, 10]	loss: 0.17189072966575622	accuracy: 94.06%	val_loss: 0.7299645398743451	val_accuracy: 76.92%
[12, 20]	loss: 0.18843016475439073	accuracy: 93.44%	val_loss: 0.8331656884402037	val_accuracy: 75.12%
[12, 30]	loss: 0.18653838634490966	accuracy: 91.88%	val_loss: 0.7995128836482763	val_accuracy: 75.52%
[12, 40]	loss: 0.14363089874386786	accuracy: 96.25%	val_loss: 0.7840799847617745	val_accuracy: 76.02%
[12, 50]	loss: 0.17881478667259215	accuracy: 93.12%	val_loss: 0.8264884324744344	val_accuracy: 77.02%
[12, 60]	loss: 0.18179772570729255	accuracy: 94.06%	val_loss: 0.8232440818101168	val_accuracy: 75.42%
[12, 70]	loss: 0.16494905091822148	accuracy: 95.31%	val_loss: 0.8706642417237163	val_accuracy: 75.12%
[12, 80]	loss: 0.17275362685322762	accuracy: 94.38%	val_loss: 0.8301319209858775	val_accuracy: 74.53%
[12, 90]	loss: 0.1899953603744507	accuracy: 93.75%	val_loss: 0.8244237517938018	val_accuracy: 74.53%
[12, 100]	loss: 0.17877879813313485	accuracy: 93.12%	val_loss: 0.8401510482653975	val_accuracy: 75.92%
[12, 110]	loss: 0.19057614281773566	accuracy: 93.44%	val_loss: 0.8433142928406596	val_accuracy: 77.42%
[12, 120]	loss: 0.12587017845362425	accuracy: 95.62%	val_loss: 0.8734995201230049	val_accuracy: 75.42%
[12, 130]	loss: 0.2095507375895977	accuracy: 92.81%	val_loss: 0.8544324692338705	val_accuracy: 77.12%
[12, 140]	loss: 0.19610246866941453	accuracy: 92.19%	val_loss: 0.8447294514626265	val_accuracy: 74.53%
[12, 150]	loss: 0.1994698651134968	accuracy: 93.12%	val_loss: 0.8893896332010627	val_accuracy: 75.82%
[12, 160]	loss: 0.2374863713979721	accuracy: 92.50%	val_loss: 0.9530644482001662	val_accuracy: 74.53%
[12, 170]	loss: 0.18425684422254562	accuracy: 93.12%	val_loss: 0.8772980635985732	val_accuracy: 76.42%
[12, 180]	loss: 0.22585366144776345	accuracy: 91.88%	val_loss: 0.8799355593509972	val_accuracy: 76.62%
[13, 10]	loss: 0.11684320718050004	accuracy: 94.69%	val_loss: 0.9171552592888474	val_accuracy: 75.42%
[13, 20]	loss: 0.09838083013892174	accuracy: 97.19%	val_loss: 0.8840879998169839	val_accuracy: 75.42%
[13, 30]	loss: 0.1121705312281847	accuracy: 95.94%	val_loss: 0.9080717815086246	val_accuracy: 76.02%
[13, 40]	loss: 0.07448841854929925	accuracy: 96.88%	val_loss: 0.8930431200424209	val_accuracy: 76.22%
[13, 50]	loss: 0.08496476169675589	accuracy: 97.19%	val_loss: 0.8916977224871516	val_accuracy: 75.52%
[13, 60]	loss: 0.10428302772343159	accuracy: 95.62%	val_loss: 0.9002576246857643	val_accuracy: 76.82%

[13, 70]	loss: 0.09401681665331126	accuracy: 97.19%	val_loss: 0.930202744435519	val_accuracy: 76.92%
[13, 80]	loss: 0.11778297163546085	accuracy: 95.94%	val_loss: 0.8923173239454627	val_accuracy: 75.02%
[13, 90]	loss: 0.17582507357001304	accuracy: 92.50%	val_loss: 0.9215852450579405	val_accuracy: 76.22%
[13, 100]	loss: 0.09754376448690891	accuracy: 96.56%	val_loss: 0.9387229764834046	val_accuracy: 76.22%
[13, 110]	loss: 0.14602591842412949	accuracy: 95.94%	val_loss: 0.9933974640443921	val_accuracy: 76.02%
[13, 120]	loss: 0.2238267980515957	accuracy: 92.19%	val_loss: 0.971290459856391	val_accuracy: 74.03%
[13, 130]	loss: 0.22838102877140046	accuracy: 89.69%	val_loss: 0.9918082654476166	val_accuracy: 76.42%
[13, 140]	loss: 0.15112231075763702	accuracy: 94.38%	val_loss: 0.8657059650868177	val_accuracy: 74.43%
[13, 150]	loss: 0.14471314158290624	accuracy: 94.38%	val_loss: 0.8998843785375357	val_accuracy: 76.12%
[13, 160]	loss: 0.17268291488289833	accuracy: 92.81%	val_loss: 1.045113236643374	val_accuracy: 75.42%
[13, 170]	loss: 0.1448018416762352	accuracy: 95.00%	val_loss: 0.9915933618322015	val_accuracy: 74.23%
[13, 180]	loss: 0.15660530161112546	accuracy: 95.00%	val_loss: 0.9095904808491468	val_accuracy: 75.22%
[14, 10]	loss: 0.1092452647164464	accuracy: 96.88%	val_loss: 0.9795179031789303	val_accuracy: 75.42%
[14, 20]	loss: 0.11364899054169655	accuracy: 96.56%	val_loss: 0.879524442832917	val_accuracy: 75.22%
[14, 30]	loss: 0.09839461799710988	accuracy: 96.25%	val_loss: 0.86506960215047	val_accuracy: 76.22%
[14, 40]	loss: 0.08732141219079495	accuracy: 96.88%	val_loss: 0.8699559001252055	val_accuracy: 77.02%
[14, 50]	loss: 0.11512785851955414	accuracy: 96.56%	val_loss: 0.9614036418497562	val_accuracy: 74.93%
[14, 60]	loss: 0.07953784661367536	accuracy: 98.12%	val_loss: 1.0151421763002872	val_accuracy: 75.02%
[14, 70]	loss: 0.13037337847054004	accuracy: 95.94%	val_loss: 0.9893855499103665	val_accuracy: 74.23%
[14, 80]	loss: 0.07448555342853069	accuracy: 97.81%	val_loss: 1.0185956303030252	val_accuracy: 76.22%
[14, 90]	loss: 0.08787544537335634	accuracy: 97.81%	val_loss: 0.9538600156083703	val_accuracy: 77.12%
[14, 100]	loss: 0.07245448045432568	accuracy: 98.44%	val_loss: 0.9358690204098821	val_accuracy: 76.62%
[14, 110]	loss: 0.050964575819671155	accuracy: 98.12%	val_loss: 0.9907288872636855	val_accuracy: 77.22%
[14, 120]	loss: 0.11003897897899151	accuracy: 96.56%	val_loss: 0.8850810807198286	val_accuracy: 77.02%
[14, 130]	loss: 0.0825831675902009	accuracy: 96.88%	val_loss: 0.8891354277729988	val_accuracy: 75.62%
[14, 140]	loss: 0.0966179920360446	accuracy: 97.19%	val_loss: 1.086267976090312	val_accuracy: 75.42%
[14, 150]	loss: 0.17631169483065606	accuracy: 94.06%	val_loss: 1.3416821043938398	val_accuracy: 68.53%
[14, 160]	loss: 0.32007497027516363	accuracy: 90.00%	val_loss: 1.2209116648882627	val_accuracy: 72.13%
[14, 170]	loss: 0.34369288086891175	accuracy: 88.44%	val_loss: 1.230240099132061	val_accuracy: 69.03%
[14, 180]	loss: 0.38305840343236924	accuracy: 88.44%	val_loss: 1.3252412304282188	val_accuracy: 67.43%
[15, 10]	loss: 0.49573964923620223	accuracy: 85.94%	val_loss: 1.1905236234888434	val_accuracy: 69.13%
[15, 20]	loss: 0.5185202077031136	accuracy: 81.56%	val_loss: 1.1884257979691029	val_accuracy: 67.03%
[15, 30]	loss: 0.45216861218214033	accuracy: 85.62%	val_loss: 1.0893761916086078	val_accuracy: 68.53%
[15, 40]	loss: 0.4720574736595154	accuracy: 85.31%	val_loss: 0.9197953986003995	val_accuracy: 72.23%
[15, 50]	loss: 0.41242883801460267	accuracy: 83.75%	val_loss: 1.0666353879496455	val_accuracy: 71.93%
[15, 60]	loss: 0.395389649271965	accuracy: 83.44%	val_loss: 0.976802496239543	val_accuracy: 72.63%
[15, 70]	loss: 0.2921566590666771	accuracy: 90.31%	val_loss: 0.8956059133633971	val_accuracy: 73.73%
[15, 80]	loss: 0.31494794860482217	accuracy: 85.94%	val_loss: 0.8800910674035549	val_accuracy: 75.52%
[15, 90]	loss: 0.33348817825317384	accuracy: 87.81%	val_loss: 0.8789001628756523	val_accuracy: 73.73%
[15, 100]	loss: 0.2743725150823593	accuracy: 88.75%	val_loss: 0.9440681338310242	val_accuracy: 74.53%
[15, 110]	loss: 0.2581175535917282	accuracy: 92.81%	val_loss: 0.8643502136692405	val_accuracy: 74.03%
[15, 120]	loss: 0.24159217476844788	accuracy: 91.88%	val_loss: 0.883890375494957	val_accuracy: 74.43%

[15, 130]	loss: 0.21111626625061036	accuracy: 92.19%	val_loss: 0.8452020939439535	val_accuracy: 75.62%
[15, 140]	loss: 0.20894777700304984	accuracy: 91.25%	val_loss: 0.9492220850661397	val_accuracy: 73.63%
[15, 150]	loss: 0.21466444283723832	accuracy: 92.19%	val_loss: 0.8630329100415111	val_accuracy: 76.32%
[15, 160]	loss: 0.1834508553147316	accuracy: 93.12%	val_loss: 0.9224801063537598	val_accuracy: 74.93%
[15, 170]	loss: 0.2106838312000036	accuracy: 93.12%	val_loss: 0.859103687107563	val_accuracy: 75.12%
[15, 180]	loss: 0.24233685657382012	accuracy: 90.94%	val_loss: 0.9086406696587801	val_accuracy: 74.23%
[16, 10]	loss: 0.08733085431158542	accuracy: 98.12%	val_loss: 0.9539161436259747	val_accuracy: 75.12%
[16, 20]	loss: 0.06425436902791262	accuracy: 98.75%	val_loss: 0.855805394705385	val_accuracy: 74.93%
[16, 30]	loss: 0.09046242497861386	accuracy: 96.56%	val_loss: 0.9085665065795183	val_accuracy: 75.92%
[16, 40]	loss: 0.08927472587674856	accuracy: 96.56%	val_loss: 0.9482476953417063	val_accuracy: 73.93%
[16, 50]	loss: 0.09410128332674503	accuracy: 97.19%	val_loss: 0.9348852597177029	val_accuracy: 75.22%
[16, 60]	loss: 0.07975856978446245	accuracy: 97.50%	val_loss: 0.8442891910672188	val_accuracy: 76.72%
[16, 70]	loss: 0.10170113686472178	accuracy: 96.56%	val_loss: 0.8486360562965274	val_accuracy: 77.62%
[16, 80]	loss: 0.07817066200077534	accuracy: 97.81%	val_loss: 0.8715338464826345	val_accuracy: 77.12%
[16, 90]	loss: 0.06954825576394796	accuracy: 98.44%	val_loss: 0.873411831445992	val_accuracy: 77.32%
[16, 100]	loss: 0.0755876999348402	accuracy: 97.81%	val_loss: 0.9274545875377953	val_accuracy: 76.42%
[16, 110]	loss: 0.0729253027588129	accuracy: 97.81%	val_loss: 0.8743837997317314	val_accuracy: 76.02%
[16, 120]	loss: 0.12511163540184497	accuracy: 96.56%	val_loss: 0.9129260182380676	val_accuracy: 75.62%
[16, 130]	loss: 0.07610094640403986	accuracy: 97.81%	val_loss: 1.0298645375296474	val_accuracy: 76.72%
[16, 140]	loss: 0.07126870704814792	accuracy: 98.12%	val_loss: 1.0183548610657454	val_accuracy: 76.92%
[16, 150]	loss: 0.07399176768958568	accuracy: 98.12%	val_loss: 0.8689677864313126	val_accuracy: 76.42%
[16, 160]	loss: 0.06106665097177029	accuracy: 97.50%	val_loss: 0.9331467002630234	val_accuracy: 76.02%
[16, 170]	loss: 0.05677299704402685	accuracy: 98.12%	val_loss: 1.0377621054649353	val_accuracy: 75.42%
[16, 180]	loss: 0.0910740016028285	accuracy: 96.25%	val_loss: 0.9280209178104997	val_accuracy: 75.52%
[17, 10]	loss: 0.041447667963802816	accuracy: 99.06%	val_loss: 1.1642630454152822	val_accuracy: 76.52%
[17, 20]	loss: 0.026961497962474823	accuracy: 100.00%	val_loss: 0.9391981400549412	val_accuracy: 78.02%
[17, 30]	loss: 0.04565362446010113	accuracy: 98.44%	val_loss: 0.9468518458306789	val_accuracy: 76.82%
[17, 40]	loss: 0.02720305062830448	accuracy: 99.69%	val_loss: 1.060548109933734	val_accuracy: 75.22%
[17, 50]	loss: 0.025791312195360662	accuracy: 99.69%	val_loss: 0.9041395271196961	val_accuracy: 76.62%
[17, 60]	loss: 0.030939908465370536	accuracy: 98.75%	val_loss: 0.8954761540517211	val_accuracy: 76.52%
[17, 70]	loss: 0.025526810600422323	accuracy: 98.75%	val_loss: 1.0027928375639021	val_accuracy: 77.52%
[17, 80]	loss: 0.022513466235250235	accuracy: 99.69%	val_loss: 0.9466279679909348	val_accuracy: 77.52%
[17, 90]	loss: 0.025520556280389427	accuracy: 99.38%	val_loss: 1.0115755246952176	val_accuracy: 76.82%
[17, 100]	loss: 0.026385349314659835	accuracy: 99.69%	val_loss: 1.0579945258796215	val_accuracy: 76.82%
[17, 110]	loss: 0.02802219456061721	accuracy: 99.06%	val_loss: 1.0378834512084723	val_accuracy: 76.52%
[17, 120]	loss: 0.03155543589964509	accuracy: 99.38%	val_loss: 0.9523917756741866	val_accuracy: 76.22%
[17, 130]	loss: 0.029002569895237685	accuracy: 99.69%	val_loss: 1.0290159238502383	val_accuracy: 76.92%
[17, 140]	loss: 0.02098571276292205	accuracy: 100.00%	val_loss: 1.0427995431236923	val_accuracy: 77.82%
[17, 150]	loss: 0.03567408134695142	accuracy: 98.75%	val_loss: 1.0344724897295237	val_accuracy: 76.12%
[17, 160]	loss: 0.023623794317245483	accuracy: 99.69%	val_loss: 1.096784452907741	val_accuracy: 74.63%
[17, 170]	loss: 0.027889063768088817	accuracy: 99.06%	val_loss: 1.0636102235876024	val_accuracy: 76.12%
[17, 180]	loss: 0.03421892351470888	accuracy: 99.69%	val_loss: 1.0321167488582432	val_accuracy: 75.82%

[18, 10]	loss: 0.04400113318115473	accuracy: 97.81%	val_loss: 1.1583909673208836	val_accuracy: 74.93%
[18, 20]	loss: 0.042383333574980496	accuracy: 98.44%	val_loss: 1.1642248288262635	val_accuracy: 73.93%
[18, 30]	loss: 0.021744605619460344	accuracy: 100.00%	val_loss: 1.0931975021958351	val_accuracy: 74.93%
[18, 40]	loss: 0.039773119427263734	accuracy: 99.06%	val_loss: 1.047281626611948	val_accuracy: 75.02%
[18, 50]	loss: 0.03474041782319546	accuracy: 98.44%	val_loss: 1.106070772279054	val_accuracy: 76.42%
[18, 60]	loss: 0.03608318879269064	accuracy: 98.75%	val_loss: 1.1108182622119784	val_accuracy: 77.22%
[18, 70]	loss: 0.045808013435453175	accuracy: 97.81%	val_loss: 1.1047103116288781	val_accuracy: 75.82%
[18, 80]	loss: 0.05174839673563838	accuracy: 98.75%	val_loss: 1.0082464395090938	val_accuracy: 76.52%
[18, 90]	loss: 0.05389313846826553	accuracy: 98.44%	val_loss: 1.0428943894803524	val_accuracy: 76.12%
[18, 100]	loss: 0.03801048304885626	accuracy: 98.44%	val_loss: 1.1771902267355472	val_accuracy: 77.32%
[18, 110]	loss: 0.06311619449406862	accuracy: 98.12%	val_loss: 1.1226431066170335	val_accuracy: 75.72%
[18, 120]	loss: 0.07573584276251495	accuracy: 97.50%	val_loss: 0.9691665912978351	val_accuracy: 77.32%
[18, 130]	loss: 0.041198386205360295	accuracy: 98.75%	val_loss: 1.0499125043861568	val_accuracy: 77.22%
[18, 140]	loss: 0.044523959048092365	accuracy: 98.75%	val_loss: 1.1034487765282393	val_accuracy: 76.02%
[18, 150]	loss: 0.028677542693912983	accuracy: 99.06%	val_loss: 1.0832493295893073	val_accuracy: 75.42%
[18, 160]	loss: 0.033233076147735116	accuracy: 98.75%	val_loss: 1.1513370610773563	val_accuracy: 76.32%
[18, 170]	loss: 0.05017867516726256	accuracy: 98.12%	val_loss: 1.156607878394425	val_accuracy: 76.02%
[18, 180]	loss: 0.03585084555670619	accuracy: 99.38%	val_loss: 1.153944669291377	val_accuracy: 76.22%
[19, 10]	loss: 0.03282549101859331	accuracy: 98.44%	val_loss: 1.0780758364126086	val_accuracy: 76.02%
[19, 20]	loss: 0.024516710732132196	accuracy: 99.38%	val_loss: 1.1376294349320233	val_accuracy: 76.52%
[19, 30]	loss: 0.024862885614857076	accuracy: 99.38%	val_loss: 1.1609821543097496	val_accuracy: 76.92%
[19, 40]	loss: 0.035346540971659124	accuracy: 98.44%	val_loss: 1.1922883912920952	val_accuracy: 77.02%
[19, 50]	loss: 0.031171096954494715	accuracy: 99.69%	val_loss: 1.2085969243198633	val_accuracy: 76.52%
[19, 60]	loss: 0.04109780993312597	accuracy: 98.75%	val_loss: 1.297060564160347	val_accuracy: 75.52%
[19, 70]	loss: 0.05811179550364613	accuracy: 98.12%	val_loss: 1.210824977606535	val_accuracy: 76.22%
[19, 80]	loss: 0.036411972250789404	accuracy: 98.75%	val_loss: 1.0613366484176368	val_accuracy: 75.82%
[19, 90]	loss: 0.03166766818612814	accuracy: 99.06%	val_loss: 1.159876998513937	val_accuracy: 75.52%
[19, 100]	loss: 0.04073687865165994	accuracy: 98.12%	val_loss: 1.1419364856556058	val_accuracy: 75.32%
[19, 110]	loss: 0.037305661290884015	accuracy: 98.75%	val_loss: 1.0927963587455451	val_accuracy: 76.42%
[19, 120]	loss: 0.04633616032078862	accuracy: 98.75%	val_loss: 1.3095095697790384	val_accuracy: 76.42%
[19, 130]	loss: 0.0330146221909672	accuracy: 99.06%	val_loss: 1.3624939247965813	val_accuracy: 76.52%
[19, 140]	loss: 0.04918899382464588	accuracy: 98.44%	val_loss: 1.263733577914536	val_accuracy: 75.92%
[19, 150]	loss: 0.01980252468492836	accuracy: 99.69%	val_loss: 1.2078442415222526	val_accuracy: 75.22%
[19, 160]	loss: 0.04361733673140407	accuracy: 98.12%	val_loss: 1.1092464574612677	val_accuracy: 74.33%
[19, 170]	loss: 0.028691383730620145	accuracy: 99.38%	val_loss: 1.2624807432293892	val_accuracy: 75.92%
[19, 180]	loss: 0.03292649185750633	accuracy: 98.44%	val_loss: 1.590342273004353	val_accuracy: 75.42%
[20, 10]	loss: 0.02175521143944934	accuracy: 99.38%	val_loss: 1.2991049876436591	val_accuracy: 75.82%
[20, 20]	loss: 0.031927933520637455	accuracy: 99.38%	val_loss: 1.1845022700726986	val_accuracy: 75.92%
[20, 30]	loss: 0.041074052918702364	accuracy: 98.12%	val_loss: 1.232221613638103	val_accuracy: 74.13%
[20, 40]	loss: 0.0348235527984798	accuracy: 98.75%	val_loss: 1.2137351660057902	val_accuracy: 75.32%
[20, 50]	loss: 0.04165551303885877	accuracy: 99.06%	val_loss: 1.2063562273979187	val_accuracy: 73.93%
[20, 60]	loss: 0.05446967093739659	accuracy: 98.12%	val_loss: 1.411851180717349	val_accuracy: 74.93%

```

[20, 70] loss: 0.06584543150383979 | accuracy: 98.75% | val_loss: 1.2757055163383484 | val_accuracy: 74.13%
[20, 80] loss: 0.04780685137957334 | accuracy: 98.12% | val_loss: 1.2773096915334463 | val_accuracy: 73.43%
[20, 90] loss: 0.02913059606216848 | accuracy: 99.06% | val_loss: 1.2940648831427097 | val_accuracy: 74.83%
[20, 100] loss: 0.03795592403039336 | accuracy: 98.75% | val_loss: 1.4006359623817843 | val_accuracy: 75.62%
[20, 110] loss: 0.06953518060036004 | accuracy: 97.81% | val_loss: 1.2208834467455745 | val_accuracy: 77.42%
[20, 120] loss: 0.07346122339367867 | accuracy: 97.81% | val_loss: 1.1770339645445347 | val_accuracy: 76.32%
[20, 130] loss: 0.03822103044949472 | accuracy: 99.06% | val_loss: 1.242755758576095 | val_accuracy: 75.82%
[20, 140] loss: 0.03869967209175229 | accuracy: 98.75% | val_loss: 1.226339160464704 | val_accuracy: 74.23%
[20, 150] loss: 0.039054635399952534 | accuracy: 98.44% | val_loss: 1.2944226926192641 | val_accuracy: 74.83%
[20, 160] loss: 0.05315618868917227 | accuracy: 98.12% | val_loss: 1.2595800887793303 | val_accuracy: 75.72%
[20, 170] loss: 0.045203320204745975 | accuracy: 98.12% | val_loss: 1.4370029252022505 | val_accuracy: 75.22%
[20, 180] loss: 0.03638701136223972 | accuracy: 98.44% | val_loss: 1.2892804685980082 | val_accuracy: 75.82%

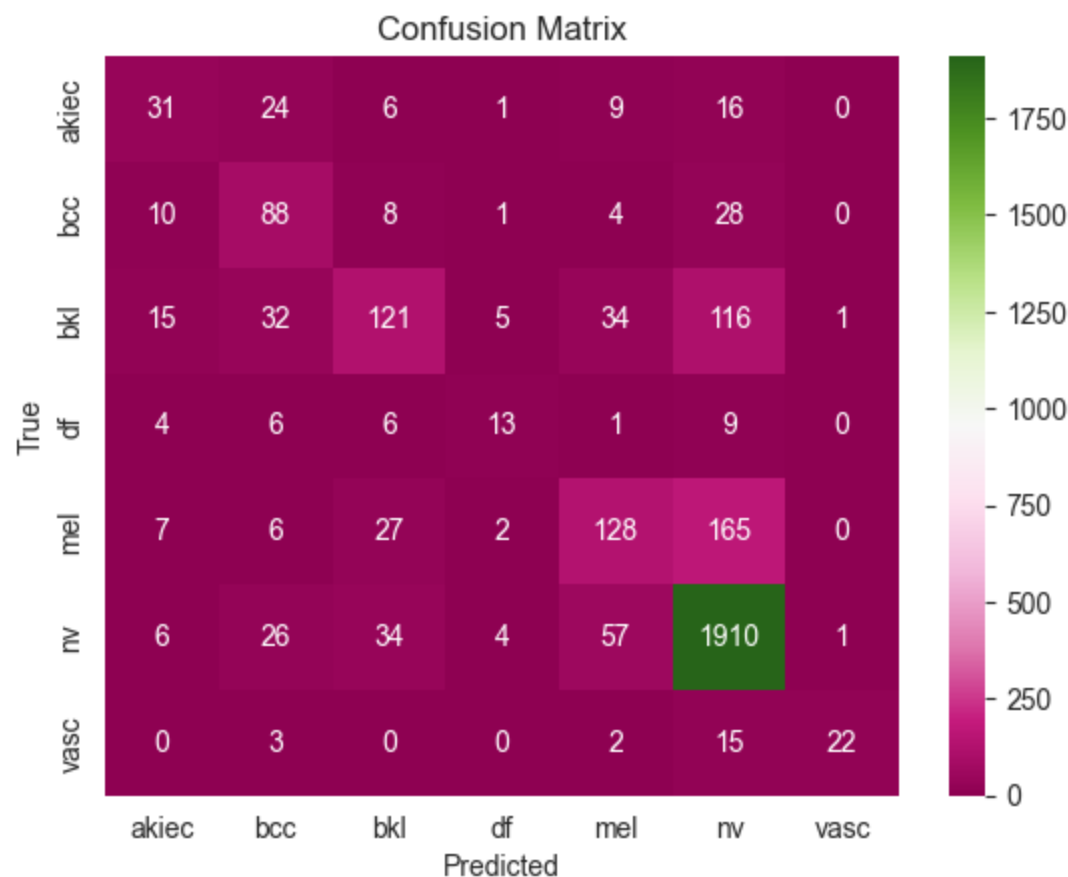
```

```

In [34]: # Get the number_to_label_array so that we can use it later to plot the confusion matrix
         complex_number_to_label_array = complex_testloader.dataset.dataset.number_to_label_array
         # Test the model
         test_net(complex_model, complex_testloader, complex_number_to_label_array, loss=complex_loss, device=device)

```

Test loss: 1.1766993999481201 | Test accuracy: 77.00%



## Observations


- We can see again that the model has a strong inclination to predict the label `nv`, but this time not so much for `mel`
- The model predictions are mostly concentrated on the diagonal, (correct predictions), which is a VERY good thing
- There is also a huge improvement on `df`, the model actually predicted the label, and it did it correctly most of the time.
- Interestingly, for the `vasc` label, the model is not as good as the Simple model, precision is still good, but the recall is not as good as the Simple model

```
In [18]: # Copying the lists so that I don't overwrite them
complex_train_loss = copy.deepcopy(train_loss)
complex_val_loss = copy.deepcopy(val_loss)
```

```
complex_train_accuracy = copy.deepcopy(train_accuracy)
complex_val_accuracy = copy.deepcopy(val_accuracy)
```

```
In [19]: x_axis = list(map(lambda x:x*10,range(len(complex_train_loss))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,complex_train_loss, label='Training Loss')
plt.plot(x_axis , complex_val_loss,label='Validation Loss')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 400.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid()
plt.show()
```



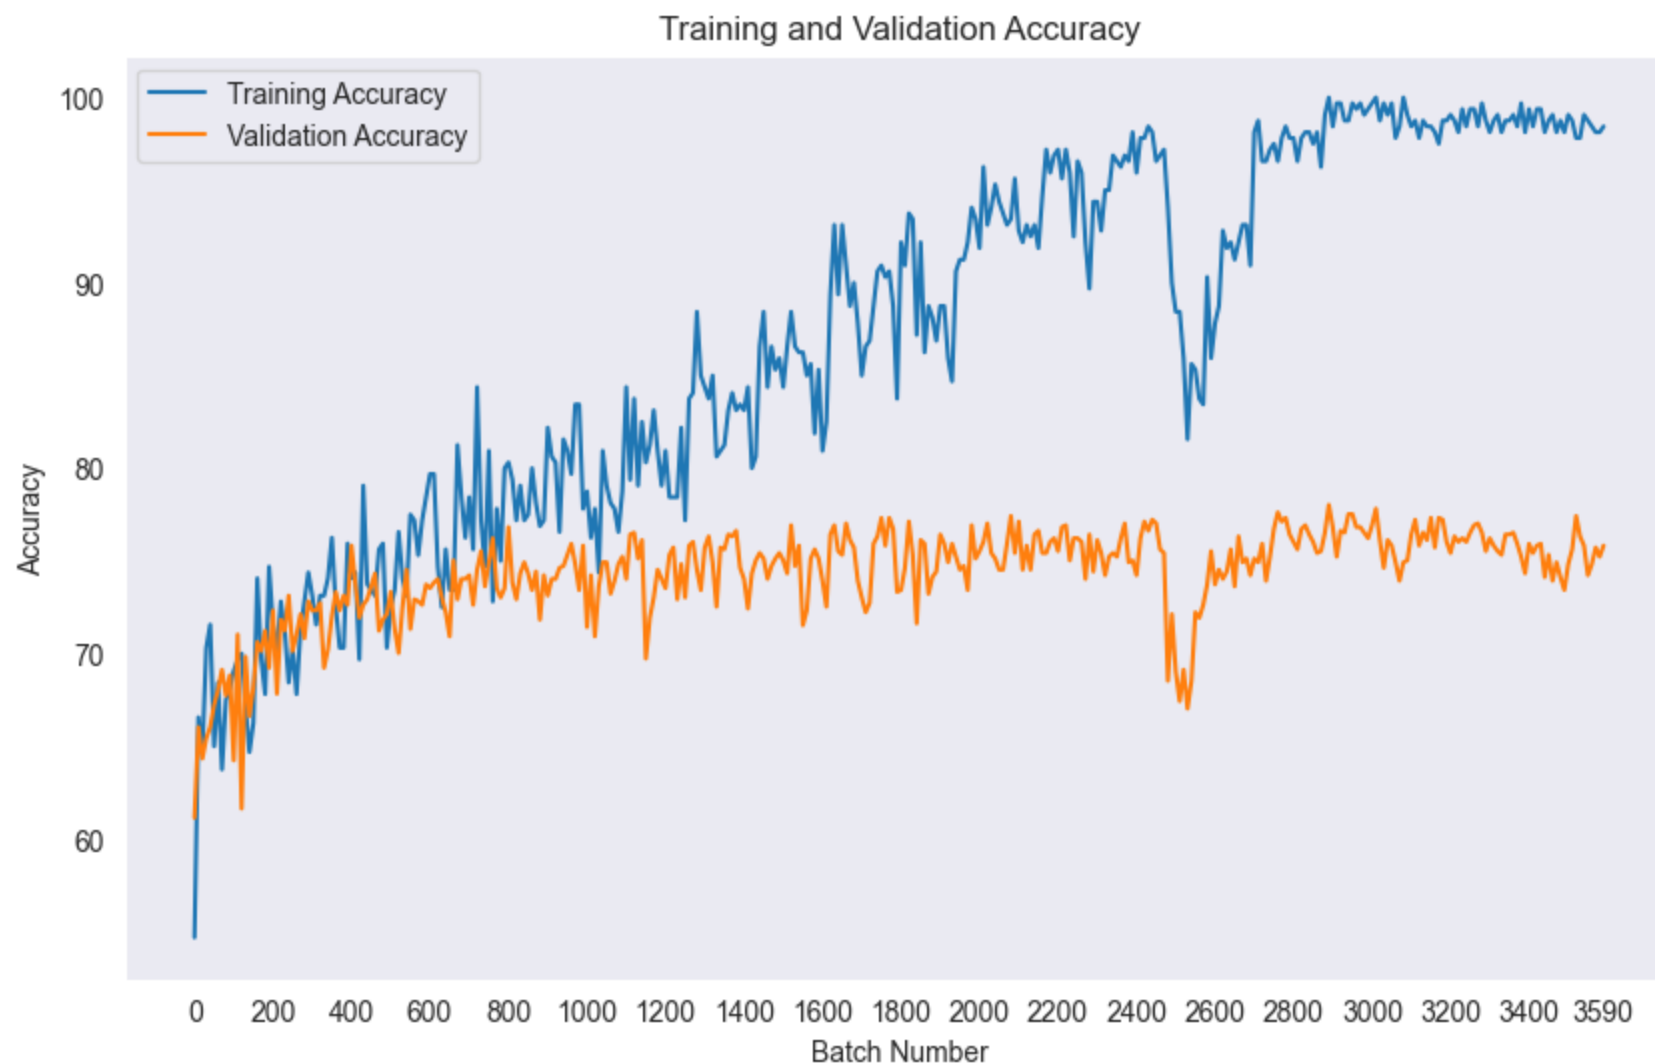


## Observations

- We can see that the model starts to overfit at around 400-500 batches, as the validation loss starts to increase, while the training loss keeps decreasing
- We can see that the graph is less "jagged", which I think means that the model isn't overshooting the minimum loss, meaning the optimizer setting and choice are way better than the ones for the Simple model.

- There is a weird spike at 2400 batches, which I think might mean that while the optimizer was steadily decreasing the loss, there was a sudden "hill" on the feature space.
- This graph indicates that the model is doing two things better than the Simple model, it is not overshooting the minimum loss (the optimizer choice and settings are way better), and the model starts overfitting faster, which means that the model is learning faster.

```
In [20]: x_axis = list(map(lambda x:x*10,range(len(complex_train_accuracy))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,complex_train_accuracy, label='Training Accuracy')
plt.plot(x_axis , complex_val_accuracy,label='Validation Accuracy')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 200.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.grid()
plt.show()
```



## Observations

- In this graph the overfitting is really visible, because the validation accuracy is not decreasing, but it is not increasing either, while the training accuracy is increasing.
- This time it also less "jagged"
- and we can still see the same weird spike at 2400 batches

```
In [21]: # I will also save the models so that I can load them later instead of training them again (takes a lot of time)
torch.save(model.state_dict(), 'results/model_states/simple_model.pth')
torch.save(complex_model.state_dict(), 'results/model_states/complex_model.pth')

# We don't need to save the dataloaders, because we use the same seed every time
# so the train, validation and test sets will be the same (seed = 42)
# But we do need to save the train, validation loss and accuracy
# so that we can plot them later without having to train the models again

torch.save(simple_train_loss, 'results/general/simple_train_loss.list.pyobj')
torch.save(simple_val_loss, 'results/general/simple_val_loss.list.pyobj')
torch.save(simple_train_accuracy, 'results/general/simple_train_accuracy.list.pyobj')
torch.save(simple_val_accuracy, 'results/general/simple_val_accuracy.list.pyobj')

torch.save(complex_train_loss, 'results/general/complex_train_loss.list.pyobj')
torch.save(complex_val_loss, 'results/general/complex_val_loss.list.pyobj')
torch.save(complex_train_accuracy, 'results/general/complex_train_accuracy.list.pyobj')
torch.save(complex_val_accuracy, 'results/general/complex_val_accuracy.list.pyobj')
```

## ResNet: Transfer Learning

```
In [12]: import torchvision.transforms as transforms
# The transforms for the ResNet model are different from the ones for the Simple and Complex models
# create the transforms as given in the assignment
data_transforms = {
    'train': torchvision.transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])]),
    'val': torchvision.transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])
},

# Create the dataset
```

```
resnet_dataset = MLProject2Dataset(data_dir='data', metadata_fname='metadata.csv')

# To apply the different transforms to the train, validation and test sets
# We need to first split the dataset into train, validation and test to separate them as python objects
# and then apply the transforms to each one of them
# The split is train:val:test = 60%:10%:30% and because we use different transform for the train
# and validation/test sets, we need to split into train:val+test = 60%:40% and then split the val+test
resnet_train_dataset, resnet_val_test_dataset = torch.utils.data.random_split(resnet_dataset, [.6, .4],
                                                                              generator=torch.Generator().manual_seed(42)
                                                                              )

# Set the transforms for the train set
resnet_train_dataset.dataset.transform = data_transforms['train']
# Split the val+test set into val and test
# the split in the original dataset is 10%:30% when the total of the two is 40%
# now that the "total" of the two is 100%, the split is 25%:75% to get the correct split
resnet_val_dataset, resnet_test_dataset = torch.utils.data.random_split(resnet_val_test_dataset, [.25, .75],
                                                                          generator=torch.Generator().manual_seed(42)
                                                                          )

# Set the transforms for the val and test sets
# I don't really need to do the transform setting for both of them, because they are the same python object
# under the hood, but I do it for readability, and just in case that the implementation of the random_split
# changes in the future
resnet_val_dataset.dataset.transform = data_transforms['val']
resnet_test_dataset.dataset.transform = data_transforms['val']

# Create the dataloaders
BATCH_SIZE = 32
resnet_trainloader = DataLoader(resnet_train_dataset, batch_size=BATCH_SIZE, shuffle=True)
resnet_valloader = DataLoader(resnet_val_dataset, batch_size=BATCH_SIZE, shuffle=True)
resnet_testloader = DataLoader(resnet_test_dataset, batch_size=BATCH_SIZE, shuffle=True)

# Get the pretrained model
resnet_model = torchvision.models.resnet34(weights="DEFAULT")
# Change the last layer to have 7 outputs instead of 1000 to match our dataset
in_features = resnet_model.fc.in_features
resnet_model.fc = nn.Linear(in_features, 7)
resnet_model = resnet_model.to(device)
# Print the model summary
summary(resnet_model, (3, 224, 224))
# Create the optimizer
resnet_optimizer = optim.SGD(resnet_model.parameters(), lr=0.001, momentum=0.9)
```

```
# Create the Loss function
resnet_loss = nn.CrossEntropyLoss()
# Train the model
train_net(resnet_model, resnet_trainloader, resnet_valloader, epochs=5, optimizer=resnet_optimizer,
          loss=resnet_loss, device=device)
# This cell took 17 minutes to run on my machine
```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
ReLU-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	36,864
BatchNorm2d-6	[-1, 64, 56, 56]	128
ReLU-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 64, 56, 56]	36,864
BatchNorm2d-9	[-1, 64, 56, 56]	128
ReLU-10	[-1, 64, 56, 56]	0
BasicBlock-11	[-1, 64, 56, 56]	0
Conv2d-12	[-1, 64, 56, 56]	36,864
BatchNorm2d-13	[-1, 64, 56, 56]	128
ReLU-14	[-1, 64, 56, 56]	0
Conv2d-15	[-1, 64, 56, 56]	36,864
BatchNorm2d-16	[-1, 64, 56, 56]	128
ReLU-17	[-1, 64, 56, 56]	0
BasicBlock-18	[-1, 64, 56, 56]	0
Conv2d-19	[-1, 64, 56, 56]	36,864
BatchNorm2d-20	[-1, 64, 56, 56]	128
ReLU-21	[-1, 64, 56, 56]	0
Conv2d-22	[-1, 64, 56, 56]	36,864
BatchNorm2d-23	[-1, 64, 56, 56]	128
ReLU-24	[-1, 64, 56, 56]	0
BasicBlock-25	[-1, 64, 56, 56]	0
Conv2d-26	[-1, 128, 28, 28]	73,728
BatchNorm2d-27	[-1, 128, 28, 28]	256
ReLU-28	[-1, 128, 28, 28]	0
Conv2d-29	[-1, 128, 28, 28]	147,456
BatchNorm2d-30	[-1, 128, 28, 28]	256
Conv2d-31	[-1, 128, 28, 28]	8,192
BatchNorm2d-32	[-1, 128, 28, 28]	256
ReLU-33	[-1, 128, 28, 28]	0
BasicBlock-34	[-1, 128, 28, 28]	0
Conv2d-35	[-1, 128, 28, 28]	147,456
BatchNorm2d-36	[-1, 128, 28, 28]	256
ReLU-37	[-1, 128, 28, 28]	0
Conv2d-38	[-1, 128, 28, 28]	147,456
BatchNorm2d-39	[-1, 128, 28, 28]	256

ReLU-40	[-1, 128, 28, 28]	0
BasicBlock-41	[-1, 128, 28, 28]	0
Conv2d-42	[-1, 128, 28, 28]	147,456
BatchNorm2d-43	[-1, 128, 28, 28]	256
ReLU-44	[-1, 128, 28, 28]	0
Conv2d-45	[-1, 128, 28, 28]	147,456
BatchNorm2d-46	[-1, 128, 28, 28]	256
ReLU-47	[-1, 128, 28, 28]	0
BasicBlock-48	[-1, 128, 28, 28]	0
Conv2d-49	[-1, 128, 28, 28]	147,456
BatchNorm2d-50	[-1, 128, 28, 28]	256
ReLU-51	[-1, 128, 28, 28]	0
Conv2d-52	[-1, 128, 28, 28]	147,456
BatchNorm2d-53	[-1, 128, 28, 28]	256
ReLU-54	[-1, 128, 28, 28]	0
BasicBlock-55	[-1, 128, 28, 28]	0
Conv2d-56	[-1, 256, 14, 14]	294,912
BatchNorm2d-57	[-1, 256, 14, 14]	512
ReLU-58	[-1, 256, 14, 14]	0
Conv2d-59	[-1, 256, 14, 14]	589,824
BatchNorm2d-60	[-1, 256, 14, 14]	512
Conv2d-61	[-1, 256, 14, 14]	32,768
BatchNorm2d-62	[-1, 256, 14, 14]	512
ReLU-63	[-1, 256, 14, 14]	0
BasicBlock-64	[-1, 256, 14, 14]	0
Conv2d-65	[-1, 256, 14, 14]	589,824
BatchNorm2d-66	[-1, 256, 14, 14]	512
ReLU-67	[-1, 256, 14, 14]	0
Conv2d-68	[-1, 256, 14, 14]	589,824
BatchNorm2d-69	[-1, 256, 14, 14]	512
ReLU-70	[-1, 256, 14, 14]	0
BasicBlock-71	[-1, 256, 14, 14]	0
Conv2d-72	[-1, 256, 14, 14]	589,824
BatchNorm2d-73	[-1, 256, 14, 14]	512
ReLU-74	[-1, 256, 14, 14]	0
Conv2d-75	[-1, 256, 14, 14]	589,824
BatchNorm2d-76	[-1, 256, 14, 14]	512
ReLU-77	[-1, 256, 14, 14]	0
BasicBlock-78	[-1, 256, 14, 14]	0
Conv2d-79	[-1, 256, 14, 14]	589,824
BatchNorm2d-80	[-1, 256, 14, 14]	512
ReLU-81	[-1, 256, 14, 14]	0



Conv2d-82	[-1, 256, 14, 14]	589,824
BatchNorm2d-83	[-1, 256, 14, 14]	512
ReLU-84	[-1, 256, 14, 14]	0
BasicBlock-85	[-1, 256, 14, 14]	0
Conv2d-86	[-1, 256, 14, 14]	589,824
BatchNorm2d-87	[-1, 256, 14, 14]	512
ReLU-88	[-1, 256, 14, 14]	0
Conv2d-89	[-1, 256, 14, 14]	589,824
BatchNorm2d-90	[-1, 256, 14, 14]	512
ReLU-91	[-1, 256, 14, 14]	0
BasicBlock-92	[-1, 256, 14, 14]	0
Conv2d-93	[-1, 256, 14, 14]	589,824
BatchNorm2d-94	[-1, 256, 14, 14]	512
ReLU-95	[-1, 256, 14, 14]	0
Conv2d-96	[-1, 256, 14, 14]	589,824
BatchNorm2d-97	[-1, 256, 14, 14]	512
ReLU-98	[-1, 256, 14, 14]	0
BasicBlock-99	[-1, 256, 14, 14]	0
Conv2d-100	[-1, 512, 7, 7]	1,179,648
BatchNorm2d-101	[-1, 512, 7, 7]	1,024
ReLU-102	[-1, 512, 7, 7]	0
Conv2d-103	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-104	[-1, 512, 7, 7]	1,024
Conv2d-105	[-1, 512, 7, 7]	131,072
BatchNorm2d-106	[-1, 512, 7, 7]	1,024
ReLU-107	[-1, 512, 7, 7]	0
BasicBlock-108	[-1, 512, 7, 7]	0
Conv2d-109	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-110	[-1, 512, 7, 7]	1,024
ReLU-111	[-1, 512, 7, 7]	0
Conv2d-112	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-113	[-1, 512, 7, 7]	1,024
ReLU-114	[-1, 512, 7, 7]	0
BasicBlock-115	[-1, 512, 7, 7]	0
Conv2d-116	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-117	[-1, 512, 7, 7]	1,024
ReLU-118	[-1, 512, 7, 7]	0
Conv2d-119	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-120	[-1, 512, 7, 7]	1,024
ReLU-121	[-1, 512, 7, 7]	0
BasicBlock-122	[-1, 512, 7, 7]	0
AdaptiveAvgPool2d-123	[-1, 512, 1, 1]	0

```

Linear-124          [-1, 7]          3,591
=====
Total params: 21,288,263
Trainable params: 21,288,263
Non-trainable params: 0
-----
Input size (MB): 0.57
Forward/backward pass size (MB): 96.28
Params size (MB): 81.21
Estimated Total Size (MB): 178.06
-----
[1, 10] loss: 1.642268431186676 | accuracy: 46.56% | val_loss: 1.3907485622912645 | val_accuracy: 65.17%
[1, 20] loss: 1.349936830997467 | accuracy: 55.62% | val_loss: 1.0695527996867895 | val_accuracy: 61.38%
[1, 30] loss: 1.027817076444626 | accuracy: 64.69% | val_loss: 1.0737572479993105 | val_accuracy: 65.47%
[1, 40] loss: 0.93227618932724 | accuracy: 66.88% | val_loss: 0.8777777878567576 | val_accuracy: 67.07%
[1, 50] loss: 0.8050221145153046 | accuracy: 73.12% | val_loss: 0.8232784047722816 | val_accuracy: 69.86%
[1, 60] loss: 0.8791698813438416 | accuracy: 68.75% | val_loss: 0.8064984260126948 | val_accuracy: 69.46%
[1, 70] loss: 0.811342766880989 | accuracy: 72.19% | val_loss: 0.7766866907477379 | val_accuracy: 71.06%
[1, 80] loss: 0.7647741436958313 | accuracy: 70.00% | val_loss: 0.8005622318014503 | val_accuracy: 70.36%
[1, 90] loss: 0.7021259933710098 | accuracy: 76.25% | val_loss: 0.7633304372429848 | val_accuracy: 70.86%
[1, 100] loss: 0.6783249080181122 | accuracy: 76.25% | val_loss: 0.7689546691253781 | val_accuracy: 72.16%
[1, 110] loss: 0.8222796976566314 | accuracy: 72.81% | val_loss: 0.7420624755322933 | val_accuracy: 74.15%
[1, 120] loss: 0.7508459329605103 | accuracy: 74.06% | val_loss: 0.712641834281385 | val_accuracy: 74.55%
[1, 130] loss: 0.7339757978916168 | accuracy: 75.00% | val_loss: 0.67294130474329 | val_accuracy: 74.95%
[1, 140] loss: 0.6550167083740235 | accuracy: 74.69% | val_loss: 0.6905804807320237 | val_accuracy: 74.95%
[1, 150] loss: 0.7602157443761826 | accuracy: 72.50% | val_loss: 0.66605294495821 | val_accuracy: 73.95%
[1, 160] loss: 0.735264602303505 | accuracy: 73.44% | val_loss: 0.6715724971145391 | val_accuracy: 75.35%
[1, 170] loss: 0.7272795557975769 | accuracy: 72.50% | val_loss: 0.6929608015343547 | val_accuracy: 75.65%
[1, 180] loss: 0.6356493890285492 | accuracy: 75.62% | val_loss: 0.6929528722539544 | val_accuracy: 75.95%
[2, 10] loss: 0.7148155957460404 | accuracy: 72.19% | val_loss: 0.6394696235656738 | val_accuracy: 75.65%
[2, 20] loss: 0.593882355093956 | accuracy: 78.12% | val_loss: 0.6394281582906842 | val_accuracy: 76.05%
[2, 30] loss: 0.7295987367630005 | accuracy: 73.12% | val_loss: 0.6325073651969433 | val_accuracy: 78.44%
[2, 40] loss: 0.6056247740983963 | accuracy: 77.19% | val_loss: 0.6356831612065434 | val_accuracy: 77.74%
[2, 50] loss: 0.62055723965168 | accuracy: 78.44% | val_loss: 0.6116708340123296 | val_accuracy: 78.44%
[2, 60] loss: 0.7556171715259552 | accuracy: 73.75% | val_loss: 0.5962635017931461 | val_accuracy: 77.64%
[2, 70] loss: 0.6275269210338592 | accuracy: 76.88% | val_loss: 0.5984478429891169 | val_accuracy: 78.84%
[2, 80] loss: 0.6145122349262238 | accuracy: 77.81% | val_loss: 0.605888256803155 | val_accuracy: 78.24%
[2, 90] loss: 0.6820860922336578 | accuracy: 73.44% | val_loss: 0.6034647654742002 | val_accuracy: 77.84%
[2, 100] loss: 0.5779212713241577 | accuracy: 79.69% | val_loss: 0.6279908828437328 | val_accuracy: 77.35%
[2, 110] loss: 0.6385987341403961 | accuracy: 75.00% | val_loss: 0.6046873684972525 | val_accuracy: 76.75%
[2, 120] loss: 0.6048689633607864 | accuracy: 78.12% | val_loss: 0.5938305016607046 | val_accuracy: 79.04%
[2, 130] loss: 0.6049979895353317 | accuracy: 77.19% | val_loss: 0.6162361348979175 | val_accuracy: 78.34%


```

[2, 140]	loss: 0.6912472933530808	accuracy: 77.50%	val_loss: 0.568815141916275	val_accuracy: 78.44%
[2, 150]	loss: 0.5483153820037842	accuracy: 79.69%	val_loss: 0.5871286988258362	val_accuracy: 78.34%
[2, 160]	loss: 0.6391437709331512	accuracy: 75.94%	val_loss: 0.5833015320822597	val_accuracy: 78.44%
[2, 170]	loss: 0.562196047604084	accuracy: 80.94%	val_loss: 0.6172289131209254	val_accuracy: 77.05%
[2, 180]	loss: 0.6487706303596497	accuracy: 75.62%	val_loss: 0.5705101564526558	val_accuracy: 80.04%
[3, 10]	loss: 0.5164233863353729	accuracy: 80.62%	val_loss: 0.6200332557782531	val_accuracy: 78.84%
[3, 20]	loss: 0.5740506201982498	accuracy: 79.06%	val_loss: 0.6363858114928007	val_accuracy: 76.75%
[3, 30]	loss: 0.6018024265766144	accuracy: 78.12%	val_loss: 0.5989496097899973	val_accuracy: 78.34%
[3, 40]	loss: 0.6214186817407608	accuracy: 77.50%	val_loss: 0.5657363114878535	val_accuracy: 79.24%
[3, 50]	loss: 0.6107508152723312	accuracy: 76.56%	val_loss: 0.5645522205159068	val_accuracy: 78.54%
[3, 60]	loss: 0.6165035307407379	accuracy: 77.81%	val_loss: 0.5697467951104045	val_accuracy: 78.84%
[3, 70]	loss: 0.5604213118553162	accuracy: 78.75%	val_loss: 0.5862571848556399	val_accuracy: 79.54%
[3, 80]	loss: 0.5533954083919526	accuracy: 79.69%	val_loss: 0.5434774262830615	val_accuracy: 80.34%
[3, 90]	loss: 0.6076442897319794	accuracy: 75.62%	val_loss: 0.5336125008761883	val_accuracy: 81.24%
[3, 100]	loss: 0.5147301197052002	accuracy: 81.56%	val_loss: 0.6021204991266131	val_accuracy: 79.14%
[3, 110]	loss: 0.5713592112064362	accuracy: 78.12%	val_loss: 0.6041157245635986	val_accuracy: 77.35%
[3, 120]	loss: 0.5880119532346726	accuracy: 79.06%	val_loss: 0.5649197357706726	val_accuracy: 79.14%
[3, 130]	loss: 0.4802358329296112	accuracy: 80.62%	val_loss: 0.5285905320197344	val_accuracy: 80.24%
[3, 140]	loss: 0.552993306517601	accuracy: 78.44%	val_loss: 0.5338416015729308	val_accuracy: 80.54%
[3, 150]	loss: 0.5499305129051208	accuracy: 79.38%	val_loss: 0.5600376049987972	val_accuracy: 78.74%
[3, 160]	loss: 0.6081220477819442	accuracy: 78.44%	val_loss: 0.5524249239824712	val_accuracy: 79.74%
[3, 170]	loss: 0.48097063302993776	accuracy: 83.12%	val_loss: 0.5882867947220802	val_accuracy: 80.04%
[3, 180]	loss: 0.539277458190918	accuracy: 80.94%	val_loss: 0.5632255030795932	val_accuracy: 78.94%
[4, 10]	loss: 0.5932123482227325	accuracy: 75.62%	val_loss: 0.5224190000444651	val_accuracy: 80.04%
[4, 20]	loss: 0.5854963183403015	accuracy: 76.56%	val_loss: 0.537953126244247	val_accuracy: 80.34%
[4, 30]	loss: 0.5422192841768265	accuracy: 79.06%	val_loss: 0.5475990492850542	val_accuracy: 79.04%
[4, 40]	loss: 0.5197598785161972	accuracy: 80.31%	val_loss: 0.5296813899185508	val_accuracy: 81.34%
[4, 50]	loss: 0.5124416083097458	accuracy: 78.44%	val_loss: 0.5186387663707137	val_accuracy: 80.64%
[4, 60]	loss: 0.5347197264432907	accuracy: 78.75%	val_loss: 0.5523659279569983	val_accuracy: 79.74%
[4, 70]	loss: 0.5630500793457032	accuracy: 78.75%	val_loss: 0.5039988532662392	val_accuracy: 81.24%
[4, 80]	loss: 0.5672984898090363	accuracy: 79.38%	val_loss: 0.5145617239177227	val_accuracy: 81.64%
[4, 90]	loss: 0.567824125289917	accuracy: 77.50%	val_loss: 0.538137529976666	val_accuracy: 80.84%
[4, 100]	loss: 0.4755730703473091	accuracy: 82.50%	val_loss: 0.563140376470983	val_accuracy: 79.44%
[4, 110]	loss: 0.5503921627998352	accuracy: 81.88%	val_loss: 0.5712568582966924	val_accuracy: 78.74%
[4, 120]	loss: 0.49675209224224093	accuracy: 83.44%	val_loss: 0.5520836571231484	val_accuracy: 79.84%
[4, 130]	loss: 0.4864620536565781	accuracy: 82.81%	val_loss: 0.5655836267396808	val_accuracy: 79.04%
[4, 140]	loss: 0.45151119530200956	accuracy: 81.25%	val_loss: 0.502715844893828	val_accuracy: 81.94%
[4, 150]	loss: 0.5292572036385537	accuracy: 78.75%	val_loss: 0.5315211932174861	val_accuracy: 80.24%
[4, 160]	loss: 0.42558020949363706	accuracy: 83.44%	val_loss: 0.49494313169270754	val_accuracy: 82.73%
[4, 170]	loss: 0.5385327190160751	accuracy: 80.31%	val_loss: 0.5412771161645651	val_accuracy: 80.24%
[4, 180]	loss: 0.4520386248826981	accuracy: 82.50%	val_loss: 0.49196505593135953	val_accuracy: 82.44%
[5, 10]	loss: 0.4870768666267395	accuracy: 83.12%	val_loss: 0.5109239374287426	val_accuracy: 82.53%

[5, 20]	loss: 0.4925625562667847	accuracy: 82.50%	val_loss: 0.543040793389082	val_accuracy: 80.24%
[5, 30]	loss: 0.4401986598968506	accuracy: 83.12%	val_loss: 0.5167415789328516	val_accuracy: 81.54%
[5, 40]	loss: 0.4419477045536041	accuracy: 83.44%	val_loss: 0.5368921086192131	val_accuracy: 80.84%
[5, 50]	loss: 0.37956525534391405	accuracy: 86.88%	val_loss: 0.5222408720292151	val_accuracy: 81.54%
[5, 60]	loss: 0.4026473075151443	accuracy: 84.38%	val_loss: 0.5177223826758564	val_accuracy: 81.74%
[5, 70]	loss: 0.49978149831295016	accuracy: 82.50%	val_loss: 0.5398647943511605	val_accuracy: 80.24%
[5, 80]	loss: 0.5197296351194381	accuracy: 81.56%	val_loss: 0.5233790478669107	val_accuracy: 80.44%
[5, 90]	loss: 0.4416507065296173	accuracy: 84.06%	val_loss: 0.5503846746869385	val_accuracy: 80.14%
[5, 100]	loss: 0.4366041451692581	accuracy: 84.06%	val_loss: 0.5024455799721181	val_accuracy: 81.74%
[5, 110]	loss: 0.4530560731887817	accuracy: 84.38%	val_loss: 0.4722372507676482	val_accuracy: 82.34%
[5, 120]	loss: 0.5389544427394867	accuracy: 79.69%	val_loss: 0.479295730125159	val_accuracy: 81.74%
[5, 130]	loss: 0.508551687002182	accuracy: 80.31%	val_loss: 0.5249498272314668	val_accuracy: 81.14%
[5, 140]	loss: 0.4603620916604996	accuracy: 83.44%	val_loss: 0.5169145530089736	val_accuracy: 80.74%
[5, 150]	loss: 0.5088219553232193	accuracy: 80.94%	val_loss: 0.5182687751948833	val_accuracy: 80.24%
[5, 160]	loss: 0.48052942752838135	accuracy: 84.38%	val_loss: 0.5208396627567708	val_accuracy: 80.84%
[5, 170]	loss: 0.46181192100048063	accuracy: 80.00%	val_loss: 0.5450489167124033	val_accuracy: 79.54%
[5, 180]	loss: 0.47742532193660736	accuracy: 82.50%	val_loss: 0.5315917734988034	val_accuracy: 80.94%

```
In [13]: # Get the number_to_label_array so that we can use it later to plot the confusion matrix
resnet_number_to_label_array = resnet_testloader.dataset.dataset.number_to_label_array
# Test the model
test_net(resnet_model, resnet_testloader, resnet_number_to_label_array, loss=resnet_loss, device=device)
```

Test loss: 0.5195795297622681 | Test accuracy: 80.93%




## Observations

- Way better performance than the Simple and Complex models, huge recall and precision for all labels
- The `bkl` label is now the biggest misclassification, for the `nv` label, which is weird in comparison to our previous results
- The precession on the `nv` label is WAY better than the Simple and Complex models, which is a good thing, because the `nv` label is the most common label in the dataset

```
In [14]: # Copying the lists so that I don't overwrite them
resnet_train_loss = copy.deepcopy(train_loss)
resnet_val_loss = copy.deepcopy(val_loss)
```

```
resnet_train_accuracy = copy.deepcopy(train_accuracy)
resnet_val_accuracy = copy.deepcopy(val_accuracy)
```

```
In [15]: x_axis = list(map(lambda x:x*10,range(len(resnet_train_loss))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,resnet_train_loss, label='Training Loss')
plt.plot(x_axis , resnet_val_loss,label='Validation Loss')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 100.0)
plt.xticks([*xticks, x_axis[-1]])
plt.ylim((0,2))
plt.xlabel('Batch Number')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid()
plt.show()
```




## Observations

- The model is WAY less "jagged", which could be the result of the optimizer choice and settings or because the number of batches is way less than the Simple and Complex models
- The validation and training curves are very close together, which is a good thing, because the model is improving on both sets
- The model doesn't overfit

```
In [16]: x_axis = list(map(lambda x:x*10,range(len(resnet_train_accuracy))))
plt.figure(figsize=(10,6))
plt.plot(x_axis ,resnet_train_accuracy, label='Training Accuracy')
plt.plot(x_axis , resnet_val_accuracy,label='Validation Accuracy')
xticks = np.arange(x_axis[0], x_axis[-1]+1, 100.0)
plt.xticks(*xticks, x_axis[-1])
plt.xlabel('Batch Number')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.grid()
plt.show()
```





## Observations

- Again, less "jagged" than the Simple and Complex models
- The model doesn't overfit
- The validation and training curves are very close together, which is a good thing, because the model is improving on both sets

- and something we can see here better than on the loss plot, is that the performance is better than the Simple and Complex models

## Bonus: Complex CNN with demographic data as input

First we need to remake the dataset as we also need to load the demographic data

```
In [17]: class MLProject2DatasetBonus(Dataset):
    def __init__(self, data_dir:str, metadata_fname='metadata.csv', transform: torchvision.transforms = None):
        self.data_dir = data_dir
        self.transform = transform

        self.df = pd.DataFrame(columns=['image_id', 'path'])

        # Get all files in data_dir
        files = glob.glob(data_dir + '/*/*.jpg')
        # If on windows, replace backslash with forward slash
        # on linux, this does nothing
        files = [file.replace(os.sep, '/') for file in files]
        # The line commented bellow was used for testing on a smaller dataset
        # files = files[:1000]

        # Add all files to dataframe
        for file in files:
            image_id = file.split('/')[-1].split('.')[0]
            path = file

            # Add to dataframe
            self.df.loc[len(self.df)] = [image_id, path]

        # Add metadata
        self.metadata = pd.read_csv(os.path.join(data_dir, metadata_fname))
        dx_categorical = pd.Categorical(self.metadata['dx'])
        self.number_to_label_array = [dx for dx in dx_categorical.categories]

        self.metadata['dx'] = dx_categorical.codes

        # Merge metadata with dataframe
```

```
self.df = self.df.merge(self.metadata, on='image_id')

# Drop the columns we don't need
self.df = self.df.drop(columns=["lesion_id", "dx_type"])

# Normalize the age column
self.df['age'] = self.df['age'].map(lambda x: x/100)
# Make 'localization' and 'sex' into one-hot columns
self.df = pd.get_dummies(self.df, columns=["localization", 'sex'])

def __len__(self):
    return len(self.df)

def __getitem__(self, idx):
    # Get image path
    image_path = self.df.iloc[idx]['path']

    # Load image
    image = PIL.Image.open(image_path)

    # Change pixel values to float
    image = image.convert('RGB')

    # Normalize image
    # image = torchvision.transforms.ToTensor()(image)

    # Get all the demographic columns
    demographic_columns = ['localization_abdomen',
        'localization_acral', 'localization_back', 'localization_chest',
        'localization_ear', 'localization_face', 'localization_foot',
        'localization_genital', 'localization_hand',
        'localization_lower extremity', 'localization_neck',
        'localization_scalp', 'localization_trunk', 'localization_unknown',
        'localization_upper extremity', 'sex_female', 'sex_male',
        'sex_unknown']

    # Get demographic vector
    demographic_vector = self.df.iloc[idx][demographic_columns].values
    demographic_vector = demographic_vector.astype('float32')
    demographic_vector = torch.tensor(demographic_vector)

    # Get Label
```

```

label = self.df.iloc[idx]['dx']

if self.transform:
    image = self.transform(image)

return image,demographic_vector, label

```

## We also need to change the NN model to accept the demographic data

```

In [23]: from torch.nn.functional import relu
class ComplexConvNNBonus(nn.Module):
    def __init__(self):
        super(ComplexConvNNBonus, self).__init__()
        self.maxpool = nn.MaxPool2d(2)
        self.flatten = nn.Flatten()
        # 100 x 125 x 3
        self.conv1 = nn.Conv2d(3, 32, 3) # 98 x 123 x 32
        self.batchnorm1 = nn.BatchNorm2d(32)
        # MaxPool 2x2 49 x 61 x 32
        self.conv2 = nn.Conv2d(32, 64, 3) # 47 x 59 x 64
        self.batchnorm2 = nn.BatchNorm2d(64)
        # MaxPool 2x2 23 x 29 x 64
        self.conv3 = nn.Conv2d(64, 128, 3) # 21 x 27 x 128
        self.batchnorm3 = nn.BatchNorm2d(128)
        # MaxPool 2x2 10 x 13 x 128
        self.conv4 = nn.Conv2d(128, 256, 3) # 8 x 11 x 256
        self.batchnorm4 = nn.BatchNorm2d(256)
        # MaxPool 2x2 4 x 5 x 256
        self.conv5 = nn.Conv2d(256, 512, 3) # 2 x 3 x 512
        self.batchnorm5 = nn.BatchNorm2d(512)
        # MaxPool 2x2 1 x 1 x 512
        self.fc = nn.Linear(512 + 128, 7)

        # Demographic data
        self.fc_demographic = nn.Linear(18,128)

    def forward(self, x,p):
        x = self.maxpool(self.batchnorm1(relu(self.conv1(x))))

```

```

x = self.maxpool(self.batchnorm2(relu(self.conv2(x))))
x = self.maxpool(self.batchnorm3(relu(self.conv3(x))))
x = self.maxpool(self.batchnorm4(relu(self.conv4(x))))
x = self.maxpool(self.batchnorm5(relu(self.conv5(x))))
x = nn.functional.adaptive_avg_pool2d(x,(1,1))
x = self.flatten(x)
# Demographic data
p = relu(self.fc_demographic(p))
# Concatenate the two vectors, the image vector and the demographic vector
x = torch.cat((x,p),1)
x = self.fc(x)
return x

```

Finally, to train,validate and test the model, we need to change the train\_net and test\_net functions to accept the demographic data

```

In [24]: def train_bonus_net(model: ComplexConvNNBonus, trainloader: DataLoader, valloader: DataLoader = None, epochs: int = 10,
            optimizer: optim.Optimizer = None, loss: nn.modules.loss = None, device: str = 'cpu',
            print_period: int = 10) -> None:
    global val_loss, train_loss, val_accuracy, train_accuracy
    val_loss = []
    train_loss = []
    val_accuracy = []
    train_accuracy = []
    for epoch in range(epochs):
        total = 0
        correct = 0
        running_loss = 0.0
        # The only difference between this function and the train_net function is
        # that we need to also get the demographic vector (p) from the dataloader
        for batch, (X, p, y) in enumerate(trainloader, 0):
            model.train()
            X = X.to(device)
            y = y.to(device)
            p = p.to(device)

            optimizer.zero_grad()

            # We need to pass both the image and the demographic vector to the model
            y_pred = model(X, p)

```

```

current_loss = loss(y_pred, y.long())

current_loss.backward()

optimizer.step()

total += y.size(0)
correct += (y_pred.argmax(1) == y).type(torch.float).sum().item()
running_loss += current_loss.item()
if batch % print_period == print_period - 1:
    avg_loss = running_loss / print_period
    avg_accuracy = correct / total * 100
    train_loss.append(avg_loss)
    train_accuracy.append(avg_accuracy)

total_val = 0
correct_val = 0
running_loss_val = 0.0

if valloader is not None:
    model.eval()
    with torch.no_grad():
        # Also need to get the demographic vector (p) from the dataloader here
        for (X_val, p_val, y_val) in valloader:
            X_val = X_val.to(device)
            y_val = y_val.to(device)
            p_val = p_val.to(device)

            y_pred_val = model(X_val, p_val)
            running_loss_val += loss(y_pred_val, y_val.long()).item()
            total_val += y_val.size(0)
            correct_val += (y_pred_val.argmax(1) == y_val).type(torch.float).sum().item()

    avg_loss_val = running_loss_val / len(valloader)
    avg_accuracy_val = correct_val / total_val * 100
    val_loss.append(avg_loss_val)
    val_accuracy.append(avg_accuracy_val)
    # writer.add_scalar('Loss/train', avg_loss, (epoch+1) * len(trainloader) + batch)

print(
    f'[{epoch + 1}, {batch + 1}] loss: {avg_loss} | accuracy: {avg_accuracy:.2f}% | '
    f' val_loss: {avg_loss_val} | val_accuracy: {avg_accuracy_val:.2f}%')

```

```

else:
    print(f'[{epoch + 1}, {batch + 1}] loss: {avg_loss} | accuracy: {avg_accuracy :.4f}%')
    running_loss = 0.0
    total = 0
    correct = 0

```

```

In [25]: def test_bonus_net(model: ComplexConvNNBonus, testloader: DataLoader, number_to_label_array: list[str],
        loss: nn.modules.loss = None, device: str = 'cpu',
        model_name="") -> None:

    model.eval()
    correct = 0
    loss_test = 0.0
    y_trues = []
    y_preds = []
    with torch.no_grad():
        # Again, the only difference between this function and the test_net function is
        # that we need to also get the demographic vector (p) from the dataloader
        for (X, p, y) in testloader:
            X = X.to(device)
            y = y.to(device)
            p = p.to(device)

            # And pass both the image and the demographic vector to the model
            y_pred_test = model(X, p)
            y_trues.extend(y.cpu().numpy())
            y_preds.extend(y_pred_test.argmax(1).cpu().numpy())
            loss_test += loss(y_pred_test, y.long())
            correct += (y_pred_test.argmax(1) == y).type(torch.float).sum().item()

    print(f"Test accuracy: {correct / len(testloader.dataset) * 100:.2f}% | "
          f" Test loss: {loss_test / len(testloader)}")
    conf_matrix = confusion_matrix(y_trues, y_preds)
    df_cm = pd.DataFrame(conf_matrix, index=number_to_label_array,
                        columns=number_to_label_array)

    # Plot the confusion matrix
    sns.heatmap(df_cm, annot=True, cmap='Blues', fmt='g')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title('Confusion Matrix')
    plt.savefig(f'results/confusion_matrices/{model_name}_confusion_matrix.png')
    plt.show()

```

# Now let's train the model and test it

as seen from the code above we are going to use the same way we did for the Complex model

```
In [26]: # Creating image side sizes
M = 100
N = 125
# Creating transforms
transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((M, N)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
# Creating dataset
dataset = MLProject2DatasetBonus(data_dir='data', metadata_fname='metadata.csv', transform=transforms)
# Splitting dataset into train, validation and test
train_dataset, val_dataset, test_dataset = torch.utils.data.random_split(dataset, [.6, .1, .3],
                                                                           generator=torch.Generator().manual_seed(
                                                                               42))

# Creating dataloaders
BATCH_SIZE = 32
trainloader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
valloader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=True)
testloader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=True)

# Creating model
model = ComplexConvNNBonus()
model = model.to(device)
# I couldn't get the summary to work for two input models
# summary(model, ((3, 100, 125), 18))

# Creating optimizer
optimizer = optim.Adam(model.parameters(), lr=1e-3)
# Creating loss function
loss = nn.CrossEntropyLoss()
# Training model
train_bonus_net(model, trainloader, valloader, epochs=20, optimizer=optimizer, loss=loss, device=device)
# This cell took 1 hour and 10 minutes to run on my machine
```



[1, 10]	loss: 1.600027358531952	accuracy: 53.44%	val_loss: 1.5364666320383549	val_accuracy: 64.64%
[1, 20]	loss: 1.0599934697151183	accuracy: 64.38%	val_loss: 1.3302472680807114	val_accuracy: 65.13%
[1, 30]	loss: 0.9950246572494507	accuracy: 66.88%	val_loss: 1.2394726481288671	val_accuracy: 65.13%
[1, 40]	loss: 0.988733834028244	accuracy: 64.69%	val_loss: 1.2943845260888338	val_accuracy: 65.33%
[1, 50]	loss: 0.8500926554203033	accuracy: 69.38%	val_loss: 1.0078584756702185	val_accuracy: 67.43%
[1, 60]	loss: 0.8730084300041199	accuracy: 68.75%	val_loss: 0.9326594062149525	val_accuracy: 65.03%
[1, 70]	loss: 0.8652713000774384	accuracy: 67.19%	val_loss: 0.8426277991384268	val_accuracy: 69.03%
[1, 80]	loss: 0.8383728981018066	accuracy: 70.94%	val_loss: 0.8461446650326252	val_accuracy: 69.83%
[1, 90]	loss: 0.809982979297638	accuracy: 68.44%	val_loss: 0.8949098861776292	val_accuracy: 69.23%
[1, 100]	loss: 0.873773169517517	accuracy: 69.06%	val_loss: 0.7640560418367386	val_accuracy: 69.83%
[1, 110]	loss: 0.8482986629009247	accuracy: 68.75%	val_loss: 0.8752862829715014	val_accuracy: 69.63%
[1, 120]	loss: 0.8278371155261993	accuracy: 70.00%	val_loss: 0.8086240328848362	val_accuracy: 70.43%
[1, 130]	loss: 0.9079543948173523	accuracy: 68.75%	val_loss: 0.883495606482029	val_accuracy: 67.33%
[1, 140]	loss: 0.7375246763229371	accuracy: 73.12%	val_loss: 0.8481611255556345	val_accuracy: 69.43%
[1, 150]	loss: 0.8354595303535461	accuracy: 69.69%	val_loss: 0.8600427834317088	val_accuracy: 67.83%
[1, 160]	loss: 0.8211566209793091	accuracy: 69.38%	val_loss: 0.8079473879188299	val_accuracy: 68.13%
[1, 170]	loss: 0.9094484210014343	accuracy: 66.25%	val_loss: 0.8698827885091305	val_accuracy: 66.63%
[1, 180]	loss: 0.7958847463130951	accuracy: 73.44%	val_loss: 0.8812715185340494	val_accuracy: 70.13%
[2, 10]	loss: 0.7058807641267777	accuracy: 73.75%	val_loss: 0.7899825861677527	val_accuracy: 68.33%
[2, 20]	loss: 0.8089695274829865	accuracy: 65.62%	val_loss: 0.7424014080315828	val_accuracy: 72.33%
[2, 30]	loss: 0.682023686170578	accuracy: 75.31%	val_loss: 0.7450373386964202	val_accuracy: 72.23%
[2, 40]	loss: 0.5885782569646836	accuracy: 78.44%	val_loss: 0.7091180719435215	val_accuracy: 74.53%
[2, 50]	loss: 0.736363884806633	accuracy: 73.75%	val_loss: 0.7450867993757129	val_accuracy: 74.53%
[2, 60]	loss: 0.6782770276069641	accuracy: 73.44%	val_loss: 0.7738254982978106	val_accuracy: 72.23%
[2, 70]	loss: 0.691400408744812	accuracy: 75.62%	val_loss: 0.7647241540253162	val_accuracy: 71.33%
[2, 80]	loss: 0.7422822713851929	accuracy: 72.19%	val_loss: 0.827694064937532	val_accuracy: 67.13%
[2, 90]	loss: 0.7414175808429718	accuracy: 73.75%	val_loss: 0.7356111742556095	val_accuracy: 73.33%
[2, 100]	loss: 0.7690630823373794	accuracy: 71.88%	val_loss: 0.7098126262426376	val_accuracy: 71.43%
[2, 110]	loss: 0.6506730079650879	accuracy: 76.88%	val_loss: 0.7227714918553829	val_accuracy: 72.53%
[2, 120]	loss: 0.5850475937128067	accuracy: 77.19%	val_loss: 0.7792424410581589	val_accuracy: 72.03%
[2, 130]	loss: 0.7382927000522613	accuracy: 71.88%	val_loss: 0.7593212053179741	val_accuracy: 71.63%
[2, 140]	loss: 0.739416629076004	accuracy: 73.12%	val_loss: 0.7597107272595167	val_accuracy: 70.63%
[2, 150]	loss: 0.6110884487628937	accuracy: 77.81%	val_loss: 0.8123133331537247	val_accuracy: 69.53%
[2, 160]	loss: 0.719591674208641	accuracy: 73.12%	val_loss: 0.751898841932416	val_accuracy: 71.43%
[2, 170]	loss: 0.796382087469101	accuracy: 74.06%	val_loss: 0.7405819613486528	val_accuracy: 72.73%
[2, 180]	loss: 0.7750088989734649	accuracy: 70.00%	val_loss: 0.734442587941885	val_accuracy: 71.83%
[3, 10]	loss: 0.6434117197990418	accuracy: 75.31%	val_loss: 0.7139282049611211	val_accuracy: 73.43%
[3, 20]	loss: 0.6765142858028412	accuracy: 72.81%	val_loss: 0.7572554582729936	val_accuracy: 72.43%
[3, 30]	loss: 0.6426381707191468	accuracy: 75.31%	val_loss: 0.7310755904763937	val_accuracy: 71.63%
[3, 40]	loss: 0.6699868679046631	accuracy: 72.19%	val_loss: 0.7127618528902531	val_accuracy: 73.13%
[3, 50]	loss: 0.6655596762895584	accuracy: 74.38%	val_loss: 0.6620474574156106	val_accuracy: 75.22%
[3, 60]	loss: 0.5848332494497299	accuracy: 77.19%	val_loss: 0.6959523484110832	val_accuracy: 74.03%

[3, 70]	loss: 0.6740255445241928	accuracy: 75.31%	val_loss: 0.643298814073205	val_accuracy: 75.32%
[3, 80]	loss: 0.5798832803964615	accuracy: 79.69%	val_loss: 0.6884009810164571	val_accuracy: 75.22%
[3, 90]	loss: 0.7163134634494781	accuracy: 74.06%	val_loss: 0.7448704056441784	val_accuracy: 71.93%
[3, 100]	loss: 0.6940701305866241	accuracy: 74.38%	val_loss: 0.7437044465914369	val_accuracy: 73.23%
[3, 110]	loss: 0.7841456353664398	accuracy: 72.19%	val_loss: 0.7799679394811392	val_accuracy: 68.93%
[3, 120]	loss: 0.6503895044326782	accuracy: 75.62%	val_loss: 0.797253024764359	val_accuracy: 72.53%
[3, 130]	loss: 0.6673334240913391	accuracy: 75.31%	val_loss: 0.716179646551609	val_accuracy: 73.33%
[3, 140]	loss: 0.6561056464910507	accuracy: 75.00%	val_loss: 0.7386818761005998	val_accuracy: 72.63%
[3, 150]	loss: 0.7767239540815354	accuracy: 70.94%	val_loss: 0.7198753887787461	val_accuracy: 73.13%
[3, 160]	loss: 0.6674642264842987	accuracy: 74.38%	val_loss: 0.6973521206527948	val_accuracy: 74.23%
[3, 170]	loss: 0.6097018480300903	accuracy: 75.62%	val_loss: 0.7555998992174864	val_accuracy: 70.93%
[3, 180]	loss: 0.6086151644587516	accuracy: 78.44%	val_loss: 0.6623652651906013	val_accuracy: 75.72%
[4, 10]	loss: 0.5753791600465774	accuracy: 80.62%	val_loss: 0.6591367209330201	val_accuracy: 75.92%
[4, 20]	loss: 0.6099459499120712	accuracy: 78.12%	val_loss: 0.6661296607926488	val_accuracy: 75.22%
[4, 30]	loss: 0.5777314335107804	accuracy: 77.50%	val_loss: 0.7194696199148893	val_accuracy: 72.33%
[4, 40]	loss: 0.5498325109481812	accuracy: 78.75%	val_loss: 0.6708477037027478	val_accuracy: 73.83%
[4, 50]	loss: 0.5657051205635071	accuracy: 77.81%	val_loss: 0.7175300316885114	val_accuracy: 74.43%
[4, 60]	loss: 0.6385161697864532	accuracy: 74.38%	val_loss: 0.696208999492228	val_accuracy: 73.53%
[4, 70]	loss: 0.6772415190935135	accuracy: 74.38%	val_loss: 0.7317825211212039	val_accuracy: 75.12%
[4, 80]	loss: 0.7162796437740326	accuracy: 71.88%	val_loss: 0.6911997087299824	val_accuracy: 73.23%
[4, 90]	loss: 0.6505649328231812	accuracy: 73.75%	val_loss: 0.7098056497052312	val_accuracy: 74.43%
[4, 100]	loss: 0.5810024708509445	accuracy: 77.81%	val_loss: 0.6591903353109956	val_accuracy: 75.32%
[4, 110]	loss: 0.6082124322652817	accuracy: 78.44%	val_loss: 0.731131729669869	val_accuracy: 73.83%
[4, 120]	loss: 0.6264097422361374	accuracy: 75.31%	val_loss: 0.7414591135457158	val_accuracy: 74.03%
[4, 130]	loss: 0.5761011600494385	accuracy: 79.38%	val_loss: 0.6744047403335571	val_accuracy: 75.32%
[4, 140]	loss: 0.5580457240343094	accuracy: 79.69%	val_loss: 0.6833520922809839	val_accuracy: 73.73%
[4, 150]	loss: 0.6898782074451446	accuracy: 75.62%	val_loss: 0.6669758642092347	val_accuracy: 74.53%
[4, 160]	loss: 0.6591581881046296	accuracy: 77.50%	val_loss: 0.6659227143973112	val_accuracy: 75.22%
[4, 170]	loss: 0.546792083978653	accuracy: 81.25%	val_loss: 0.6713760774582624	val_accuracy: 74.63%
[4, 180]	loss: 0.614773890376091	accuracy: 77.81%	val_loss: 0.6947182435542345	val_accuracy: 74.43%
[5, 10]	loss: 0.5512418001890182	accuracy: 78.12%	val_loss: 0.6076870001852512	val_accuracy: 77.62%
[5, 20]	loss: 0.5574635863304138	accuracy: 77.50%	val_loss: 0.6532285641878843	val_accuracy: 75.32%
[5, 30]	loss: 0.4658811539411545	accuracy: 83.44%	val_loss: 0.6404229854233563	val_accuracy: 75.82%
[5, 40]	loss: 0.5805713891983032	accuracy: 80.31%	val_loss: 0.6636958252638578	val_accuracy: 74.63%
[5, 50]	loss: 0.581307002902031	accuracy: 74.38%	val_loss: 0.6865639686584473	val_accuracy: 73.93%
[5, 60]	loss: 0.6404704093933106	accuracy: 77.50%	val_loss: 0.7215172359719872	val_accuracy: 73.73%
[5, 70]	loss: 0.5797306835651398	accuracy: 76.88%	val_loss: 0.6679781693965197	val_accuracy: 73.73%
[5, 80]	loss: 0.5098635196685791	accuracy: 81.56%	val_loss: 0.7056260667741299	val_accuracy: 74.73%
[5, 90]	loss: 0.5547813564538956	accuracy: 77.81%	val_loss: 0.6798406410962343	val_accuracy: 73.03%
[5, 100]	loss: 0.539436075091362	accuracy: 79.69%	val_loss: 0.6696066725999117	val_accuracy: 74.53%
[5, 110]	loss: 0.6093660026788712	accuracy: 78.12%	val_loss: 0.6645755730569363	val_accuracy: 76.72%
[5, 120]	loss: 0.5970535039901733	accuracy: 77.81%	val_loss: 0.656596701592207	val_accuracy: 75.92%

[5, 130]	loss: 0.5570680409669876	accuracy: 78.12%	val_loss: 0.6277870913036168	val_accuracy: 77.62%
[5, 140]	loss: 0.46322486251592637	accuracy: 84.06%	val_loss: 0.7099160477519035	val_accuracy: 73.83%
[5, 150]	loss: 0.5345730811357499	accuracy: 81.88%	val_loss: 0.6559119783341885	val_accuracy: 75.52%
[5, 160]	loss: 0.5392092436552047	accuracy: 80.00%	val_loss: 0.7124760719016194	val_accuracy: 74.23%
[5, 170]	loss: 0.7558134436607361	accuracy: 73.75%	val_loss: 0.7740396652370691	val_accuracy: 71.03%
[5, 180]	loss: 0.8474277496337891	accuracy: 68.75%	val_loss: 0.8452269993722439	val_accuracy: 70.63%
[6, 10]	loss: 0.7049004286527634	accuracy: 72.81%	val_loss: 0.7658079080283642	val_accuracy: 71.13%
[6, 20]	loss: 0.5711145013570785	accuracy: 78.75%	val_loss: 0.7219130583107471	val_accuracy: 71.33%
[6, 30]	loss: 0.5803933620452881	accuracy: 79.38%	val_loss: 0.7023132117465138	val_accuracy: 73.43%
[6, 40]	loss: 0.6136822134256363	accuracy: 76.88%	val_loss: 0.7027030000463128	val_accuracy: 73.03%
[6, 50]	loss: 0.5750079452991486	accuracy: 78.75%	val_loss: 0.6608401760458946	val_accuracy: 75.12%
[6, 60]	loss: 0.6049267768859863	accuracy: 79.06%	val_loss: 0.7054161047562957	val_accuracy: 75.02%
[6, 70]	loss: 0.560998198390007	accuracy: 76.88%	val_loss: 0.6648150077089667	val_accuracy: 75.52%
[6, 80]	loss: 0.6063017427921296	accuracy: 77.19%	val_loss: 0.705652670469135	val_accuracy: 73.23%
[6, 90]	loss: 0.4886783480644226	accuracy: 83.44%	val_loss: 0.629141004756093	val_accuracy: 75.32%
[6, 100]	loss: 0.573975682258606	accuracy: 79.69%	val_loss: 0.632698018103838	val_accuracy: 75.92%
[6, 110]	loss: 0.537535086274147	accuracy: 77.81%	val_loss: 0.6266384115442634	val_accuracy: 76.42%
[6, 120]	loss: 0.5222141414880752	accuracy: 79.06%	val_loss: 0.7569740293547511	val_accuracy: 73.33%
[6, 130]	loss: 0.5627941966056824	accuracy: 78.12%	val_loss: 0.6792449317872524	val_accuracy: 74.83%
[6, 140]	loss: 0.6020899415016174	accuracy: 75.62%	val_loss: 0.7952695833519101	val_accuracy: 73.53%
[6, 150]	loss: 0.5406428843736648	accuracy: 81.88%	val_loss: 0.6822060281410813	val_accuracy: 74.83%
[6, 160]	loss: 0.5975734114646911	accuracy: 76.25%	val_loss: 0.6758038504049182	val_accuracy: 73.53%
[6, 170]	loss: 0.628984010219574	accuracy: 77.19%	val_loss: 0.7333106929436326	val_accuracy: 71.83%
[6, 180]	loss: 0.5085981786251068	accuracy: 80.62%	val_loss: 0.7177181011065841	val_accuracy: 76.62%
[7, 10]	loss: 0.4698475360870361	accuracy: 81.88%	val_loss: 0.6623575333505869	val_accuracy: 75.82%
[7, 20]	loss: 0.47410196661949155	accuracy: 81.56%	val_loss: 0.6870648842304945	val_accuracy: 74.43%
[7, 30]	loss: 0.525987908244133	accuracy: 80.00%	val_loss: 0.6395727642811835	val_accuracy: 76.42%
[7, 40]	loss: 0.4238841950893402	accuracy: 84.69%	val_loss: 0.6521358881145716	val_accuracy: 76.02%
[7, 50]	loss: 0.36750797033309934	accuracy: 88.12%	val_loss: 0.6782672135159373	val_accuracy: 76.02%
[7, 60]	loss: 0.46455671191215514	accuracy: 83.44%	val_loss: 0.6648417022079229	val_accuracy: 74.53%
[7, 70]	loss: 0.476781764626503	accuracy: 83.12%	val_loss: 0.6851651407778263	val_accuracy: 74.53%
[7, 80]	loss: 0.4675774693489075	accuracy: 82.19%	val_loss: 0.7294744923710823	val_accuracy: 75.02%
[7, 90]	loss: 0.5176719903945923	accuracy: 78.12%	val_loss: 0.7227505864575505	val_accuracy: 72.93%
[7, 100]	loss: 0.40918833762407303	accuracy: 82.81%	val_loss: 0.7678159587085247	val_accuracy: 75.12%
[7, 110]	loss: 0.5342970699071884	accuracy: 80.31%	val_loss: 0.6671551773324609	val_accuracy: 75.22%
[7, 120]	loss: 0.5337684422731399	accuracy: 78.44%	val_loss: 0.6306691793724895	val_accuracy: 75.42%
[7, 130]	loss: 0.5160357832908631	accuracy: 82.19%	val_loss: 0.6964177442714572	val_accuracy: 75.92%
[7, 140]	loss: 0.5079063445329666	accuracy: 81.56%	val_loss: 0.6737385168671608	val_accuracy: 75.02%
[7, 150]	loss: 0.4522669970989227	accuracy: 82.81%	val_loss: 0.7518010744825006	val_accuracy: 74.63%
[7, 160]	loss: 0.5126410275697708	accuracy: 80.62%	val_loss: 0.6659566452726722	val_accuracy: 76.82%
[7, 170]	loss: 0.5384398549795151	accuracy: 80.62%	val_loss: 0.6552640157751739	val_accuracy: 75.72%
[7, 180]	loss: 0.44321165680885316	accuracy: 81.88%	val_loss: 0.6691079363226891	val_accuracy: 74.83%

[8, 10]	loss: 0.4196830540895462	accuracy: 85.00%	val_loss: 0.6271042944863439	val_accuracy: 75.92%
[8, 20]	loss: 0.3618864774703979	accuracy: 87.19%	val_loss: 0.678341337479651	val_accuracy: 76.22%
[8, 30]	loss: 0.4094381779432297	accuracy: 84.38%	val_loss: 0.6710457680746913	val_accuracy: 76.12%
[8, 40]	loss: 0.3964301273226738	accuracy: 85.00%	val_loss: 0.6726554149063304	val_accuracy: 76.52%
[8, 50]	loss: 0.44439859986305236	accuracy: 84.06%	val_loss: 0.6356390006840229	val_accuracy: 76.82%
[8, 60]	loss: 0.39001850187778475	accuracy: 84.38%	val_loss: 0.6728248903527856	val_accuracy: 76.92%
[8, 70]	loss: 0.35536017939448356	accuracy: 85.62%	val_loss: 0.6414234051480889	val_accuracy: 78.32%
[8, 80]	loss: 0.3863085836172104	accuracy: 84.69%	val_loss: 0.7230667434632778	val_accuracy: 74.63%
[8, 90]	loss: 0.44836956858634947	accuracy: 82.19%	val_loss: 0.7686876077204943	val_accuracy: 75.72%
[8, 100]	loss: 0.39942285418510437	accuracy: 85.00%	val_loss: 0.6650196835398674	val_accuracy: 77.42%
[8, 110]	loss: 0.37210783958435056	accuracy: 85.00%	val_loss: 0.7296867137774825	val_accuracy: 76.02%
[8, 120]	loss: 0.40246464386582376	accuracy: 85.00%	val_loss: 0.6499871769919991	val_accuracy: 77.32%
[8, 130]	loss: 0.44884715527296065	accuracy: 80.62%	val_loss: 0.7273526852950454	val_accuracy: 76.02%
[8, 140]	loss: 0.50336854159832	accuracy: 79.06%	val_loss: 0.7018412481993437	val_accuracy: 75.82%
[8, 150]	loss: 0.563972195982933	accuracy: 76.88%	val_loss: 0.7847473900765181	val_accuracy: 72.23%
[8, 160]	loss: 0.5599155187606811	accuracy: 78.75%	val_loss: 0.7441395176574588	val_accuracy: 75.62%
[8, 170]	loss: 0.5767714649438858	accuracy: 77.50%	val_loss: 0.7738063391298056	val_accuracy: 74.03%
[8, 180]	loss: 0.5017509669065475	accuracy: 80.62%	val_loss: 0.7064375756308436	val_accuracy: 76.02%
[9, 10]	loss: 0.43977062255144117	accuracy: 87.50%	val_loss: 0.6722211260348558	val_accuracy: 74.93%
[9, 20]	loss: 0.37991843521595003	accuracy: 87.19%	val_loss: 0.6723858956247568	val_accuracy: 77.02%
[9, 30]	loss: 0.3472294881939888	accuracy: 86.88%	val_loss: 0.7013766746968031	val_accuracy: 74.23%
[9, 40]	loss: 0.39184014201164247	accuracy: 87.50%	val_loss: 0.6683580875396729	val_accuracy: 75.02%
[9, 50]	loss: 0.34109464585781096	accuracy: 89.06%	val_loss: 0.7126669371500611	val_accuracy: 76.22%
[9, 60]	loss: 0.3965762183070183	accuracy: 84.69%	val_loss: 0.6923612430691719	val_accuracy: 74.73%
[9, 70]	loss: 0.345405350625515	accuracy: 88.75%	val_loss: 0.6846401486545801	val_accuracy: 77.42%
[9, 80]	loss: 0.34689733684062957	accuracy: 89.06%	val_loss: 0.6771679576486349	val_accuracy: 76.62%
[9, 90]	loss: 0.42212819755077363	accuracy: 85.00%	val_loss: 0.77456527762115	val_accuracy: 73.63%
[9, 100]	loss: 0.41223659217357633	accuracy: 85.62%	val_loss: 0.684301596134901	val_accuracy: 75.02%
[9, 110]	loss: 0.4135635569691658	accuracy: 84.38%	val_loss: 0.733644743449986	val_accuracy: 77.12%
[9, 120]	loss: 0.4589883297681808	accuracy: 80.94%	val_loss: 0.6411635018885136	val_accuracy: 78.32%
[9, 130]	loss: 0.4231394648551941	accuracy: 83.12%	val_loss: 0.6586901163682342	val_accuracy: 76.82%
[9, 140]	loss: 0.35671835392713547	accuracy: 86.25%	val_loss: 0.6682518608868122	val_accuracy: 76.32%
[9, 150]	loss: 0.40543562471866607	accuracy: 82.81%	val_loss: 0.6657161023467779	val_accuracy: 76.02%
[9, 160]	loss: 0.44502412974834443	accuracy: 84.38%	val_loss: 0.6779968859627843	val_accuracy: 76.32%
[9, 170]	loss: 0.4723042041063309	accuracy: 82.19%	val_loss: 0.7123349439352751	val_accuracy: 76.92%
[9, 180]	loss: 0.3587679088115692	accuracy: 86.25%	val_loss: 0.7267803568392992	val_accuracy: 76.42%
[10, 10]	loss: 0.2652467742562294	accuracy: 90.94%	val_loss: 0.7651275275275111	val_accuracy: 76.72%
[10, 20]	loss: 0.3085550218820572	accuracy: 89.06%	val_loss: 0.6819978971034288	val_accuracy: 77.32%
[10, 30]	loss: 0.26298754662275314	accuracy: 91.25%	val_loss: 0.6771595245227218	val_accuracy: 76.92%
[10, 40]	loss: 0.2864482387900352	accuracy: 91.56%	val_loss: 0.7023721495643258	val_accuracy: 77.92%
[10, 50]	loss: 0.2894578091800213	accuracy: 89.69%	val_loss: 0.6528916535899043	val_accuracy: 76.72%
[10, 60]	loss: 0.29932766407728195	accuracy: 86.88%	val_loss: 0.7003838075324893	val_accuracy: 76.52%

[10, 70]	loss: 0.28687879741191863	accuracy: 90.94%	val_loss: 0.6820931085385382	val_accuracy: 76.52%
[10, 80]	loss: 0.20995983853936195	accuracy: 91.25%	val_loss: 0.7022450808435678	val_accuracy: 75.72%
[10, 90]	loss: 0.3049706771969795	accuracy: 89.38%	val_loss: 0.7323399372398853	val_accuracy: 76.52%
[10, 100]	loss: 0.28534088730812074	accuracy: 88.75%	val_loss: 0.743275043554604	val_accuracy: 75.42%
[10, 110]	loss: 0.42470761239528654	accuracy: 85.00%	val_loss: 0.755277271848172	val_accuracy: 74.53%
[10, 120]	loss: 0.2628707975149155	accuracy: 89.69%	val_loss: 0.7454596841707826	val_accuracy: 75.22%
[10, 130]	loss: 0.29969759434461596	accuracy: 88.44%	val_loss: 0.7553270636126399	val_accuracy: 76.72%
[10, 140]	loss: 0.339595702290535	accuracy: 87.81%	val_loss: 0.7095815455541015	val_accuracy: 75.22%
[10, 150]	loss: 0.2565939098596573	accuracy: 90.94%	val_loss: 0.7399863516911864	val_accuracy: 76.72%
[10, 160]	loss: 0.3255498006939888	accuracy: 86.56%	val_loss: 0.6923547028563917	val_accuracy: 76.52%
[10, 170]	loss: 0.33107619136571886	accuracy: 88.12%	val_loss: 0.7063004714436829	val_accuracy: 76.52%
[10, 180]	loss: 0.37328106313943865	accuracy: 86.25%	val_loss: 0.7244191849604249	val_accuracy: 77.72%
[11, 10]	loss: 0.23699921518564224	accuracy: 92.19%	val_loss: 0.6994529254734516	val_accuracy: 76.22%
[11, 20]	loss: 0.19321148693561555	accuracy: 94.69%	val_loss: 0.6725504514761269	val_accuracy: 76.92%
[11, 30]	loss: 0.17280559837818146	accuracy: 94.69%	val_loss: 0.7016802290454507	val_accuracy: 77.42%
[11, 40]	loss: 0.19073916971683502	accuracy: 94.06%	val_loss: 0.6837505605071783	val_accuracy: 77.52%
[11, 50]	loss: 0.1732227656841277	accuracy: 94.06%	val_loss: 0.7281720843166113	val_accuracy: 77.82%
[11, 60]	loss: 0.20086867287755011	accuracy: 93.75%	val_loss: 0.8141607344150543	val_accuracy: 76.92%
[11, 70]	loss: 0.18132874891161918	accuracy: 93.75%	val_loss: 0.6769054336473346	val_accuracy: 77.92%
[11, 80]	loss: 0.18260411620140077	accuracy: 94.06%	val_loss: 0.7786827320232987	val_accuracy: 77.22%
[11, 90]	loss: 0.1683257095515728	accuracy: 95.00%	val_loss: 0.7892782250419259	val_accuracy: 77.52%
[11, 100]	loss: 0.21508159041404723	accuracy: 92.81%	val_loss: 0.8135646088048816	val_accuracy: 74.03%
[11, 110]	loss: 0.2414264515042305	accuracy: 91.25%	val_loss: 0.7246375661343336	val_accuracy: 77.12%
[11, 120]	loss: 0.23812475875020028	accuracy: 91.56%	val_loss: 0.7816328657791018	val_accuracy: 77.42%
[11, 130]	loss: 0.22823114097118377	accuracy: 91.56%	val_loss: 0.8199614631012082	val_accuracy: 76.82%
[11, 140]	loss: 0.2579861059784889	accuracy: 89.38%	val_loss: 0.7780319773592055	val_accuracy: 76.42%
[11, 150]	loss: 0.19401577785611152	accuracy: 93.12%	val_loss: 0.8049830263480544	val_accuracy: 76.42%
[11, 160]	loss: 0.28488488495349884	accuracy: 90.94%	val_loss: 0.7683712933212519	val_accuracy: 75.62%
[11, 170]	loss: 0.2517699547111988	accuracy: 90.62%	val_loss: 0.8790919883176684	val_accuracy: 75.42%
[11, 180]	loss: 0.30687594711780547	accuracy: 90.62%	val_loss: 0.8386523788794875	val_accuracy: 76.12%
[12, 10]	loss: 0.20538760498166084	accuracy: 91.88%	val_loss: 0.8313548341393471	val_accuracy: 75.12%
[12, 20]	loss: 0.20682285353541374	accuracy: 93.12%	val_loss: 0.8037257669493556	val_accuracy: 74.83%
[12, 30]	loss: 0.16413215845823287	accuracy: 94.38%	val_loss: 0.8888843050226569	val_accuracy: 75.12%
[12, 40]	loss: 0.1717342808842659	accuracy: 94.38%	val_loss: 0.8134537218138576	val_accuracy: 76.12%
[12, 50]	loss: 0.2200840160250664	accuracy: 90.94%	val_loss: 0.8182730153203011	val_accuracy: 76.82%
[12, 60]	loss: 0.1734132543206215	accuracy: 94.38%	val_loss: 0.8216928420588374	val_accuracy: 77.12%
[12, 70]	loss: 0.16775813177227974	accuracy: 94.69%	val_loss: 0.7985498560592532	val_accuracy: 75.12%
[12, 80]	loss: 0.17223066426813602	accuracy: 94.69%	val_loss: 0.8317371513694525	val_accuracy: 76.72%
[12, 90]	loss: 0.20045031681656839	accuracy: 92.19%	val_loss: 0.7533682556822896	val_accuracy: 76.42%
[12, 100]	loss: 0.13226275499910117	accuracy: 93.44%	val_loss: 0.8042139266617596	val_accuracy: 76.52%
[12, 110]	loss: 0.16592935882508755	accuracy: 92.50%	val_loss: 0.7878779862076044	val_accuracy: 75.82%
[12, 120]	loss: 0.15494932159781455	accuracy: 94.06%	val_loss: 0.8088584383949637	val_accuracy: 77.02%

[12, 130]	loss: 0.21586611047387122	accuracy: 93.12%	val_loss: 0.7606378900818527	val_accuracy: 78.22%
[12, 140]	loss: 0.20897438153624534	accuracy: 90.31%	val_loss: 0.78163778828457	val_accuracy: 78.22%
[12, 150]	loss: 0.2290512040257454	accuracy: 91.88%	val_loss: 0.819371581543237	val_accuracy: 76.12%
[12, 160]	loss: 0.3064619742333889	accuracy: 90.31%	val_loss: 0.8328586360439658	val_accuracy: 74.23%
[12, 170]	loss: 0.24032047465443612	accuracy: 92.50%	val_loss: 0.8869847813621163	val_accuracy: 76.22%
[12, 180]	loss: 0.2212035447359085	accuracy: 93.44%	val_loss: 0.9608374908566475	val_accuracy: 75.52%
[13, 10]	loss: 0.1657975737005472	accuracy: 95.31%	val_loss: 0.8687936258502305	val_accuracy: 75.22%
[13, 20]	loss: 0.1142881765961647	accuracy: 97.19%	val_loss: 0.8778231926262379	val_accuracy: 75.32%
[13, 30]	loss: 0.15053380317986012	accuracy: 94.06%	val_loss: 0.8628737274557352	val_accuracy: 77.32%
[13, 40]	loss: 0.1312802590429783	accuracy: 95.94%	val_loss: 0.851649516262114	val_accuracy: 76.92%
[13, 50]	loss: 0.12966367974877357	accuracy: 95.94%	val_loss: 0.8538257488980889	val_accuracy: 75.72%
[13, 60]	loss: 0.17059290669858457	accuracy: 94.38%	val_loss: 0.8371227793395519	val_accuracy: 76.62%
[13, 70]	loss: 0.20424699932336807	accuracy: 93.44%	val_loss: 0.8920517922379076	val_accuracy: 74.73%
[13, 80]	loss: 0.22964655831456185	accuracy: 91.56%	val_loss: 0.9698551334440708	val_accuracy: 73.13%
[13, 90]	loss: 0.2772284343838692	accuracy: 91.56%	val_loss: 1.0300401648273692	val_accuracy: 73.63%
[13, 100]	loss: 0.320549875497818	accuracy: 89.38%	val_loss: 1.0628969371318817	val_accuracy: 73.73%
[13, 110]	loss: 0.279242941737175	accuracy: 90.31%	val_loss: 0.9148641191422939	val_accuracy: 74.23%
[13, 120]	loss: 0.21283315643668174	accuracy: 93.44%	val_loss: 0.8747545233927667	val_accuracy: 73.13%
[13, 130]	loss: 0.2408073790371418	accuracy: 91.88%	val_loss: 0.8546751998364925	val_accuracy: 74.33%
[13, 140]	loss: 0.23353382125496863	accuracy: 89.38%	val_loss: 0.8975943131372333	val_accuracy: 73.73%
[13, 150]	loss: 0.2260521851480007	accuracy: 91.88%	val_loss: 0.8666637884452939	val_accuracy: 74.53%
[13, 160]	loss: 0.2999465361237526	accuracy: 89.38%	val_loss: 0.8713734513148665	val_accuracy: 75.52%
[13, 170]	loss: 0.22250073850154878	accuracy: 91.88%	val_loss: 0.849951739422977	val_accuracy: 74.93%
[13, 180]	loss: 0.2395026318728924	accuracy: 91.56%	val_loss: 0.808620247989893	val_accuracy: 74.53%
[14, 10]	loss: 0.1387147031724453	accuracy: 94.69%	val_loss: 0.896369218826294	val_accuracy: 76.32%
[14, 20]	loss: 0.18682505339384078	accuracy: 92.50%	val_loss: 0.8351997132413089	val_accuracy: 74.93%
[14, 30]	loss: 0.15346062630414964	accuracy: 96.25%	val_loss: 0.8807376576587558	val_accuracy: 76.82%
[14, 40]	loss: 0.12959166020154952	accuracy: 95.62%	val_loss: 0.9002801412716508	val_accuracy: 75.22%
[14, 50]	loss: 0.1992579497396946	accuracy: 92.19%	val_loss: 0.858563058078289	val_accuracy: 77.12%
[14, 60]	loss: 0.14723162911832333	accuracy: 95.62%	val_loss: 0.849678510800004	val_accuracy: 75.22%
[14, 70]	loss: 0.12563102692365646	accuracy: 97.19%	val_loss: 0.8416464924812317	val_accuracy: 75.02%
[14, 80]	loss: 0.12905420791357755	accuracy: 95.94%	val_loss: 0.9583410825580359	val_accuracy: 75.72%
[14, 90]	loss: 0.11398136485368013	accuracy: 97.19%	val_loss: 0.8649833714589477	val_accuracy: 74.93%
[14, 100]	loss: 0.1214464595541358	accuracy: 97.50%	val_loss: 0.8758415351621807	val_accuracy: 75.02%
[14, 110]	loss: 0.10476127117872239	accuracy: 95.62%	val_loss: 0.8747840784490108	val_accuracy: 77.52%
[14, 120]	loss: 0.09825150780379772	accuracy: 96.88%	val_loss: 0.95148931350559	val_accuracy: 77.62%
[14, 130]	loss: 0.1550653487443924	accuracy: 95.00%	val_loss: 0.8520962786860764	val_accuracy: 77.42%
[14, 140]	loss: 0.15480977054685355	accuracy: 94.69%	val_loss: 0.9800817584618926	val_accuracy: 76.52%
[14, 150]	loss: 0.13967944867908955	accuracy: 96.25%	val_loss: 0.9137758919969201	val_accuracy: 75.42%
[14, 160]	loss: 0.08330148160457611	accuracy: 97.19%	val_loss: 0.8602373311296105	val_accuracy: 77.22%
[14, 170]	loss: 0.17380470596253872	accuracy: 93.75%	val_loss: 0.8336730417795479	val_accuracy: 77.02%
[14, 180]	loss: 0.18031338304281236	accuracy: 93.44%	val_loss: 0.8407296189107001	val_accuracy: 77.82%

[15, 10]	loss: 0.06968428660184145	accuracy: 99.06%	val_loss: 0.8138083266094327	val_accuracy: 76.92%
[15, 20]	loss: 0.04828843008726835	accuracy: 98.75%	val_loss: 0.8360381592065096	val_accuracy: 77.52%
[15, 30]	loss: 0.09433960411697626	accuracy: 97.81%	val_loss: 0.8093052364420146	val_accuracy: 78.22%
[15, 40]	loss: 0.05804901393130422	accuracy: 98.44%	val_loss: 0.9318131108302623	val_accuracy: 78.12%
[15, 50]	loss: 0.06223299130797386	accuracy: 97.50%	val_loss: 0.8772317264229059	val_accuracy: 77.82%
[15, 60]	loss: 0.06388536272570491	accuracy: 97.81%	val_loss: 0.839064225088805	val_accuracy: 76.72%
[15, 70]	loss: 0.062367072235792874	accuracy: 97.50%	val_loss: 0.9185316506773233	val_accuracy: 77.02%
[15, 80]	loss: 0.10290309395641088	accuracy: 96.88%	val_loss: 0.9800427239388227	val_accuracy: 75.22%
[15, 90]	loss: 0.06279950998723507	accuracy: 98.75%	val_loss: 0.9693804122507572	val_accuracy: 77.52%
[15, 100]	loss: 0.0738335620611906	accuracy: 97.19%	val_loss: 0.9911299990490079	val_accuracy: 76.42%
[15, 110]	loss: 0.06199179217219353	accuracy: 97.50%	val_loss: 0.9420087006874382	val_accuracy: 77.02%
[15, 120]	loss: 0.04007026813924312	accuracy: 99.06%	val_loss: 0.9336351584643126	val_accuracy: 76.32%
[15, 130]	loss: 0.05007385211065411	accuracy: 98.75%	val_loss: 0.8773122727870941	val_accuracy: 76.72%
[15, 140]	loss: 0.06392019037157297	accuracy: 97.50%	val_loss: 0.9391215741634369	val_accuracy: 77.42%
[15, 150]	loss: 0.060480866208672525	accuracy: 98.44%	val_loss: 0.908179531339556	val_accuracy: 76.52%
[15, 160]	loss: 0.06580758700147271	accuracy: 97.81%	val_loss: 0.9304434796795249	val_accuracy: 75.52%
[15, 170]	loss: 0.11849150462076068	accuracy: 95.62%	val_loss: 0.9393503395840526	val_accuracy: 75.72%
[15, 180]	loss: 0.09932373594492674	accuracy: 97.19%	val_loss: 0.9363704091520049	val_accuracy: 75.92%
[16, 10]	loss: 0.07658144813030958	accuracy: 97.50%	val_loss: 1.0580135877244174	val_accuracy: 75.62%
[16, 20]	loss: 0.047242914652451874	accuracy: 98.44%	val_loss: 0.986137212254107	val_accuracy: 75.12%
[16, 30]	loss: 0.054080390557646754	accuracy: 98.12%	val_loss: 1.0006640711799264	val_accuracy: 76.12%
[16, 40]	loss: 0.040455558244138955	accuracy: 99.06%	val_loss: 1.0031658420339227	val_accuracy: 76.52%
[16, 50]	loss: 0.04063398689031601	accuracy: 98.12%	val_loss: 0.9571258444339037	val_accuracy: 77.12%
[16, 60]	loss: 0.031700071226805446	accuracy: 99.06%	val_loss: 0.9438794832676649	val_accuracy: 77.82%
[16, 70]	loss: 0.050132475793361664	accuracy: 98.12%	val_loss: 0.9544144524261355	val_accuracy: 78.12%
[16, 80]	loss: 0.03819299442693591	accuracy: 98.75%	val_loss: 0.9203571733087301	val_accuracy: 77.62%
[16, 90]	loss: 0.039338710531592366	accuracy: 99.06%	val_loss: 0.9402273967862129	val_accuracy: 76.32%
[16, 100]	loss: 0.042382915318012235	accuracy: 98.75%	val_loss: 1.045508943784833	val_accuracy: 77.12%
[16, 110]	loss: 0.07298169238492846	accuracy: 97.81%	val_loss: 0.953766711987555	val_accuracy: 75.72%
[16, 120]	loss: 0.048561707325279715	accuracy: 98.75%	val_loss: 0.9943855572491884	val_accuracy: 76.52%
[16, 130]	loss: 0.053991684969514606	accuracy: 98.12%	val_loss: 0.9673583619296551	val_accuracy: 76.42%
[16, 140]	loss: 0.08306734450161457	accuracy: 96.88%	val_loss: 0.9140247164759785	val_accuracy: 76.82%
[16, 150]	loss: 0.053898188192397355	accuracy: 97.81%	val_loss: 0.922791556455195	val_accuracy: 76.72%
[16, 160]	loss: 0.060224310401827096	accuracy: 98.75%	val_loss: 0.9749775305390358	val_accuracy: 76.52%
[16, 170]	loss: 0.07571202209219337	accuracy: 97.19%	val_loss: 1.0525017064064741	val_accuracy: 76.62%
[16, 180]	loss: 0.08377583138644695	accuracy: 97.19%	val_loss: 1.0250891977921128	val_accuracy: 75.22%
[17, 10]	loss: 0.031532193161547185	accuracy: 99.38%	val_loss: 1.0122089106589556	val_accuracy: 77.32%
[17, 20]	loss: 0.05025607838761061	accuracy: 98.44%	val_loss: 0.9769855616614223	val_accuracy: 77.12%
[17, 30]	loss: 0.03934240099042654	accuracy: 98.44%	val_loss: 0.9585792911238968	val_accuracy: 76.72%
[17, 40]	loss: 0.0316031813621521	accuracy: 98.75%	val_loss: 1.0800800444558263	val_accuracy: 77.12%
[17, 50]	loss: 0.039478436158969996	accuracy: 99.38%	val_loss: 1.0158325349912047	val_accuracy: 77.72%
[17, 60]	loss: 0.016844837041571737	accuracy: 100.00%	val_loss: 0.9743807064369321	val_accuracy: 77.22%

[17, 70]	loss: 0.04107870860025287	accuracy: 99.06%	val_loss: 0.9758161171339452	val_accuracy: 78.12%
[17, 80]	loss: 0.037335478514432904	accuracy: 99.06%	val_loss: 0.9972122824983671	val_accuracy: 76.72%
[17, 90]	loss: 0.056230609957128766	accuracy: 99.06%	val_loss: 0.9688113705487922	val_accuracy: 76.62%
[17, 100]	loss: 0.07869342397898435	accuracy: 97.50%	val_loss: 1.132048450410366	val_accuracy: 75.22%
[17, 110]	loss: 0.07434711847454309	accuracy: 96.56%	val_loss: 1.0922920489683747	val_accuracy: 74.83%
[17, 120]	loss: 0.10736699122935534	accuracy: 96.25%	val_loss: 1.056011184118688	val_accuracy: 75.02%
[17, 130]	loss: 0.08934989795088769	accuracy: 96.88%	val_loss: 1.0902506867423654	val_accuracy: 76.22%
[17, 140]	loss: 0.09988730512559414	accuracy: 96.88%	val_loss: 1.044658021768555	val_accuracy: 75.82%
[17, 150]	loss: 0.059851275850087406	accuracy: 97.50%	val_loss: 1.080518466886133	val_accuracy: 74.43%
[17, 160]	loss: 0.1720011468976736	accuracy: 94.69%	val_loss: 1.2388830911368132	val_accuracy: 73.13%
[17, 170]	loss: 0.35937573984265325	accuracy: 89.69%	val_loss: 1.0235389303416014	val_accuracy: 73.93%
[17, 180]	loss: 0.49498351365327836	accuracy: 85.00%	val_loss: 1.5344038009643555	val_accuracy: 71.13%
[18, 10]	loss: 0.27713136896491053	accuracy: 90.94%	val_loss: 1.1420267820358276	val_accuracy: 72.03%
[18, 20]	loss: 0.2836598392575979	accuracy: 90.94%	val_loss: 1.0542653407901525	val_accuracy: 74.33%
[18, 30]	loss: 0.24672612585127354	accuracy: 92.81%	val_loss: 1.2944398950785398	val_accuracy: 72.13%
[18, 40]	loss: 0.2642848134040833	accuracy: 91.56%	val_loss: 1.0194151289761066	val_accuracy: 73.03%
[18, 50]	loss: 0.24620438143610954	accuracy: 90.94%	val_loss: 0.9962736181914806	val_accuracy: 74.53%
[18, 60]	loss: 0.24055227190256118	accuracy: 90.62%	val_loss: 0.8965057395398617	val_accuracy: 74.83%
[18, 70]	loss: 0.21650175675749778	accuracy: 92.19%	val_loss: 0.9912459054030478	val_accuracy: 74.73%
[18, 80]	loss: 0.19642213732004166	accuracy: 91.88%	val_loss: 0.9550196873024106	val_accuracy: 74.93%
[18, 90]	loss: 0.19558829292654992	accuracy: 93.44%	val_loss: 0.9317968487157486	val_accuracy: 74.53%
[18, 100]	loss: 0.16296535283327102	accuracy: 94.06%	val_loss: 0.9167855037376285	val_accuracy: 75.12%
[18, 110]	loss: 0.12574661746621132	accuracy: 95.62%	val_loss: 0.9504032181575894	val_accuracy: 74.33%
[18, 120]	loss: 0.18772648870944977	accuracy: 92.81%	val_loss: 0.9152087979018688	val_accuracy: 74.63%
[18, 130]	loss: 0.10575482733547688	accuracy: 96.88%	val_loss: 0.8950594253838062	val_accuracy: 75.42%
[18, 140]	loss: 0.15047001596540213	accuracy: 93.12%	val_loss: 0.8619090062566102	val_accuracy: 76.22%
[18, 150]	loss: 0.1725964542478323	accuracy: 93.12%	val_loss: 0.9407825428061187	val_accuracy: 77.12%
[18, 160]	loss: 0.17610104344785213	accuracy: 93.75%	val_loss: 0.8940273718908429	val_accuracy: 73.43%
[18, 170]	loss: 0.13622300736606122	accuracy: 95.94%	val_loss: 0.9119768664240837	val_accuracy: 75.02%
[18, 180]	loss: 0.09605786856263876	accuracy: 97.19%	val_loss: 0.9495667042210698	val_accuracy: 75.02%
[19, 10]	loss: 0.06632480975240469	accuracy: 97.50%	val_loss: 0.943167963065207	val_accuracy: 74.33%
[19, 20]	loss: 0.06850914703682065	accuracy: 97.19%	val_loss: 0.944533015601337	val_accuracy: 74.83%
[19, 30]	loss: 0.04172752490267158	accuracy: 99.06%	val_loss: 0.9115141718648374	val_accuracy: 75.72%
[19, 40]	loss: 0.05607278198003769	accuracy: 98.12%	val_loss: 0.961899477057159	val_accuracy: 76.52%
[19, 50]	loss: 0.044240222033113244	accuracy: 99.06%	val_loss: 1.0268766712397337	val_accuracy: 76.02%
[19, 60]	loss: 0.029015377536416052	accuracy: 99.38%	val_loss: 0.9898779196664691	val_accuracy: 75.52%
[19, 70]	loss: 0.03324980125762522	accuracy: 98.75%	val_loss: 1.0303873028606176	val_accuracy: 75.42%
[19, 80]	loss: 0.03493372611701488	accuracy: 99.38%	val_loss: 1.0429119123145938	val_accuracy: 74.73%
[19, 90]	loss: 0.0852946074679494	accuracy: 95.94%	val_loss: 0.9945209515281022	val_accuracy: 75.72%
[19, 100]	loss: 0.08185158669948578	accuracy: 96.25%	val_loss: 0.9322339901700616	val_accuracy: 77.62%
[19, 110]	loss: 0.042408393882215026	accuracy: 97.81%	val_loss: 0.9589993450790644	val_accuracy: 76.62%
[19, 120]	loss: 0.12130515240132808	accuracy: 95.94%	val_loss: 1.1311148665845394	val_accuracy: 76.12%



```

[19, 130] loss: 0.0667321115732193 | accuracy: 97.19% | val_loss: 0.9950402351096272 | val_accuracy: 75.82%
[19, 140] loss: 0.055365361366420986 | accuracy: 99.06% | val_loss: 0.8922380227595568 | val_accuracy: 76.12%
[19, 150] loss: 0.05348079111427069 | accuracy: 98.44% | val_loss: 0.9776623672805727 | val_accuracy: 77.62%
[19, 160] loss: 0.04324442422948778 | accuracy: 98.75% | val_loss: 0.9514461010694504 | val_accuracy: 77.12%
[19, 170] loss: 0.038050378765910865 | accuracy: 98.75% | val_loss: 0.8962106574326754 | val_accuracy: 76.22%
[19, 180] loss: 0.026427942141890527 | accuracy: 99.69% | val_loss: 0.9542742553167045 | val_accuracy: 76.82%
[20, 10] loss: 0.03255069092847407 | accuracy: 99.38% | val_loss: 1.0315787252038717 | val_accuracy: 77.42%
[20, 20] loss: 0.025865146541036665 | accuracy: 99.06% | val_loss: 1.0078303813934326 | val_accuracy: 76.82%
[20, 30] loss: 0.022528247395530342 | accuracy: 99.69% | val_loss: 0.9665228994563222 | val_accuracy: 77.62%
[20, 40] loss: 0.020528855826705696 | accuracy: 100.00% | val_loss: 0.9136578496545553 | val_accuracy: 77.62%
[20, 50] loss: 0.015688369143754245 | accuracy: 99.69% | val_loss: 0.9663622798398137 | val_accuracy: 78.62%
[20, 60] loss: 0.02058320282958448 | accuracy: 99.38% | val_loss: 0.9400372400414199 | val_accuracy: 78.32%
[20, 70] loss: 0.019112458685413003 | accuracy: 99.69% | val_loss: 0.9026935258880258 | val_accuracy: 77.82%
[20, 80] loss: 0.017089965753257275 | accuracy: 100.00% | val_loss: 0.9114672979339957 | val_accuracy: 78.32%
[20, 90] loss: 0.03204528233036399 | accuracy: 98.75% | val_loss: 0.9636399227310903 | val_accuracy: 78.22%
[20, 100] loss: 0.022423171543050558 | accuracy: 99.38% | val_loss: 0.9707640549167991 | val_accuracy: 77.02%
[20, 110] loss: 0.024497645953670145 | accuracy: 99.38% | val_loss: 0.9674847521819174 | val_accuracy: 78.12%
[20, 120] loss: 0.022212970815598965 | accuracy: 99.69% | val_loss: 1.0025746310129762 | val_accuracy: 78.02%
[20, 130] loss: 0.03648566054180265 | accuracy: 98.75% | val_loss: 0.9578100708313286 | val_accuracy: 76.42%
[20, 140] loss: 0.025730058387853205 | accuracy: 99.38% | val_loss: 0.9828051831573248 | val_accuracy: 77.02%
[20, 150] loss: 0.042630175570957364 | accuracy: 98.75% | val_loss: 0.9330647288879845 | val_accuracy: 77.42%
[20, 160] loss: 0.015606430545449256 | accuracy: 100.00% | val_loss: 0.9695392269641161 | val_accuracy: 76.62%
[20, 170] loss: 0.018678332108538597 | accuracy: 99.69% | val_loss: 1.0043926732614636 | val_accuracy: 77.52%
[20, 180] loss: 0.025479060038924216 | accuracy: 99.06% | val_loss: 1.0422351500019431 | val_accuracy: 77.92%

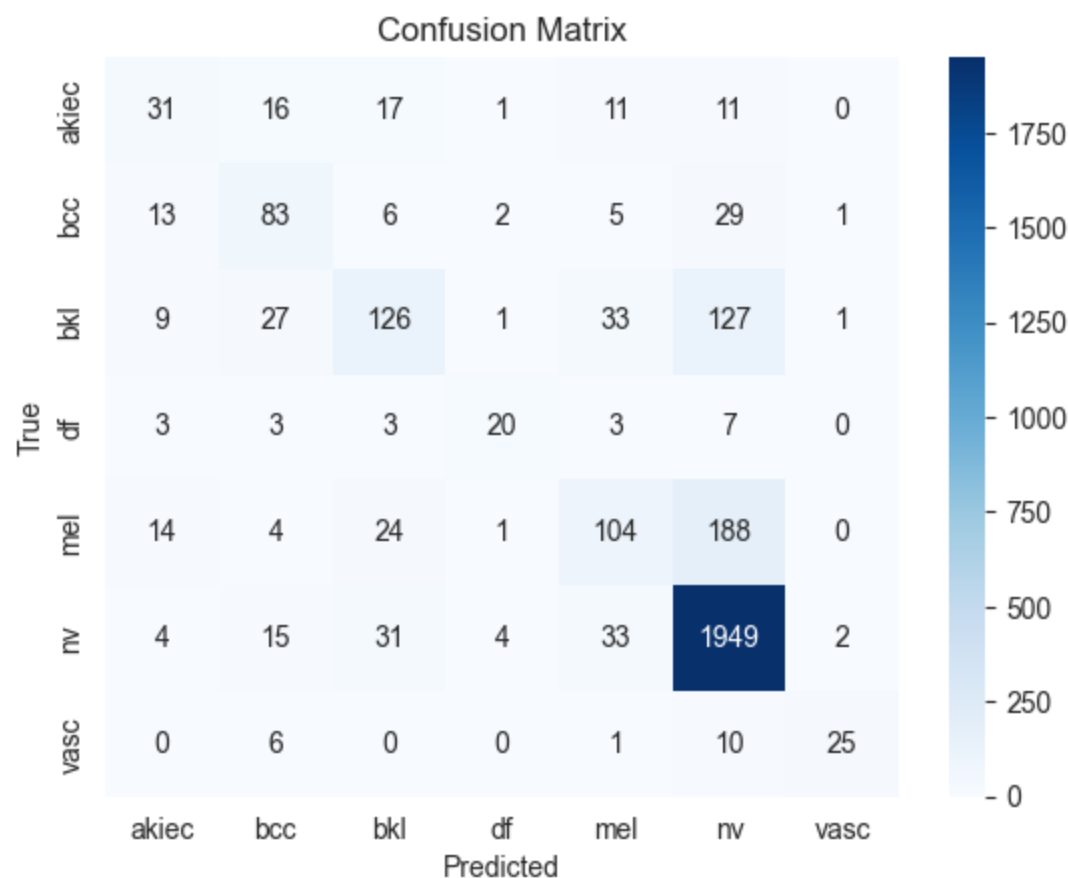
```

```

In [27]: number_to_label_array = testloader.dataset.dataset.number_to_label_array
         # Testing model
         test_bonus_net(model, testloader, number_to_label_array, loss=loss, device=device, model_name="bonus")

```

Test accuracy: 77.83% | Test loss: 1.185665249824524




## Observations

- I would say that the model is performing better than the Complex model, the recall on `nv` is better, and the precision on `bkl`, `mel` is better
- the diagonal (correct predictions) is more concentrated than the Complex model, which is a good thing
- the precision on `nv` is also way better, meaning the demographic data is helping the model to distinguish between `nv` and the other labels

```
In [28]: # Plotting losses
x_axis = list(map(lambda x: x * 10, range(len(train_loss))))
plt.figure(figsize=(10, 6))
```


```
plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, val_loss, label='Validation Loss')
xticks = np.arange(x_axis[0], x_axis[-1] + 1, 400.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid()
plt.savefig('results/losses/bonus_loss.png')
plt.show()
```



## Observations

- overfitting starts at around 1400 batches this time
- it's more "jagged" than the Complex model, maybe decreasing the learning rate might help
- we have the weird spike but this time instead of 2400 batches, it's at 3000, probably the reason for the offset is the same reason as for the offset of the overfitting line. I guess that the learning rate is too high and affects the model's performance.

```
In [29]: # Plotting accuracies
x_axis = list(map(lambda x: x * 10, range(len(train_accuracy))))
plt.figure(figsize=(10, 6))
plt.plot(x_axis, train_accuracy, label='Training Accuracy')
plt.plot(x_axis, val_accuracy, label='Validation Accuracy')
xticks = np.arange(x_axis[0], x_axis[-1] + 1, 200.0)
plt.xticks([*xticks, x_axis[-1]])
plt.xlabel('Batch Number')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.grid()
plt.savefig('results/accuracies/bonus_accuracy.png')
plt.show()
```



## Observations

- The model starts to overfit at around 1400 batches, we saw that on the loss plot as well
- we can see the weird spike we observed on the loss plot
- accuracy stays idle at around 80%